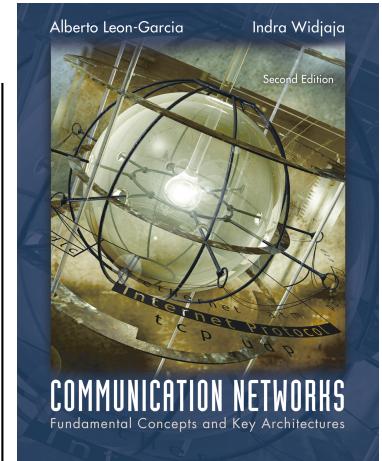
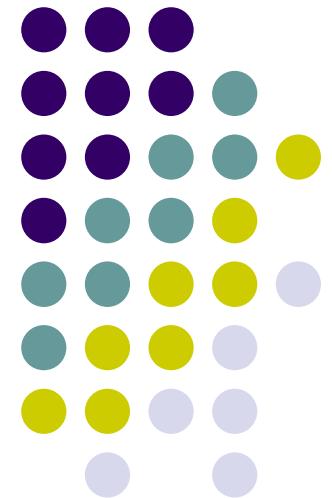


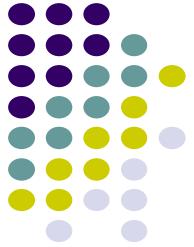
Chapter 2

Applications and Layered Architectures



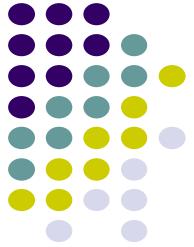
Protocols, Services & Layering





Layers, Services & Protocols

- The overall communications process between two or more machines connected across one or more networks is very complex
- *Layering* partitions related communications functions into groups that are manageable
- Each layer provides a **service** to the layer above
- Each layer operates according to a *protocol*

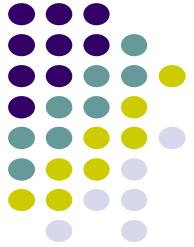


Layer

- A set of related communication functions that can be managed and grouped together

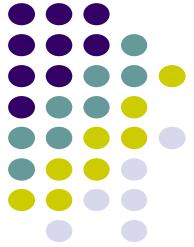
Examples of Layers

- Application Layer: communications functions that are used by application programs
 - HTTP (Browsing), DNS (URL to IP address), SMTP (email)
- Transport Layer: end-to-end communications *between two processes* in two machines
 - TCP Reliable Stream Service, UDP Datagram Service
- Network Layer: communications *between machines*
 - Internet Protocol (IP) best effort service



Protocols

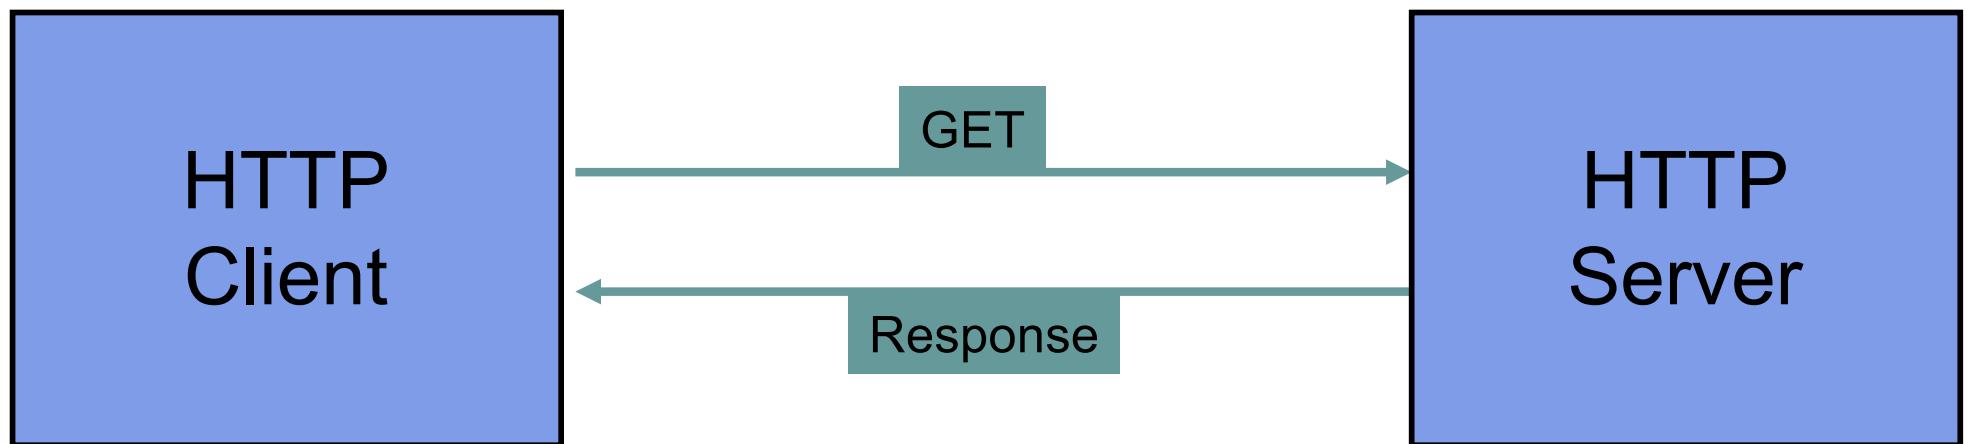
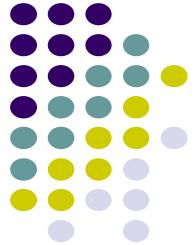
- A *protocol* is a set of rules that governs how two or more communicating entities *in a layer* are to interact
- *Messages* that can be sent and received
- *Actions* that are to be taken when a certain event occurs, e.g. sending or receiving messages, expiry of timers
- **The purpose of a protocol is to provide a service to the layer above**



Example: HTTP

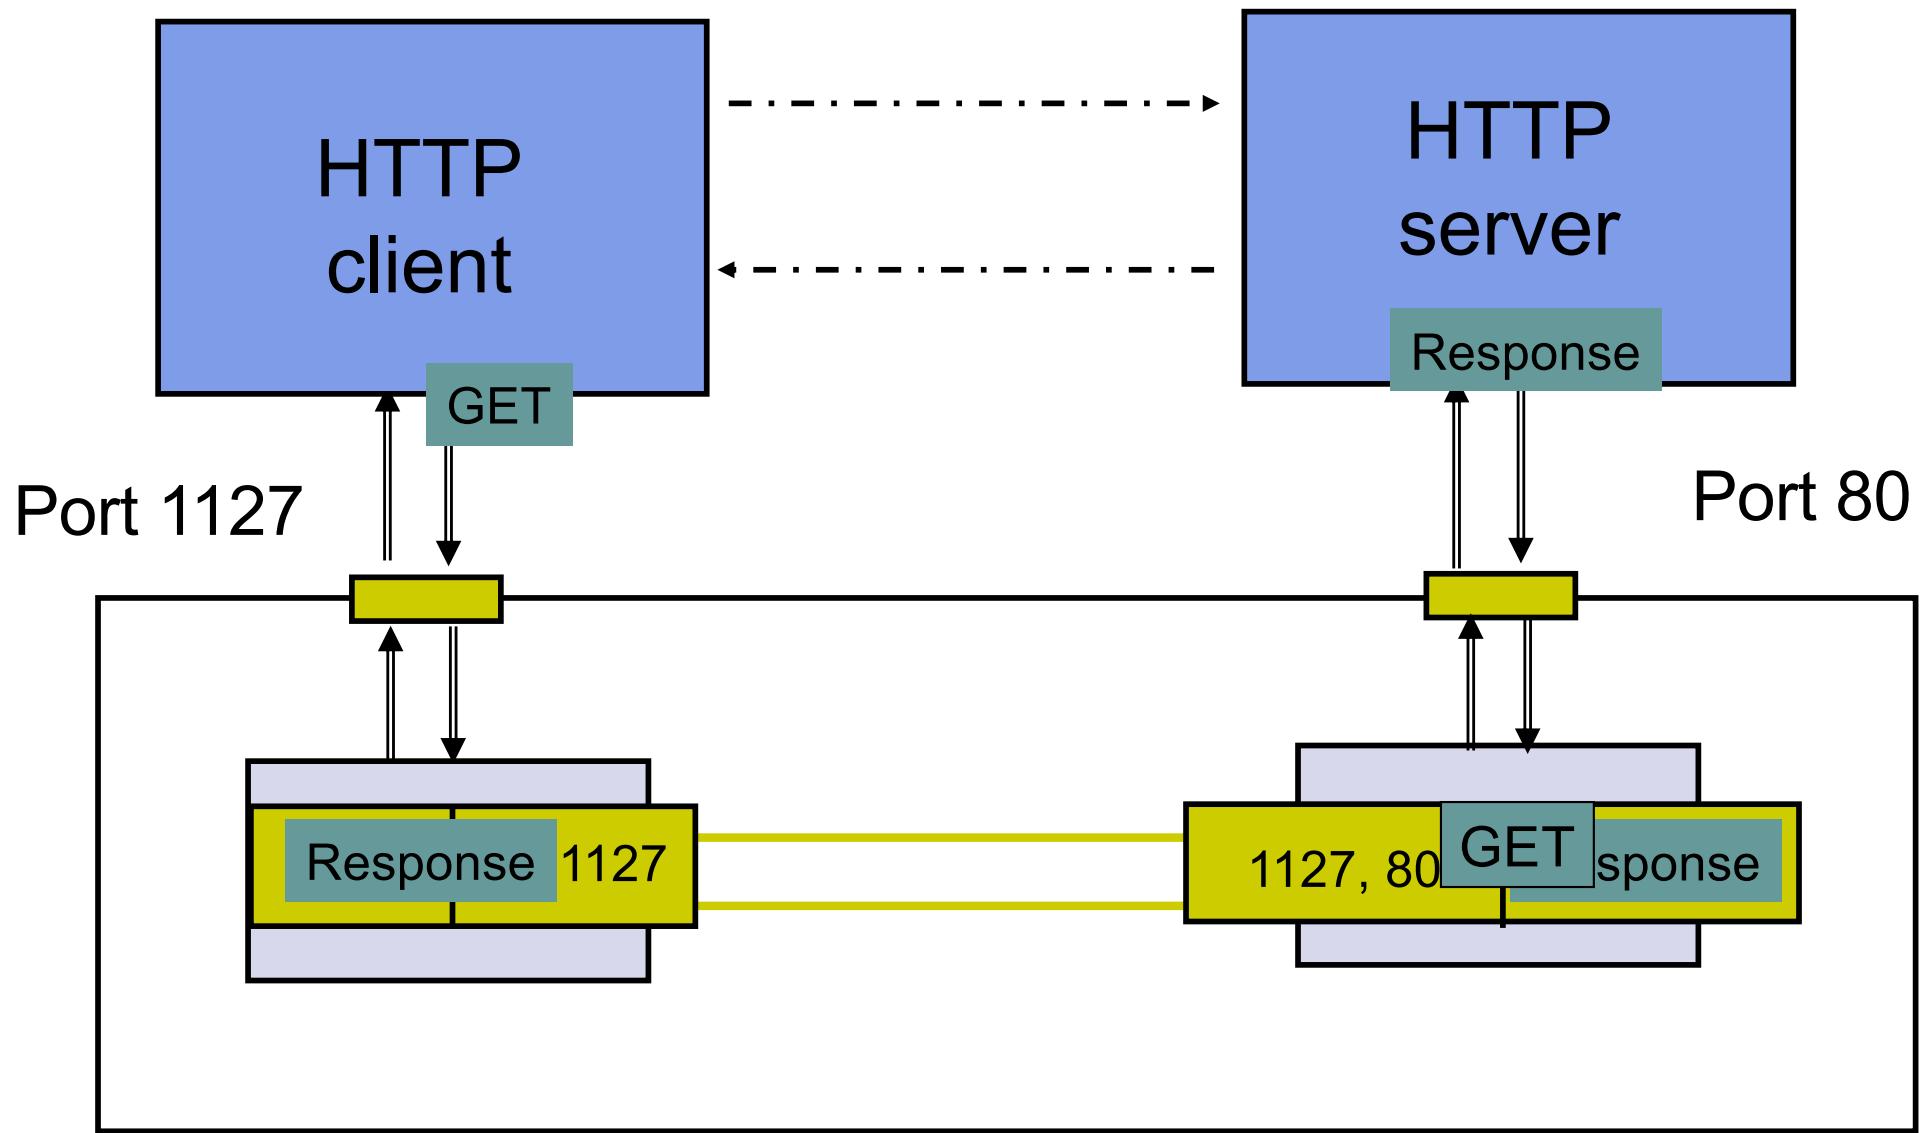
- HTTP is an application layer protocol
- Retrieves documents on behalf of a browser application program
- HTTP specifies fields in request messages and response messages
 - Request types; Response codes
 - Content type, options, cookies, ...
- HTTP specifies actions to be taken upon receipt of certain messages

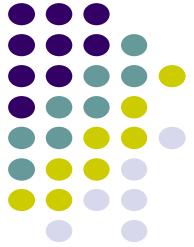
HTTP Protocol



- HTTP provides service to a browsing application
- HTTP messages exchanged between HTTP client and HTTP server
- HTTP client and server are processes running in two different machines across the Internet
- HTTP uses the reliable stream transfer service provided by TCP (Transmission Control Protocol)

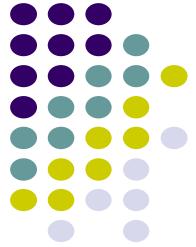
HTTP uses service of TCP





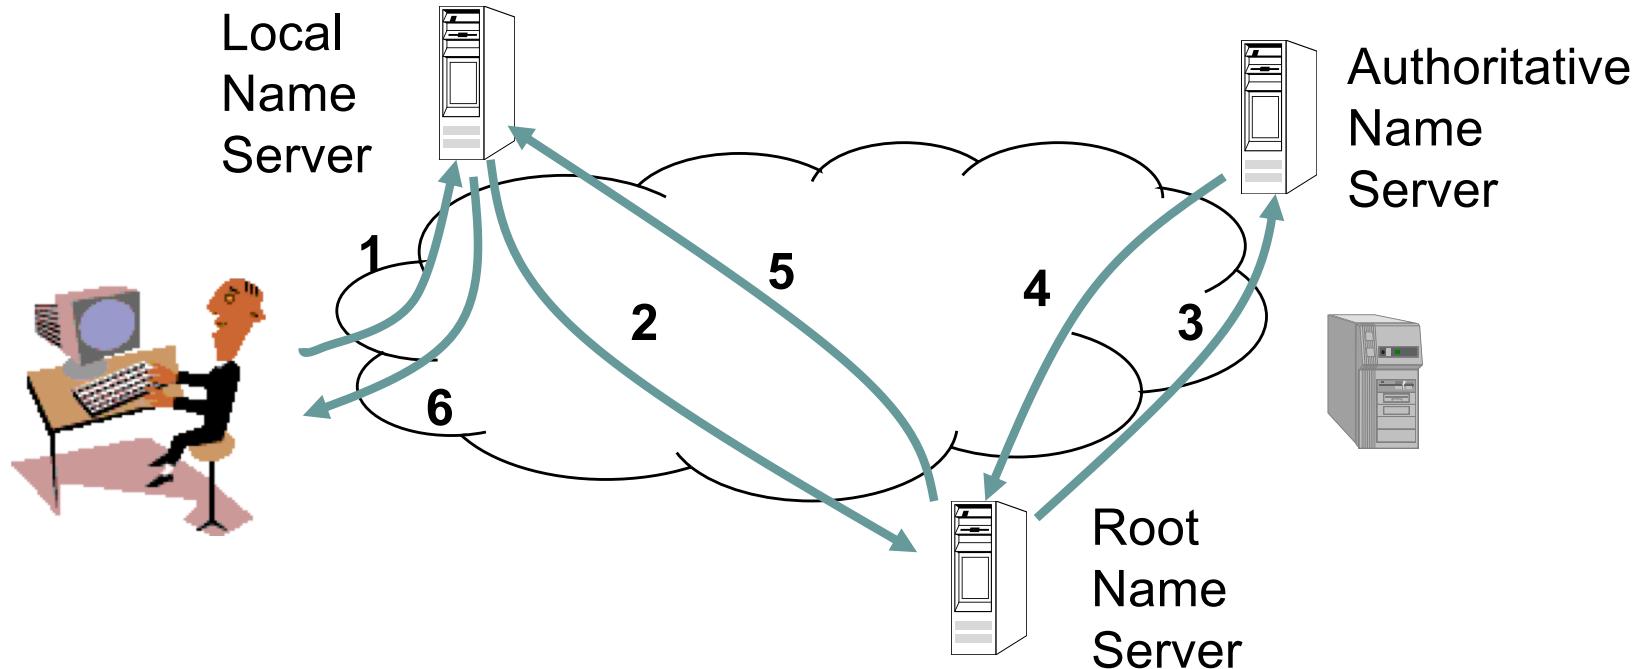
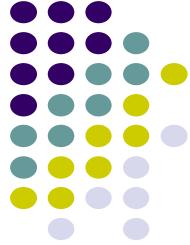
Example: TCP

- TCP is a transport layer protocol
- Provides *reliable byte stream service between two processes* in two computers across the Internet
- Sequence numbers keep track of the bytes that have been transmitted and received
- Error detection and retransmission used to recover from transmission errors, losses, duplications
- TCP is *connection-oriented*: the sender and receiver must first establish an association and set initial sequence numbers before data is transferred
- Connection ID is specified uniquely by
(send port #, send IP address, receive port #, receiver IP address)
- HTTP server has well-known port # 80

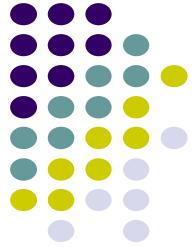


Example: DNS Protocol

- DNS (Domain Name System) is an application layer protocol
- DNS is a distributed database that resides in multiple machines in the Internet
- DNS protocol allows queries of different types
 - Name-to-address or Address-to-name
 - Mail exchange
- DNS usually involves short messages and so uses service provided by UDP
- Well-known port 53

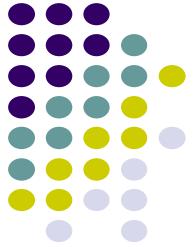


- Local Name Server: resolve frequently-used names
 - University department, ISP
 - Contacts Root Name server if it cannot resolve query
- Root Name Servers: 13 globally (actually 100's)
 - Resolves query or refers query to Authoritative Name Server
- Authoritative Name Server: last resort
 - Every machine must register its address with at least two authoritative name servers

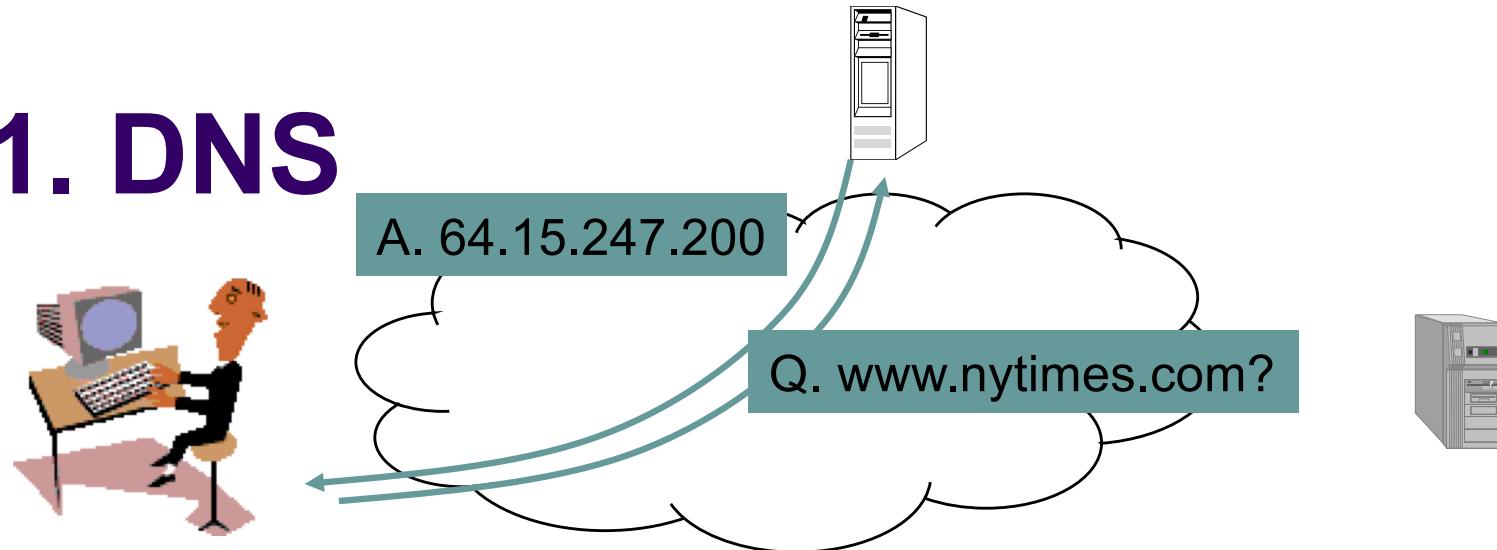


Web Browsing Application

- World Wide Web allows users to access resources (i.e. documents) located in computers connected to the Internet
- Documents prepared using HyperText Markup Language (HTML)
- A browser application program is used to access the web
- The browser displays HTML documents that include *links* to other documents
- Each link references a *Uniform Resource Locator* (URL) that gives the name of the machine and the location of the given document
- Let's see what happens when a user clicks on a link

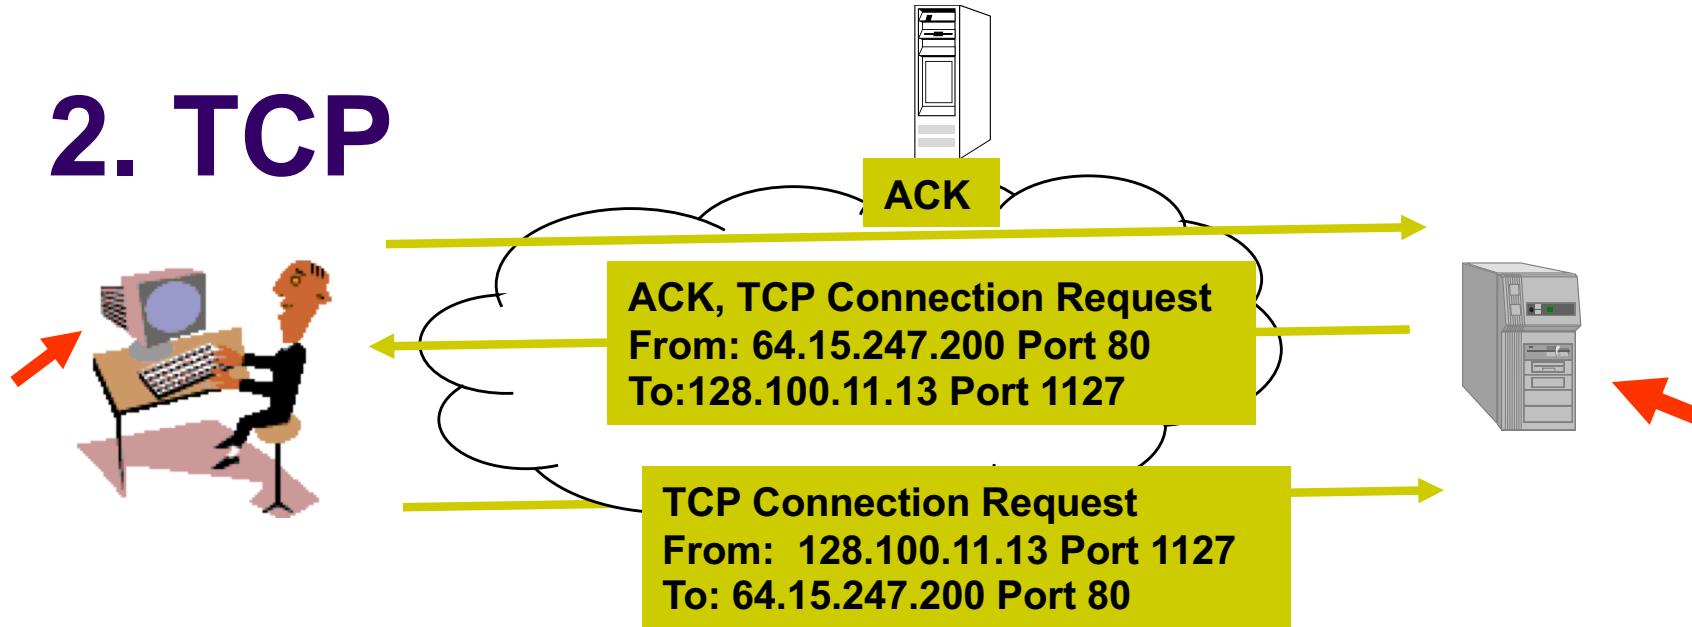
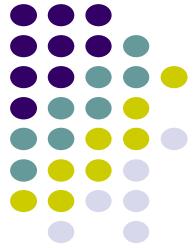


1. DNS



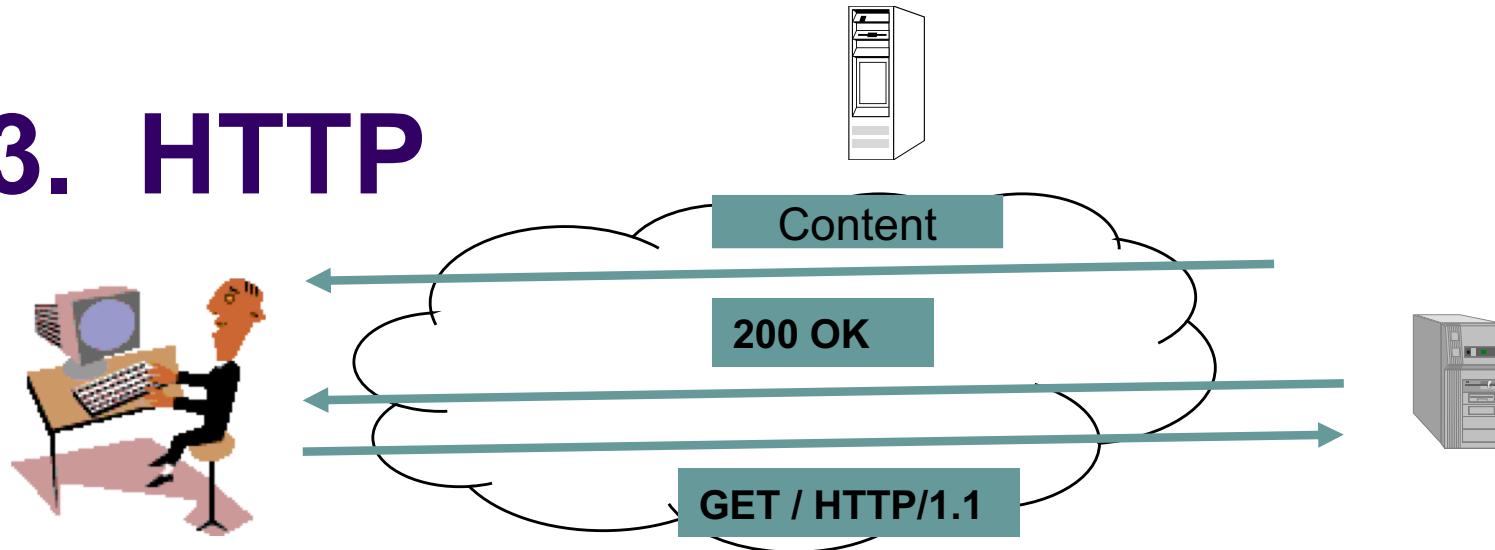
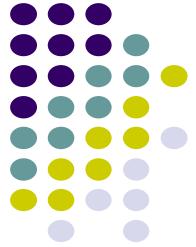
- User clicks on <http://www.nytimes.com/>
- URL contains Internet name of machine (www.nytimes.com), but not Internet address
- Internet needs Internet address to send packets to a machine
- Browser software uses DNS protocol to send query for Internet address
- DNS system responds with Internet address

2. TCP

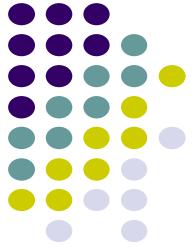


- Browser uses HTTP to send request for document
- HTTP server waits for requests by listening to a well-known port number (80 for HTTP)
- HTTP client sends request messages through an “ephemeral port number,” e.g. 1127
- HTTP needs a TCP connection between the HTTP client and HTTP server to transfer messages reliably

3. HTTP

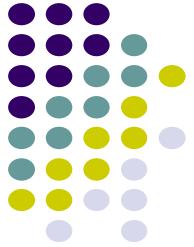


- HTTP client sends its request message: “GET ...”
 - HTTP server sends a status response: “200 OK”
 - HTTP server sends requested file
 - Browser displays document
-
- Clicking a link sets off a chain of events across the Internet!
 - DNS, TCP, HTTP, Browser display



Example: UDP

- UDP is a transport layer protocol
- Provides *best-effort datagram service* between two processes in two computers across the Internet
- Port numbers distinguish various processes in the same machine
- UDP is *connectionless*
- Datagram is sent immediately
- Quick, simple, but not reliable
- Used in DNS, real-time voice,...

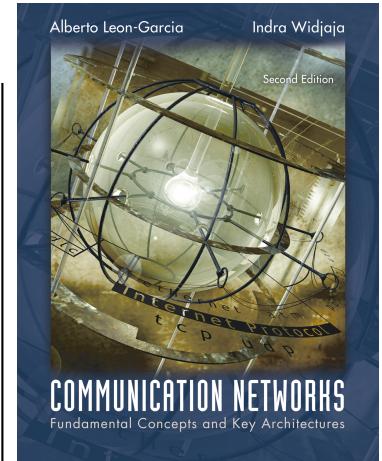


Summary

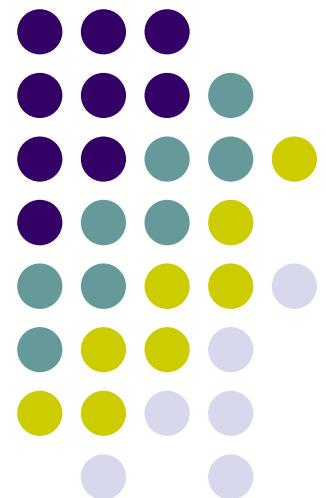
- Layers: related communications functions
 - Application Layer: HTTP, DNS
 - Transport Layer: TCP, UDP
 - Network Layer: IP
- Services: a protocol provides a communications service to the layer above
 - TCP provides connection-oriented reliable stream service
 - UDP provides connectionless datagram service
- Each layer builds on services of lower layers
 - HTTP builds on top of TCP
 - DNS builds on top of UDP
 - TCP and UDP build on top of IP

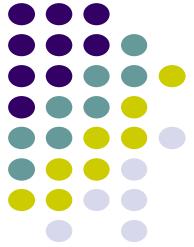
Chapter 2

Applications and Layered Architectures



OSI Reference Model

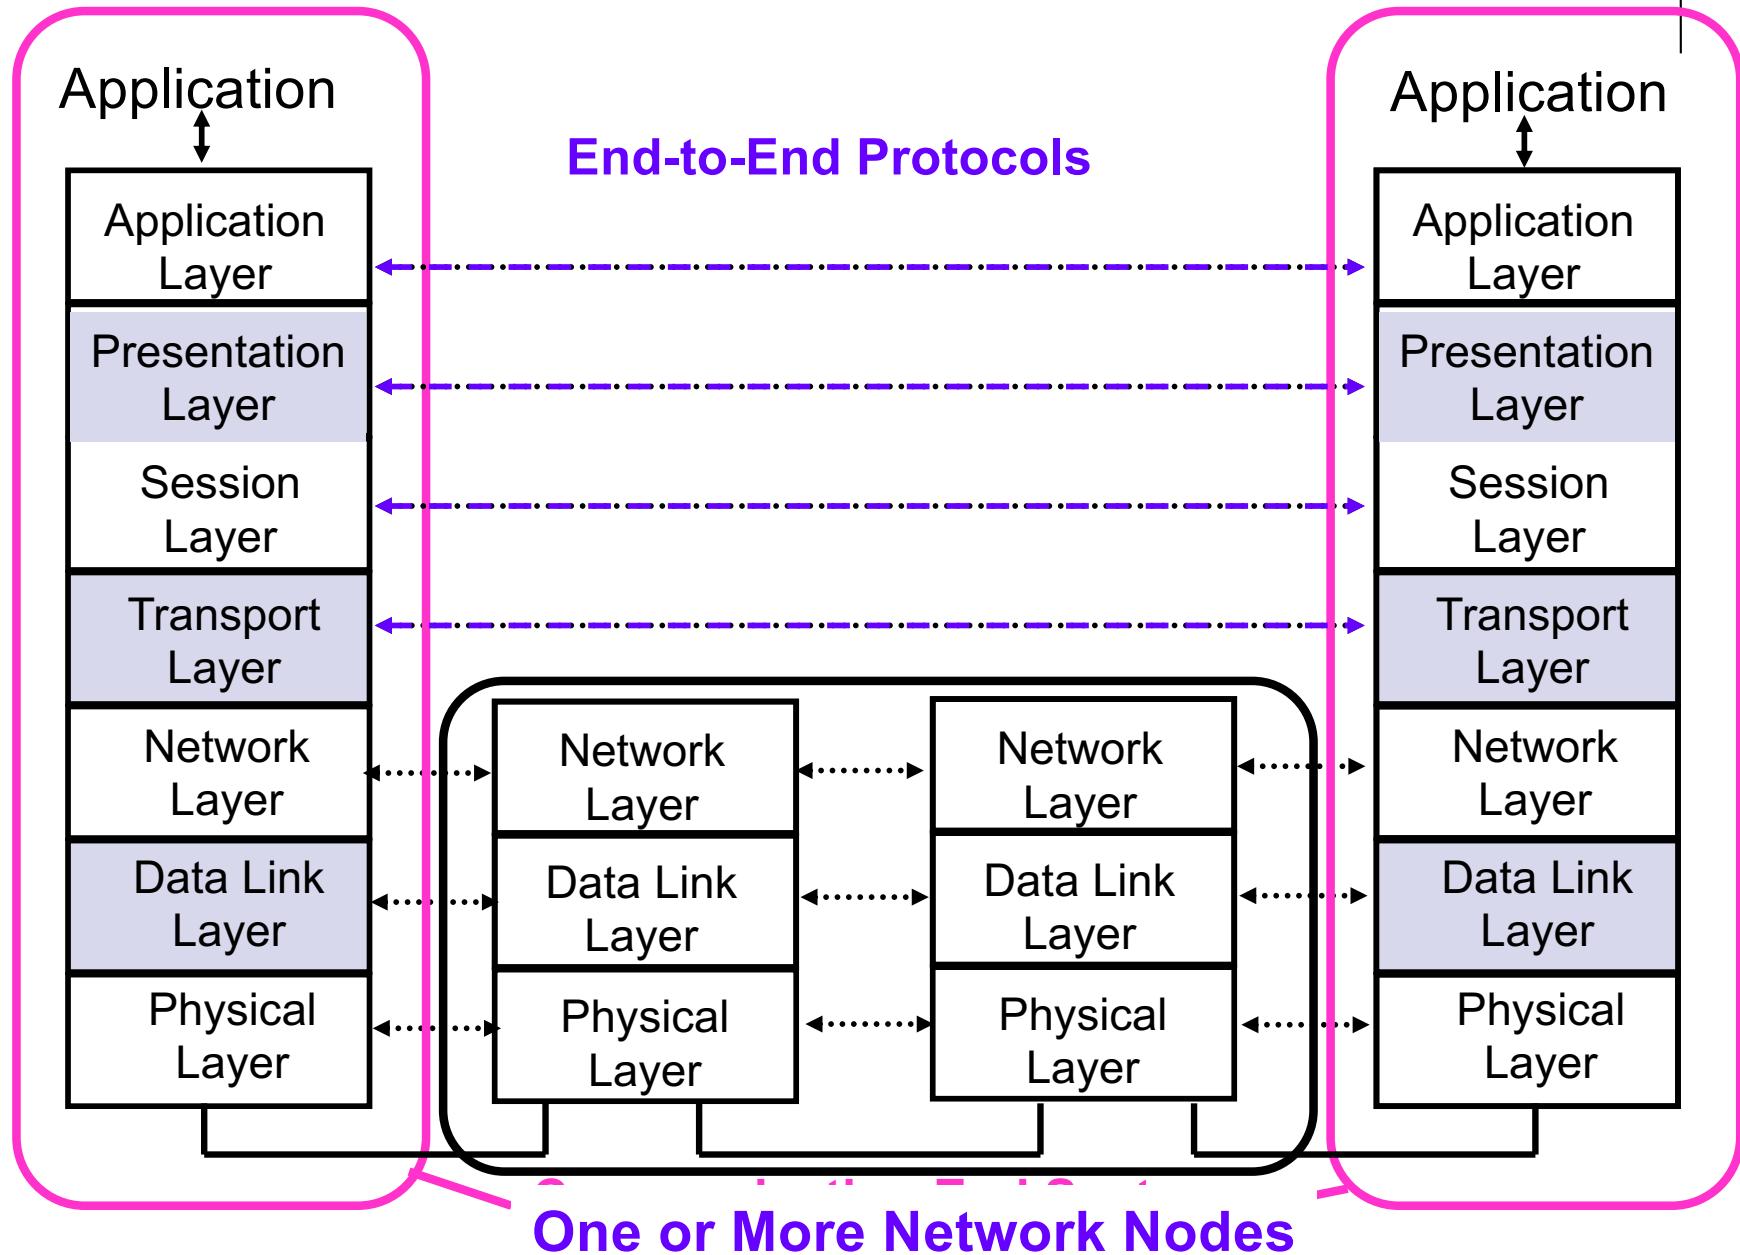




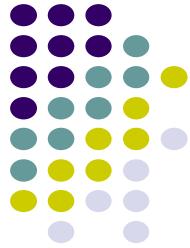
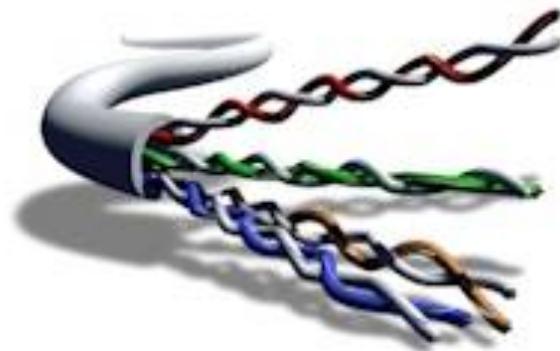
OSI Reference Model

- A seven-layer abstract reference model for a network architecture
- Provides a framework for the development of protocols
- Unified view of layers, protocols, and services which is used in the development of new protocols
- Detailed standards were developed for each layer, but most of these are not in use
- TCP/IP protocols preempted deployment of OSI protocols

7-Layer OSI Reference Model

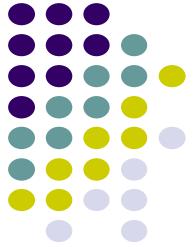


Physical Layer



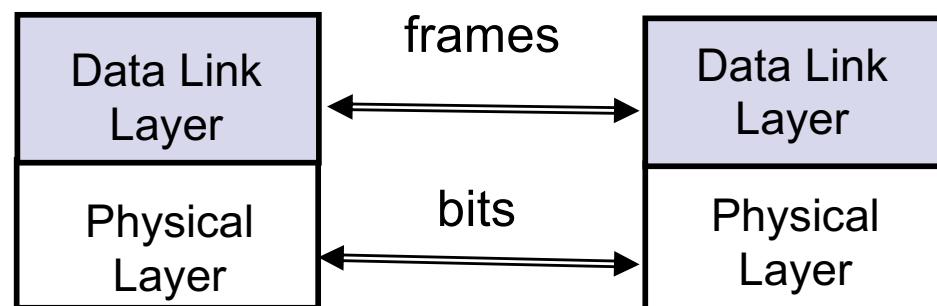
- Transfers **bits** across link
- Definition & specification of the physical aspects of a communications link
 - Mechanical: cable, plugs, pins...
 - Electrical/optical: modulation, signal strength, voltage levels, bit times, ...
 - functional/procedural: how to activate, maintain, and deactivate physical links...
- Ethernet, DSL, cable modem, cell phone,...
- Twisted-pair cable, coaxial cable optical fiber, radio, infrared, ...

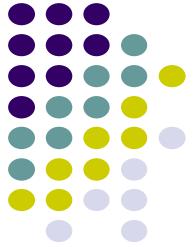




Data Link Layer

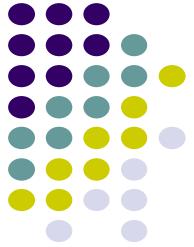
- Transfers **frames** between *directly connected* machines
- Groups bits into frames
- Detection of bit errors; Retransmission of frames
- Activation, maintenance, & deactivation of data link connections
- Access control for local area networks (WIFI, Ethernet)
- Flow control between machines of different speeds





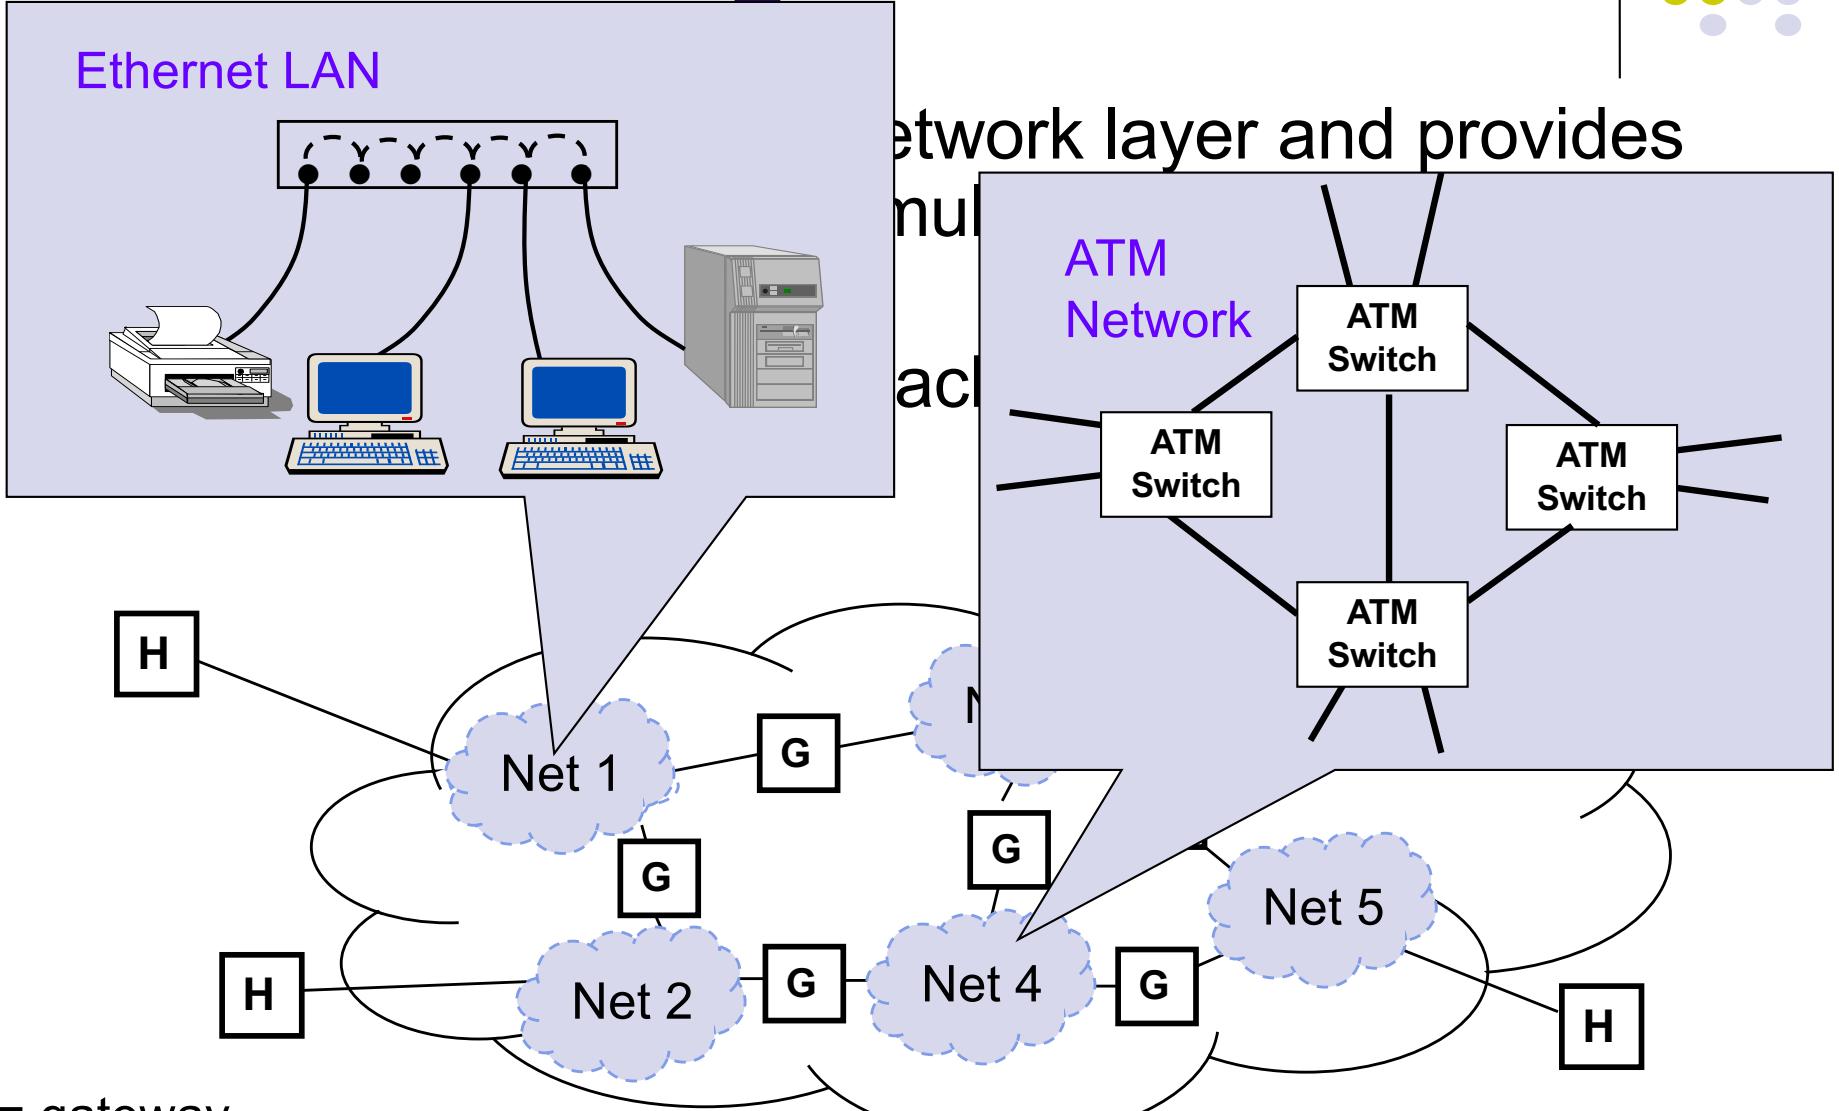
Network Layer

- Transfers *packets* across multiple links and/or multiple networks
- Addressing must scale to large networks
- Nodes *jointly* execute routing algorithm to determine paths across the network
- Forwarding transfers packet along hops in path
- Congestion control to deal with traffic surges
- Connection setup, maintenance, and teardown when connection-based



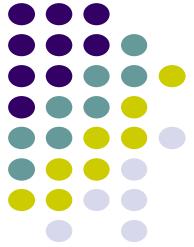
Internetworking

- Ethernet LAN
 - Local network layer and provides multiple access



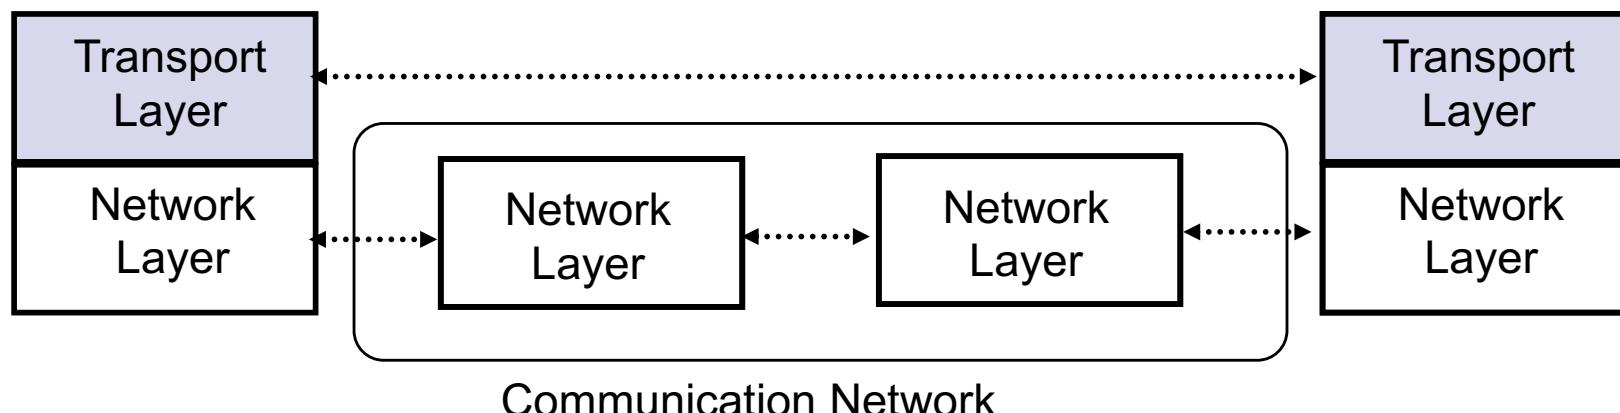
G = gateway

H = host

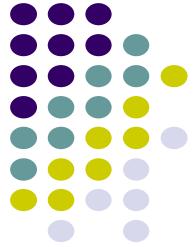


Transport Layer

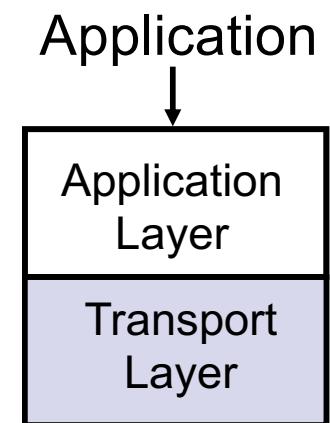
- Transfers data end-to-end from process in a machine to process in another machine
- Reliable stream transfer or quick-and-simple single-block transfer
- Port numbers enable multiplexing
- Message segmentation and reassembly
- Connection setup, maintenance, and release



Application & Upper Layers



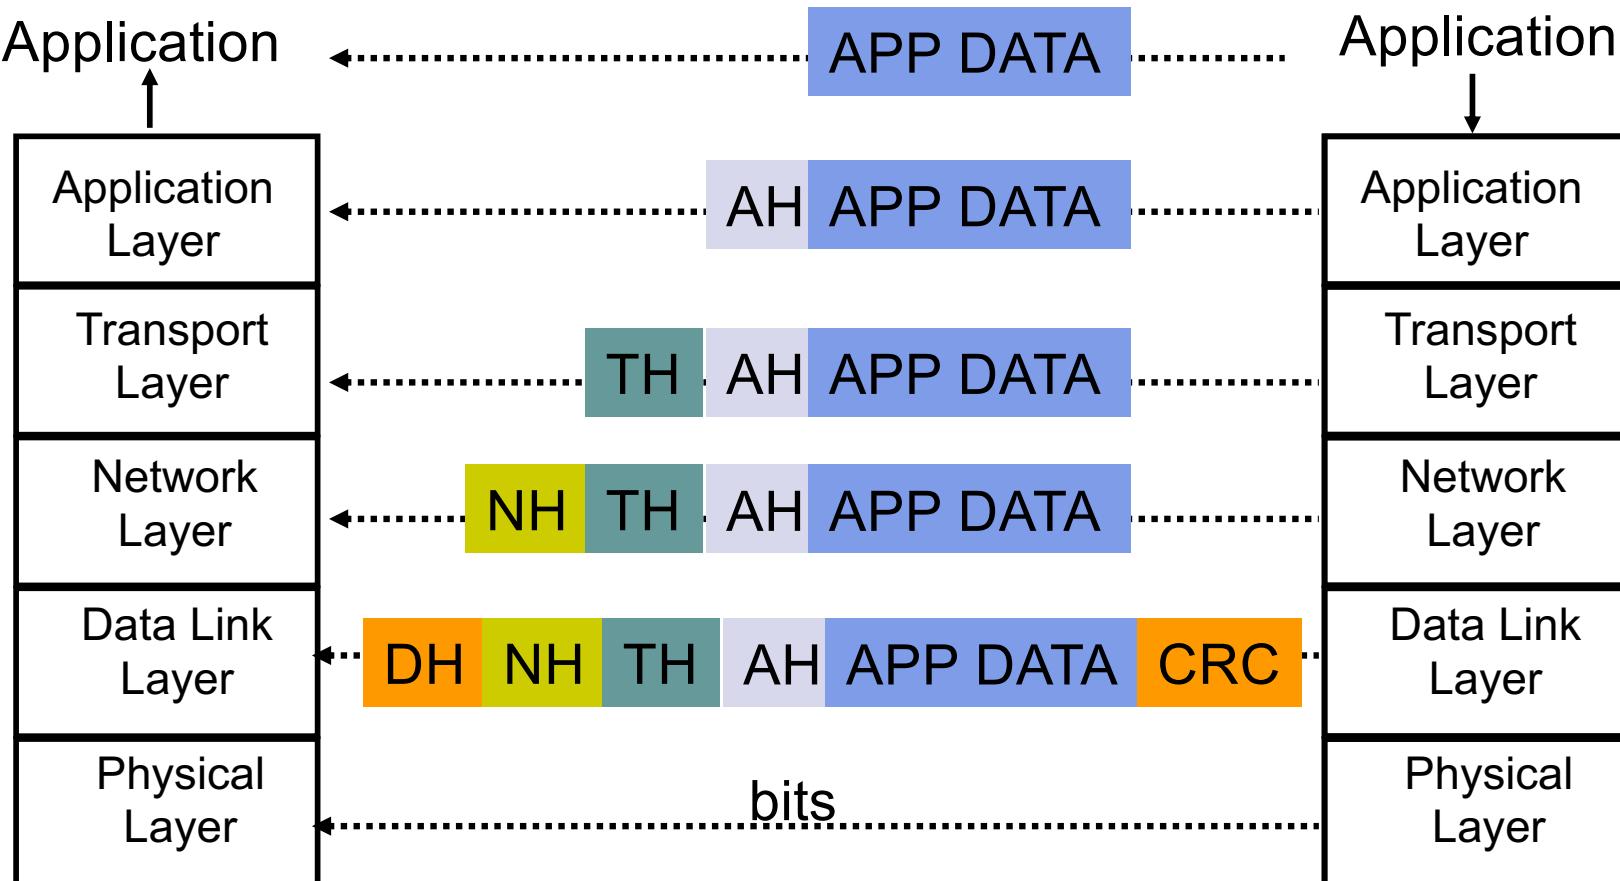
- Application Layer: Provides services that are frequently required by applications: DNS, web access, file transfer, email...
- Presentation Layer: machine-independent representation of data...
- Session Layer: dialog management, recovery from errors, ...
Incorporated into Application Layer



Headers & Trailers

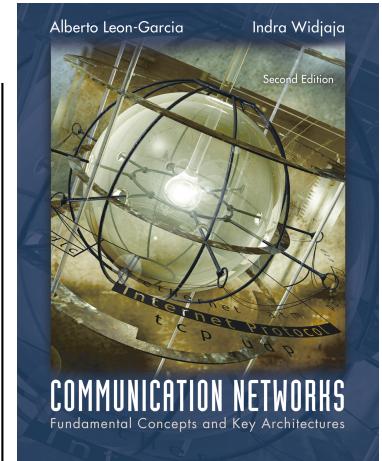


- Each protocol uses a header that carries addresses, sequence numbers, flag bits, length indicators, etc...
- CRC check bits may be appended for error detection

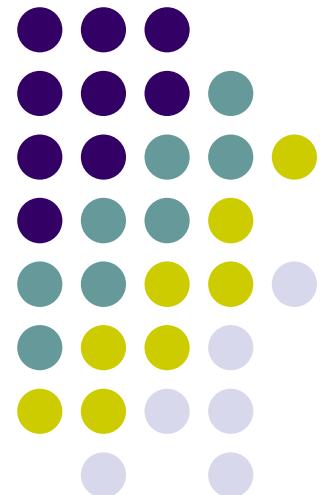


Chapter 2

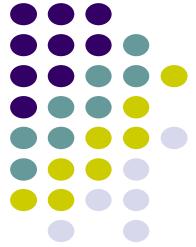
Applications and Layered Architectures



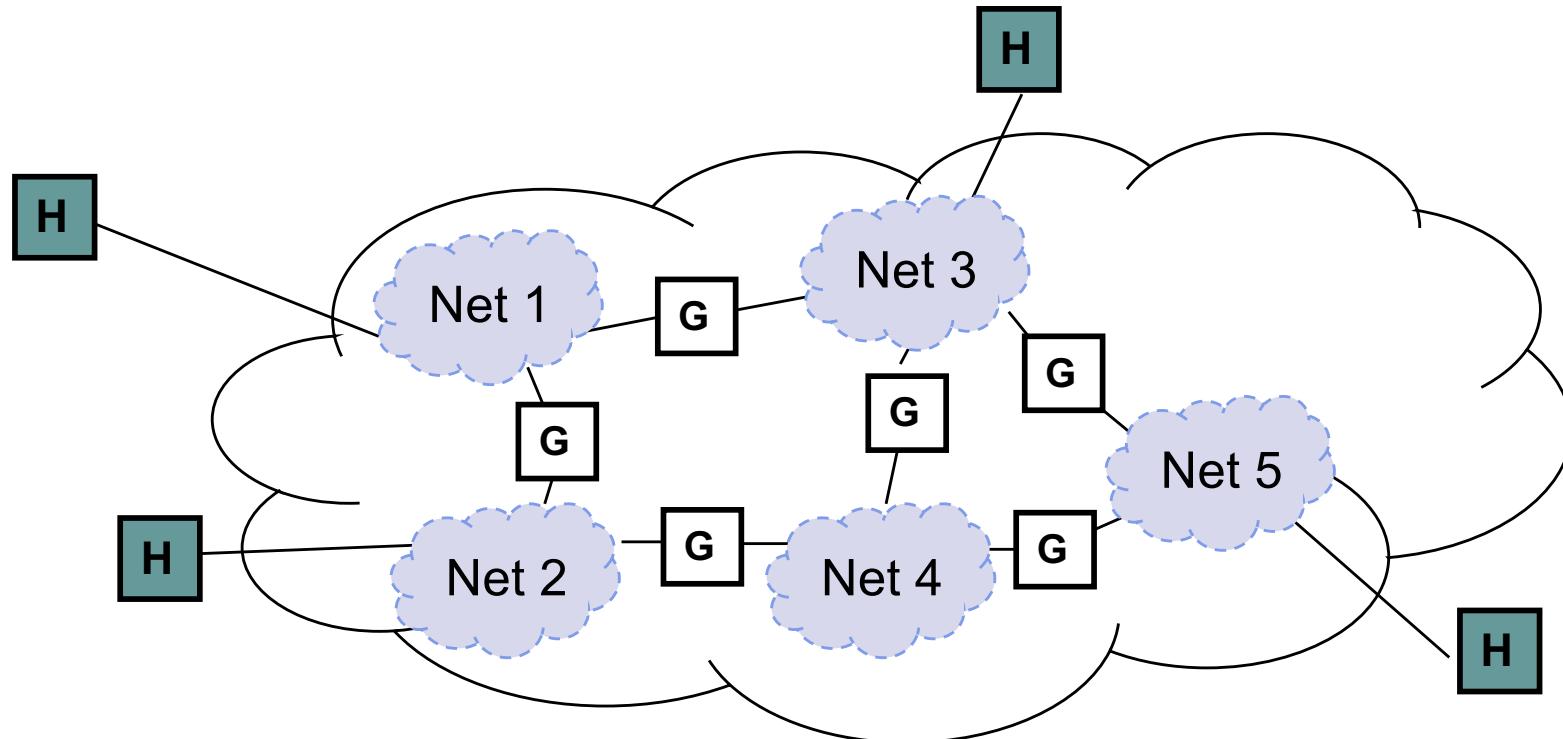
TCP/IP Architecture
How the Layers Work Together



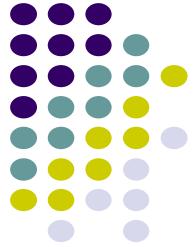
Why Internetworking?



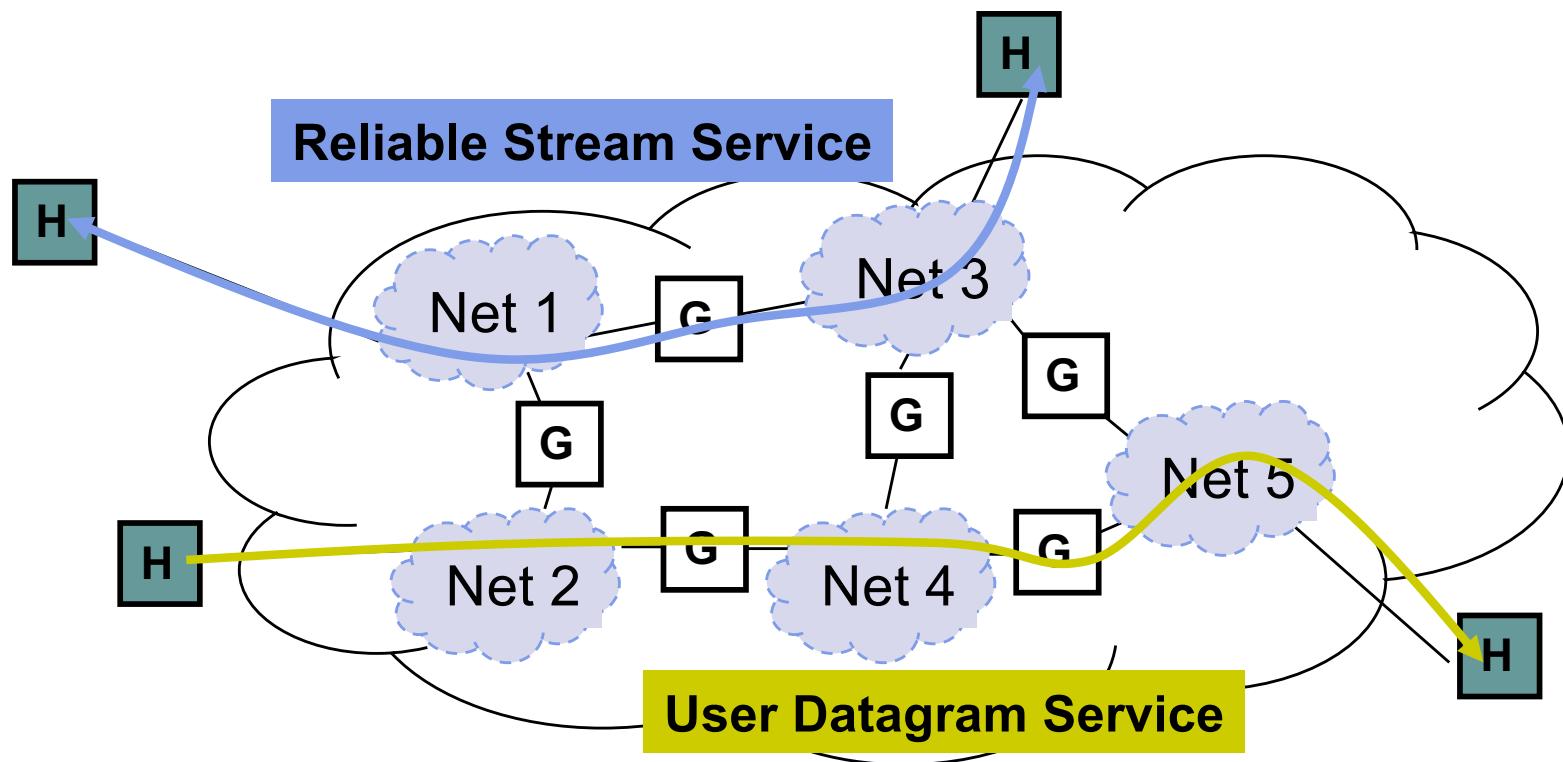
- To build a “network of networks” or internet
 - operating over multiple, coexisting, different network technologies
 - providing ubiquitous connectivity through IP packet transfer
 - achieving huge economies of scale



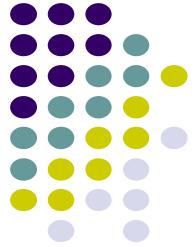
Why Internetworking?



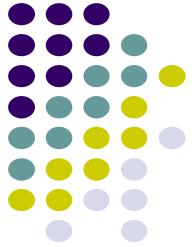
- To provide *universal communication services*
 - independent of underlying network technologies
 - providing common interface to user applications



Why Internetworking?

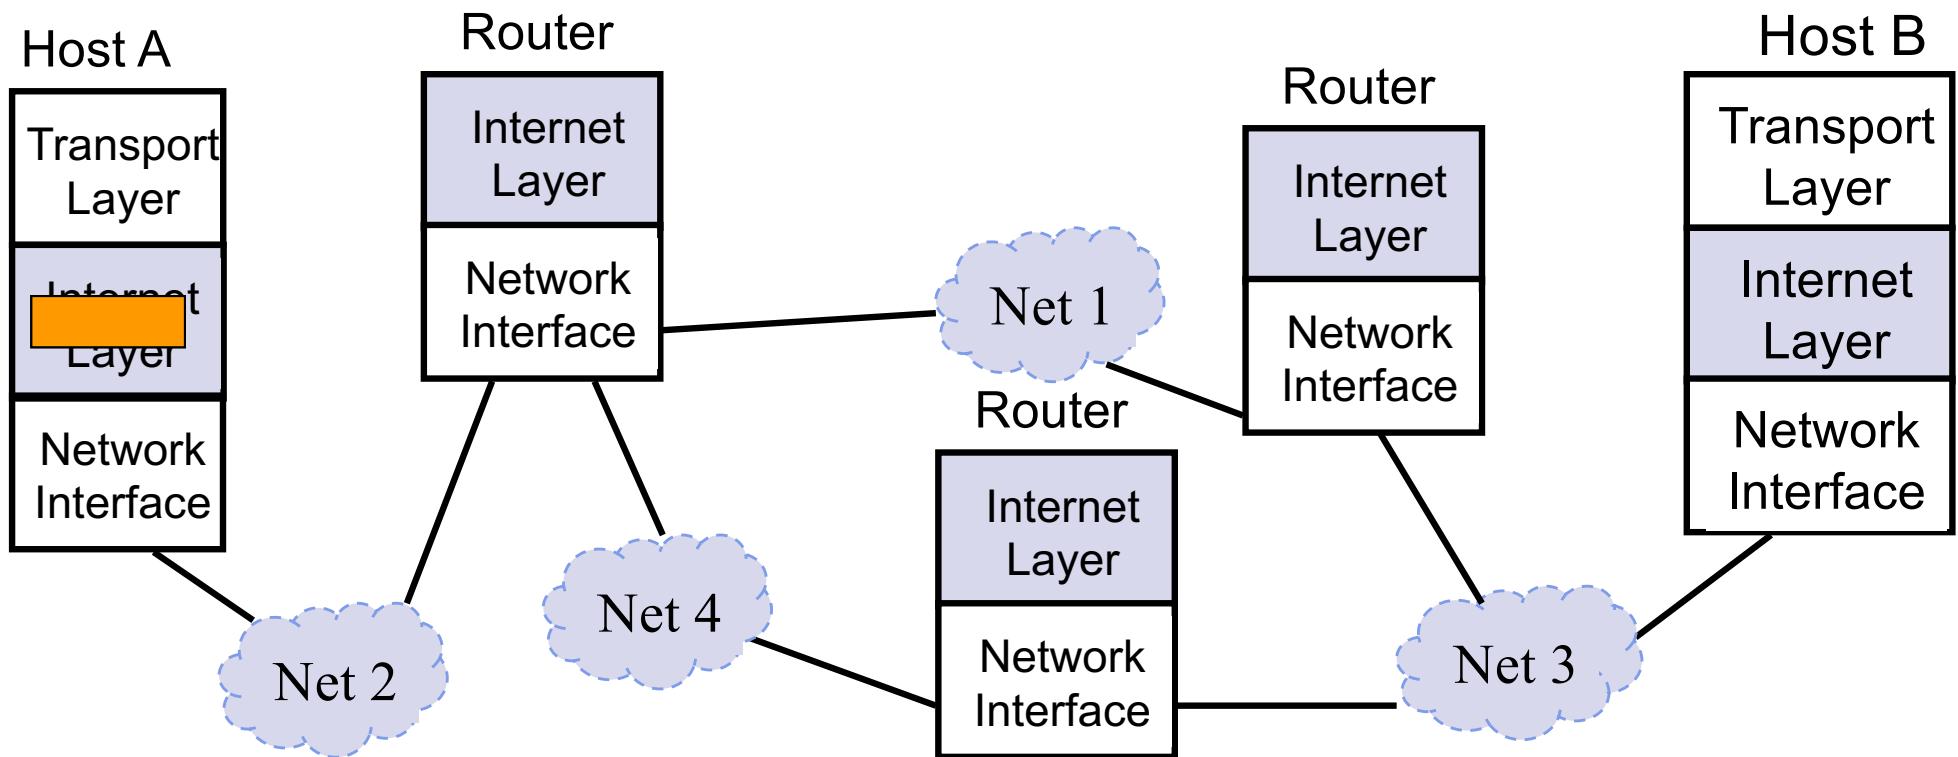


- To provide *distributed applications*
 - Any application designed to operate based on Internet communication services immediately operates across the entire Internet
 - Rapid deployment of new applications
 - Email, WWW, Peer-to-peer
 - Applications independent of network technology
 - New networks can be introduced below
 - Old network technologies can be retired

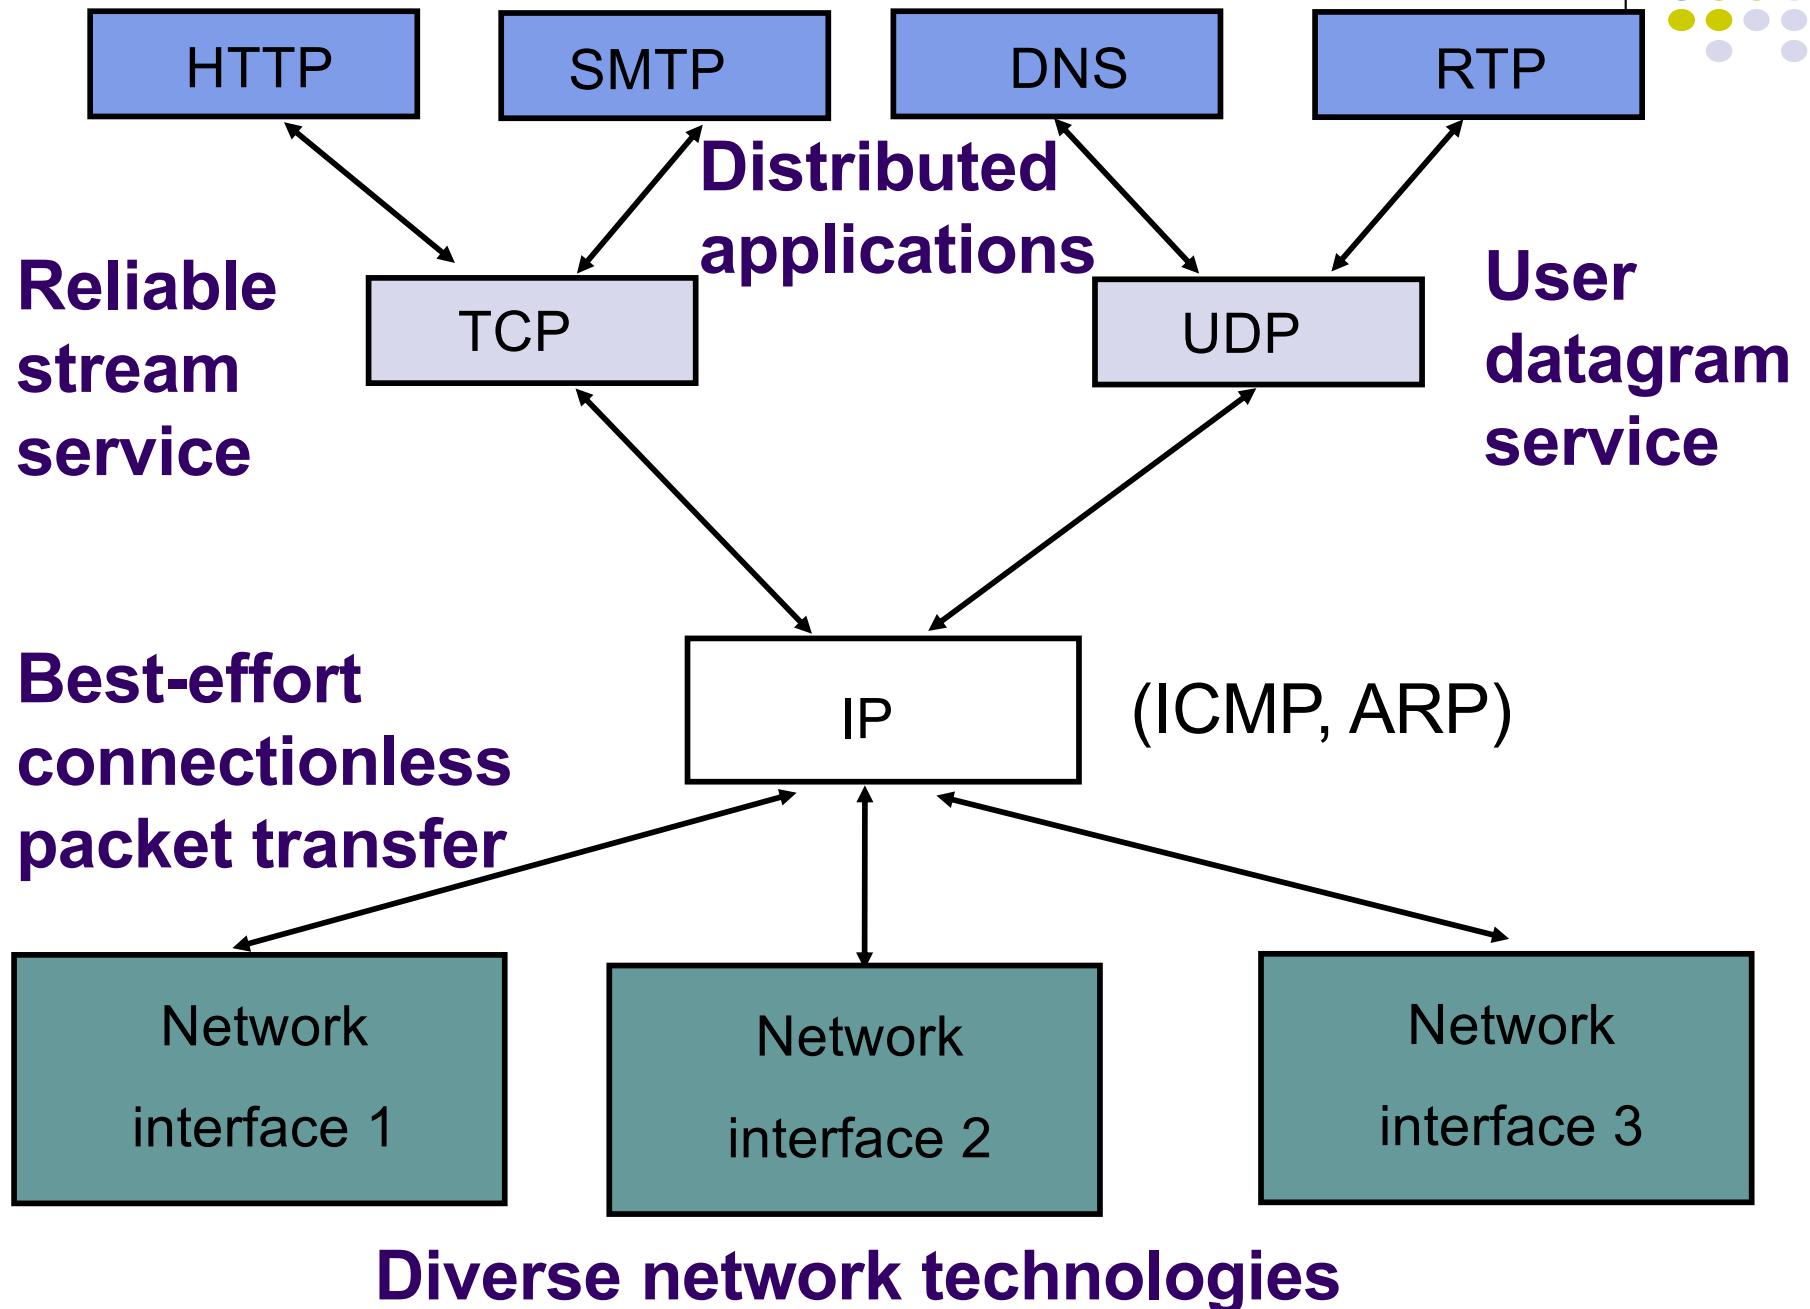
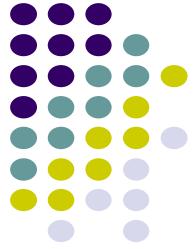


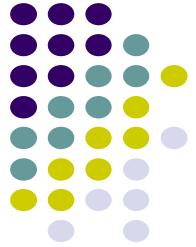
Internet Protocol Approach

- IP packets transfer information across Internet
Host A IP → router → router... → router → Host B IP
- IP layer in each router determines next hop (router)
- Network interfaces transfer IP packets across networks



TCP/IP Protocol Suite





Internet Names & Addresses

Internet Names

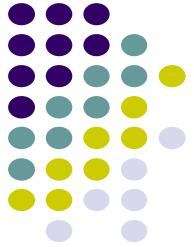
- Each host has a unique name
 - Independent of physical location
 - Facilitate memorization by humans
 - Domain Name
 - Organization under single administrative unit
- Host Name
 - Name given to host computer
- User Name
 - Name assigned to user

leongarcia@comm.utoronto.ca

Internet Addresses

- Each host has globally unique *logical* 32 bit IP address
- Separate address for each physical connection to a network
- Routing decision is done based on destination IP address
- **IP address has two parts:**
 - *netid* and *hostid*
 - *netid* unique
 - *netid* facilitates routing
- Dotted Decimal Notation:
int1.int2.int3.int4
(int_j = jth octet)
128.100.10.13

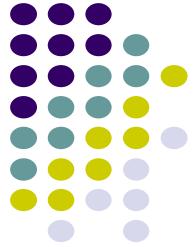
DNS resolves IP name to IP address



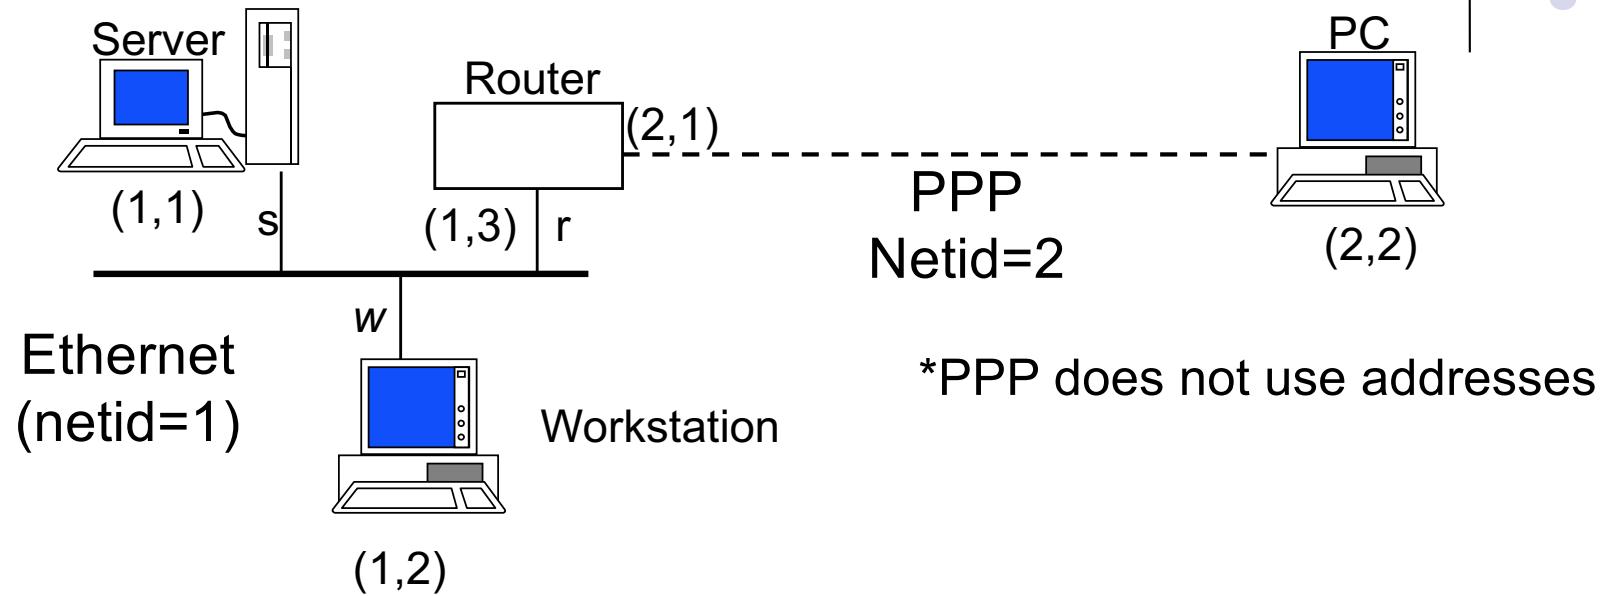
Physical Addresses

- LANs (and other networks) assign **physical addresses** to the physical attachment to the network
- The network uses its own address to transfer packets or frames to the appropriate destination
- IP address needs to be resolved to physical address at each IP network interface
- Example: Ethernet uses 48-bit addresses
 - Each Ethernet network interface card (NIC) has globally unique Medium Access Control (MAC) or physical address
 - First 24 bits identify NIC manufacturer; second 24 bits are serial number
 - 00:90:27:96:68:07 12 hex numbers

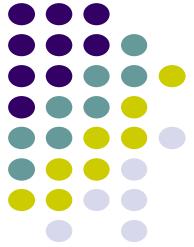
 Intel



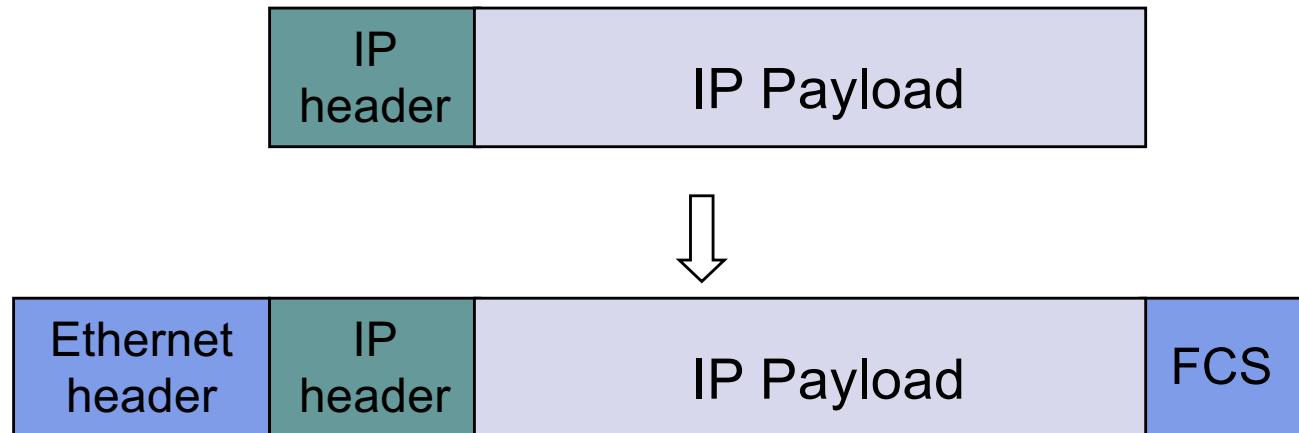
Example internet



	netid	hostid	Physical address
server	1	1	s
workstation	1	2	w
router	1	3	r
router	2	1	-
PC	2	2	-

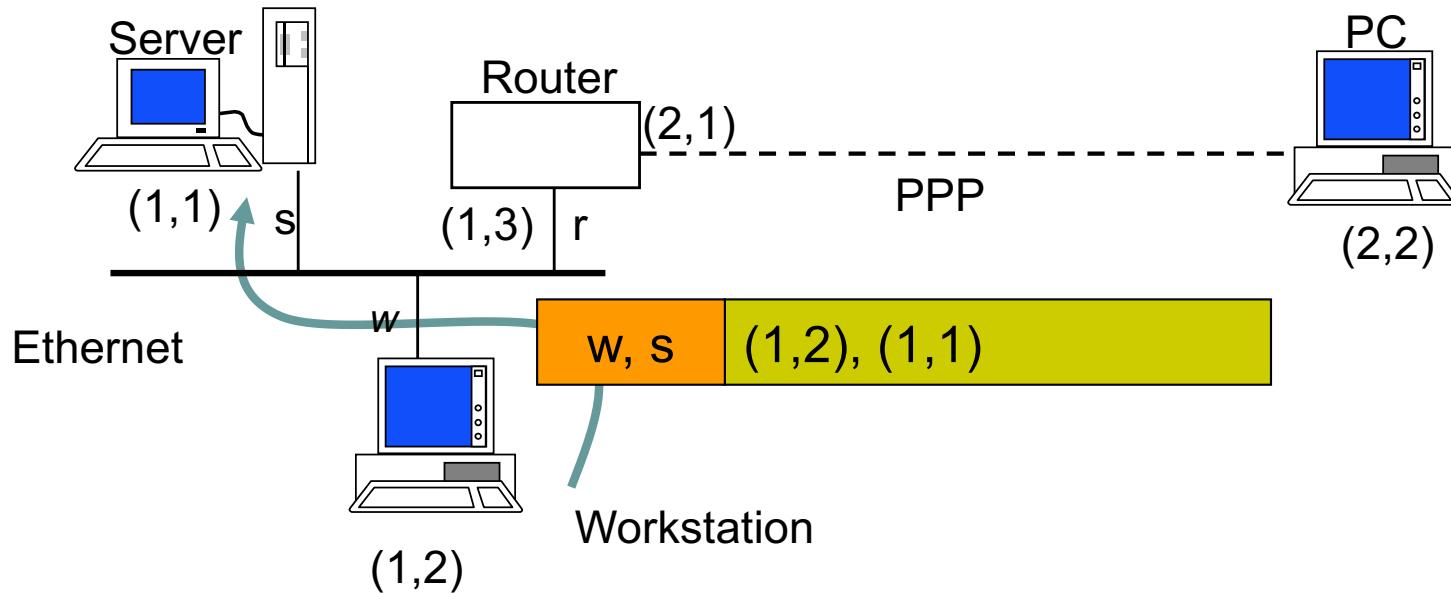
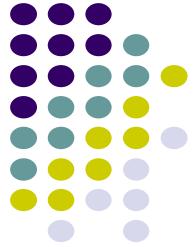


Encapsulation



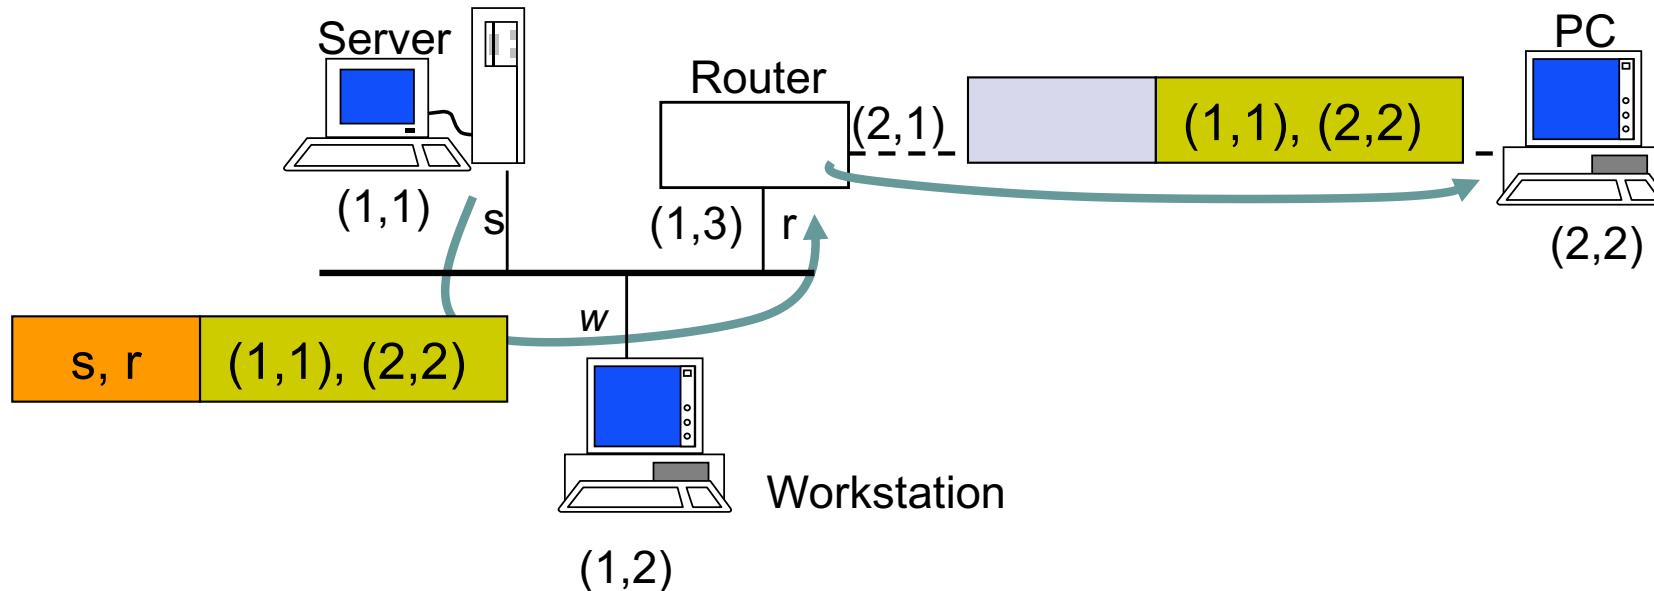
- Ethernet header contains:
 - source and destination physical addresses
 - network protocol type (e.g. IP)

IP packet from workstation to server



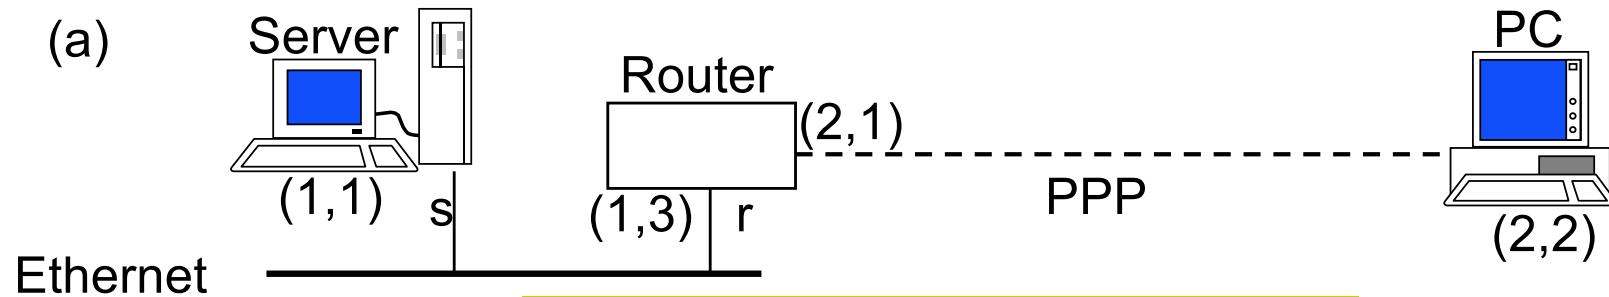
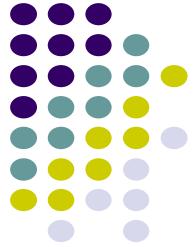
1. IP packet has (1,2) IP address for source and (1,1) IP address for destination
2. IP table at workstation indicates (1,1) connected to same network, so IP packet is encapsulated in Ethernet frame with addresses w and s
3. Ethernet frame is broadcast by workstation NIC and captured by server NIC
4. NIC examines protocol type field and then delivers packet to its IP layer

IP packet from server to PC



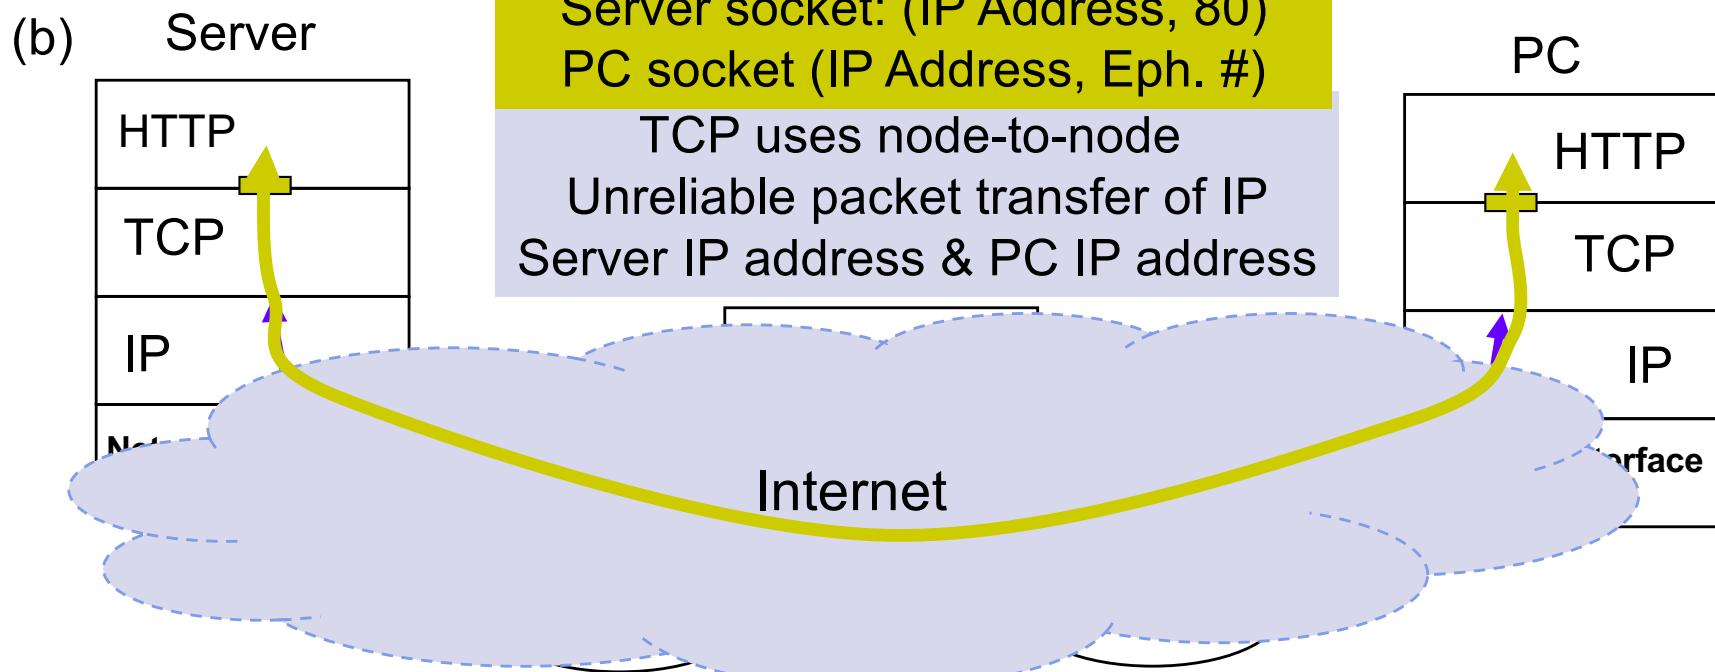
1. IP packet has (1,1) and (2,2) as IP source and destination addresses
2. IP table at server indicates packet should be sent to router, so IP packet is encapsulated in Ethernet frame with addresses s and r
3. Ethernet frame is broadcast by server NIC and captured by router NIC
4. NIC examines protocol type field and then delivers packet to its IP layer
5. IP layer examines IP packet destination address and determines IP packet should be routed to (2,2)
6. Router's table indicates (2,2) is directly connected via PPP link
7. IP packet is encapsulated in PPP frame and delivered to PC
8. PPP at PC examines protocol type field and delivers packet to PC IP layer

How the layers work together



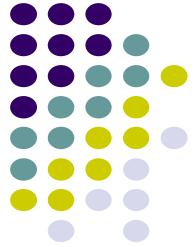
Ethernet

HTTP uses process-to-process
Reliable byte stream transfer of
TCP connection:
Server socket: (IP Address, 80)
PC socket (IP Address, Eph. #)



Internet

Encapsulation



TCP Header contains source & destination port numbers

IP Header contains source and destination IP addresses;
transport protocol type

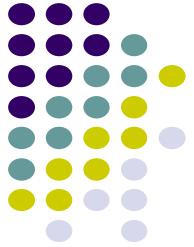
Ethernet Header contains source & destination MAC addresses;
network protocol type

HTTP Request

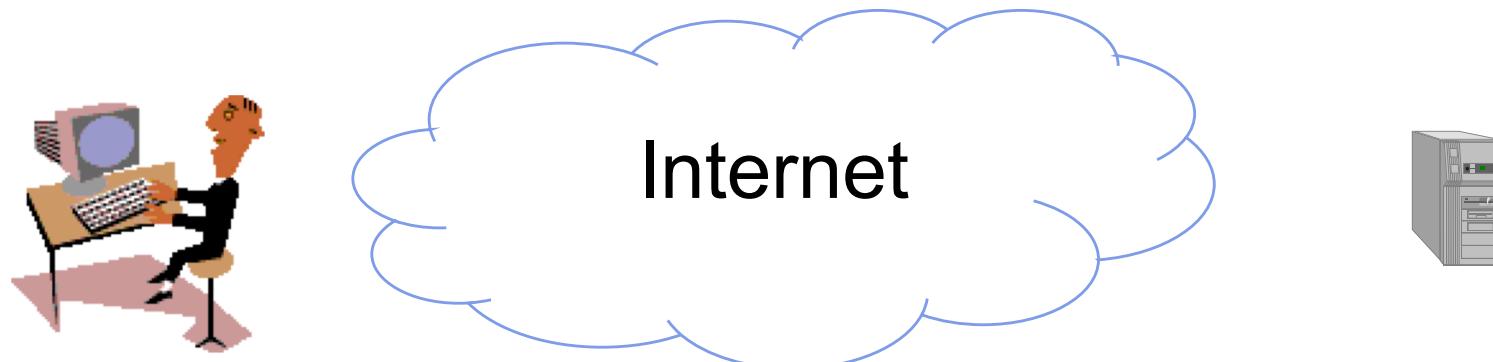
TCP header HTTP Request

IP header TCP header HTTP Request

Ethernet header IP header TCP header HTTP Request FCS



How the layers work together: Network Analyzer Example



- User clicks on <http://www.nytimes.com/>
- *Ethereal* network analyzer captures all frames observed by its Ethernet NIC
- Sequence of frames and contents of frame can be examined in detail down to individual bytes

Ether

Top Pane shows frame/packet sequence

Middle Pane shows encapsulation for a given frame

Bottom Pane shows hex & text

No.	Time	Source	Destination	Protocol	Info
1	0.000000	128.100.11.13	128.100.100.128	DNS	Standard query A www.nytimes.com
2	0.129976	128.100.100.128	128.100.11.13	DNS	Standard query response A 64.15.247.200 A 64.15.247.24
3	0.131524	128.100.11.13	64.15.247.200	TCP	1127 > http [SYN] Seq=1396200325 Ack=3638689752 Win=16384 Len=0
4	0.168286	64.15.247.200	128.100.11.13	TCP	http > 1127 [SYN] Seq=1396200325 Ack=3638689753 Win=16384 Len=0
5	0.168320	128.100.11.13	64.15.247.200	TCP	1127 > http [ACK] Seq=3638689753 Ack=1396200326 Win=17920
6	0.168688	128.100.11.13	64.15.247.200	HTTP	GET / HTTP/1.1
7	0.205439	64.15.247.200	128.100.11.13	TCP	http > 1127 [ACK] Seq=1396200326 Ack=3638690402 Win=32768
8	0.236676	64.15.247.200	128.100.11.13	HTTP	HTTP/1.1 200 OK

Frame 1 (75 bytes on wire, 75 bytes captured)
Ethernet II, Src: 00:90:27:96:b8:07, Dst: 00:e0:52:ea:b5:00
Internet Protocol, Src Addr: 128.100.11.13 (128.100.11.13), Dst Addr: 128.100.100.128 (128.100.100.128)
User Datagram Protocol, Src Port: 1126 (1126), Dst Port: domain (53)
Domain Name System (query)

0000	00 e0 52 ea b5 00 00 90	27 96 b8 07 08 00 45 00	..R.....'.....E.
0010	00 3d 54 41 00 00 80 11	76 19 80 64 0b 0d 80 64	.=TA... v..d...d
0020	64 80 04 66 00 35 00 29	49 83 00 a5 01 00 00 01	d..f.5.) I.....
0030	00 00 00 00 00 03 77	77 77 07 6e 79 74 69 6dw ww.nytim
0040	65 73 03 63 6f 6d 00 00	01 00 01	es.com.. ...

Filter:

Top pane: frame sequence



The screenshot shows the Wireshark interface with three main components highlighted by red boxes:

- DNS Query**: Points to the first two frames (DNS requests) in the sequence.
- TCP Connection Setup**: Points to the SYN and ACK segments of the TCP handshake (frames 3 and 4).
- HTTP Request & Response**: Points to the HTTP GET request and its corresponding response (frames 6 and 7).

The packet list pane displays the following sequence of frames:

No.	Time	Source	Destination	Protocol	Information
1	0.000000	128.100.11.13	128.100.100.128	DNS	Standard query A www.nytimes.com
2	0.129976	128.100.100.128	128.100.11.13	DNS	Standard query response
3	0.131124	128.100.11.13	64.15.247.200	TCP	1127 > http [SYN] Seq=535
4	0.168286	64.15.247.200	128.100.11.13	TCP	http > 1127 [SYN, ACK] Seq=1 Ack=535 Win=17
5	0.168320	128.100.11.13	64.15.247.200	TCP	1127 > http [ACK] Seq=36386 Ack=1396200326 Win=17
6	0.168688	128.100.11.13	64.15.247.200	HTTP	GET / HTTP/1.1
7	0.205439	64.15.247.200	128.100.11.13	TCP	http > 1127 [ACK] Seq=1396200326 Ack=3638690402 Win=32
8	0.236676	64.15.247.200	128.100.11.13	HTTP	HTTP/1.1 200 OK

The details pane shows the analysis for Frame 1:

- Frame 1 (75 bytes on wire, 75 bytes captured)
- Ethernet II, Src: 00:90:27:96:b8:07, Dst: 00:e0:52:ea:b5:00
- Internet Protocol, Src Addr: 128.100.11.13 (128.100.11.13), Dst Addr: 128.100.100.128 (128.100.100.128)
- User Datagram Protocol, Src Port: 1126 (1126), Dst Port: domain (53)
- Domain Name System (query)

The bytes pane shows the raw hex and ASCII data for the DNS query frame (Frame 1). The ASCII output includes the query for "www.nytimes.com".

Filter: Reset File: nytimespackets

Middle pane: Encapsulation



nytimespackets - Ethereal

Ethernet Frame

Protocol Type

Ethernet Destination and Source Addresses

The screenshot shows the middle pane of the Ethereal network traffic analyzer. The title bar reads "nytimespackets - Ethereal". The main window displays a list of captured network packets. A red box highlights the first packet, which is an Ethernet II frame. The red box has three callout arrows pointing to specific fields: "Protocol Type" points to the "Type: IP (0x0800)" field; "Ethernet Destination and Source Addresses" points to the source and destination MAC addresses; and "Ethernet Frame" points to the overall structure of the frame in the list view.

No. Time Source Destination Protocol Length Info

6 0.168688 128.100.11.13 64.15.247.200 HTTP 54 GET /index.html

Ethernet II, Src: 00:90:27:96:b8:07, Dst: 00:e0:52:ea:b5:00
Destination: 00:e0:52:ea:b5:00 (Foundry_ea:b5:00)
Source: 00:90:27:96:b8:07 (Intel_96:b8:07)
Type: IP (0x0800)

Internet Protocol Version 4, Src Addr . 128.100.11.13 (128.100.11.13), Dst Addr . 64.15.247.200 (64.15.247.200)
Version: 4 Header Len: 20 bytes
Differentiated Services Field: 0x00 (DSCP 0x00: default; ECN 0x00)
Total Length: 54
Identification: 0
Flags: Frag 0
Time to live: 128
Protocol: TCP (0x06)
Header checksum: 0xe0b8 (correct)
Source: 128.100.11.13 (128.100.11.13)
Destination: 64.15.247.200 (64.15.247.200)

Transmission Control Protocol, Src Port: 1127 (1127), Dst Port: http (80), seq: 3638689753, ACK: 139620032

Hypertext Transfer Protocol

0000 00 e0 52 ea b5 00 00 90 27 96 b8 07 08 00 45 00 .R..... E.
0010 02 b1 54 45 40 00 80 06 e0 b8 80 64 0b 0d 40 0f ..TE@... d..@.
0020 f7 c8 04 67 00 50 d8 e1 ff d9 53 38 53 86 50 18 ...g.P.. S85.P.
0030 43 a4 87 81 00 00 47 45 54 20 2f 20 48 54 54 50 C....GE T / HTTP
0040 2f 31 2e 31 0d 0a 41 63 63 65 70 74 3a 20 69 6d /1.i..Ac cept: im

Filter: File: nytimespackets

Middleman Network Encapsulation



And a lot of
other stuff!

IP Packet

IP Source and
Destination
Addresses

Protocol Type

The screenshot shows a Wireshark capture window titled "nytimespackets". The interface list pane shows a single frame (Frame 6) with the following details:

- Frame 6 (703 bytes on wire, 703 bytes captured)
- Ethernet II, Src: Intel_96:b8:07 (00:09:27:96:b8:07), Dst: Foundry_ea:b5:00 (00:e0:52:ea:b5:00)
- Internet Protocol, Src Addr: 128.100.11.13 (128.100.11.13), Dst Addr: 64.15.247.200 (64.15.247.200)
- Differentiated Services Field: 0x00 (DSCP 0x00: Default; ECN: 0x00)
- Total Length: 689
- Identification: 0x5445
- Flags: 0x04
- Fragment offset: 0
- Time to live: 128
- Protocol: TCP (0x06)
- Header checksum: 0xe0b8 (correct)
- Source: 128.100.11.13 (128.100.11.13)
- Destination: 64.15.247.200 (64.15.247.200)

The bytes pane shows the raw hex and ASCII data of the packet.

Middle pane: Encapsulation



nytimespackets - Ethereal

No. Time Source Destination Protocol Info

6 0.168688 128.100.11.13 64.15.247.200 HTTP GET / HTT

Frame 6 (703 bytes on wire, 703 bytes captured)
Ethernet II, Src: 00:90:27:96:b8:07, Dst: 00:e0:52:ea:b5:00
Internet Protocol Version 4, Src Address: 128.100.11.13 (128.100.11.13), Dst Address: 64.15.247.200 (64.15.247.200)
Transmission Control Protocol, src Port: 1127 (1127), dst Port: http (80), seq: 3638689753, Ack: 139620032
Source port: 1127 (1127)
Destination port: http (80)
Sequence number: 3638689753
Next sequence number: 3638690402
Acknowledgement number: 1396200326
Header length: 20 bytes
Flags: 0x0018 (PSH, ACK)
Window size: 17316
Checksum: 0x2721 (correct)
HTTP/1.1
GET / HTTP/1.1\r\nAccept: image/gif, image/x-xpixmap, image/jpeg, application/pjpeg, application/vnd.ms-powerpoint, application/
Accept-Language: en-us\r\nAccept-Encoding: gzip, deflate\r\nUser-Agent: Mozilla/4.0 (compatible; MSIE 6.0; Windows NT 5.0)\r\nHost: www.nytimes.com\r\nConnection: Keep-Alive\r\nCookie: RMID=80e7478f5a393db9fc19f2c4; NYT-S=1002xv091grjagxb2AZ90xq41qdE, 174380X001E287C92Q2Qme5mO8R6\r\n\r\n

0000 00 e0 52 ea b5 00 00 90 27 96 b8 07 00
0010 02 b1 54 45 40 00 80 06 e0 b8 80 64 0b 0
0020 f7 c8 04 67 00 50 d8 e1 ff d9 53 38 53 8
0030 43 a4 87 81 00 00 47 45 54 20 2f 20 48 5
0040 2f 31 2e 31 0d 0a 41 63 63 65 70 74 3a 2

E.
G.
P.
TP
im

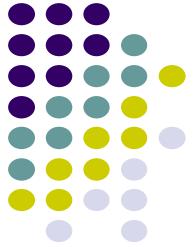
Filter: nytimespackets

TCP Segment

Source and Destination Port Numbers

GET

HTTP Request

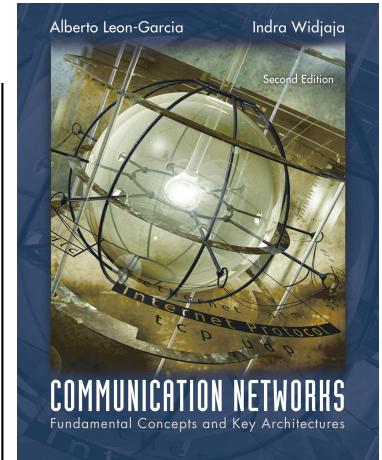


Summary

- Encapsulation is key to layering
- IP provides for transfer of packets across diverse networks
- TCP and UDP provide universal communications services across the Internet
- Distributed applications that use TCP and UDP can operate over the entire Internet
- Internet names, IP addresses, port numbers, sockets, connections, physical addresses

Chapter 2

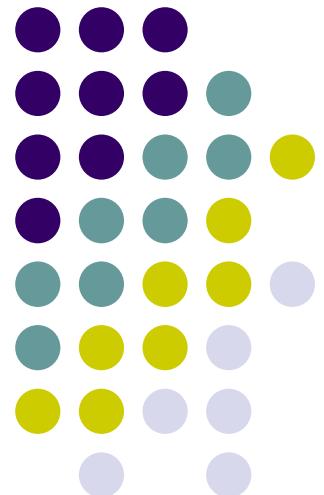
Applications and Layered Architectures

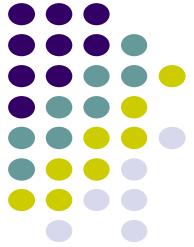


Application Layer Protocols

HTTP

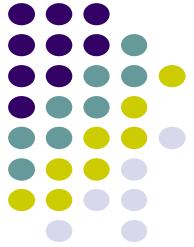
RTP & SIP





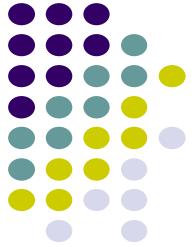
HTTP Protocol

- HTTP servers use well-known port 80
- Client request / Server reply
- Stateless: server does not keep any information about client
- HTTP 1.0 new TCP connection per request/reply (non-persistent)
- HTTP 1.1 persistent operation is default



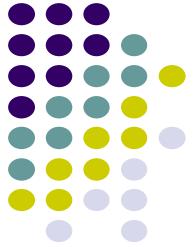
HTTP Message Formats

- HTTP messages written in ASCII text
- Request Message Format
 - *Request Line (Each line ends with carriage return)*
 - Method URL HTTP-Version \r\n
 - Method specifies action to apply to object
 - URL specifies object
 - *Header Lines (Ea. line ends with carriage return)*
 - Attribute Name: Attribute Value
 - E.g. type of client, content, identity of requester, ...
 - Last header line has extra carriage return)
 - *Entity Body (Content)*
 - Additional information to server



HTTP Request Methods

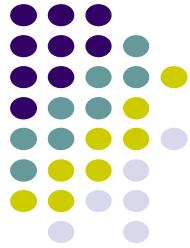
Request method	Meaning
GET	Retrieve information (object) identified by the URL.
HEAD	Retrieve meta-information about the object, but do not transfer the object; Can be used to find out if a document has changed.
POST	Send information to a URL (using the entity body) and retrieve result; used when a user fills out a form in a browser.
PUT	Store information in location named by URL
DELETE	Remove object identified by URL
TRACE	Trace HTTP forwarding through proxies, tunnels, etc.
OPTIONS	Used to determine the capabilities of the server, or characteristics of a named resource.



Universal Resource Locator

- Absolute URL
 - scheme://hostname[:port]/path
 - <http://www.nytimes.com/>
- Relative URL
 - /path
 - /

HTTP Request Message



© nytimespackets - Ethereal

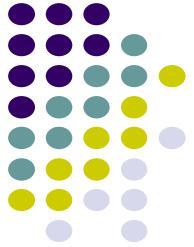
File Edit Capture Display Tools Help

No.	Time	Source	Destination	Protocol	Info
1	0.000000	128.100.11.13	128.100.100.128	DNS	Standard query A www.nytimes.com
2	0.129976	128.100.100.128	128.100.11.13	DNS	Standard query response A 64.15.247.200 A 64
3	0.131524	128.100.11.13	64.15.247.200	TCP	1127 > http [SYN] Seq=3638689752 Ack=0 Win=1
4	0.168286	64.15.247.200	128.100.11.13	TCP	http > 1127 [SYN, ACK] seq=1396200325 Ack=36
5	0.168320	128.100.11.13	64.15.247.200	TCP	1127 > http [ACK] Seq=3638689753 Ack=1396200
6	0.168688	128.100.11.13	64.15.247.200	HTTP	GET / HTTP/1.1
7	0.205439	64.15.247.200	128.100.11.13	TCP	http > 1127 [ACK] Seq=1396200326 Ack=3638690
8	0.236676	64.15.247.200	128.100.11.13	HTTP	HTTP/1.1 200 OK

Frame 6 (703 bytes on wire, 703 bytes captured)
Ethernet II, Src: 00:90:27:96:b8:07, Dst: 00:e0:52:ea:b5:00
Internet Protocol, Src Addr: 128.100.11.13 (128.100.11.13), Dst Addr: 64.15.247.200 (64.15.247.200)
Transmission Control Protocol, Src Port: 1127 (1127), Dst Port: http (80), Seq: 3638689753, Ack: 139620032
HyperText Transfer Protocol
GET / HTTP/1.1\r\nAccept: image/gif, image/x-bitmap, image/jpeg, image/pjpeg, application/vnd.ms-powerpoint, application/
Accept-Language: en-us\r\nAccept-Encoding: gzip, deflate\r\nUser-Agent: Mozilla/4.0 (compatible; MSIE 6.0; Windows NT 5.0)\r\nHost: www.nytimes.com\r\nConnection: Keep-Alive\r\nCookie: RMID=80e7478f5a393db9fc19f2c4; NYT-S=1002xv091grjagxb2AZ9oxq4lqdEe/irdKSU3XUnLr287eqe2QOMe5m08R6\r\n\r\n

Hex	Dec	ASCII
0000	00 e0 52 ea b5 00 00 90	..R..... E.
0010	02 b1 54 45 40 00 80 06	..TE@... d..@.
0020	f7 c8 04 67 00 50 d8 e1	...g.P.. .S85.P.
0030	54 20 2f 20 48 54 54 50	C....GE T / HTTP
0040	63 65 70 74 3a 20 69 6d	/1.1..Ac cept: im
0050	20 69 6d 61 67 65 2f 78	age/gif, image/x

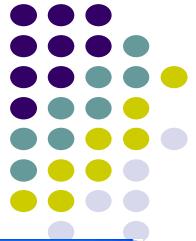
Filter: Reset File: nytimespackets



HTTP Response Message

- Response Message Format
 - *Status Line*
 - *HTTP-Version Status-Code Message*
 - Status Code: 3-digit code indicating result
 - E.g. HTTP/1.0 200 OK
 - *Headers Section*
 - Information about object transferred to client
 - E.g. server type, content length, content type, ...
 - *Content*
 - Object (document)

HTTP Response Message



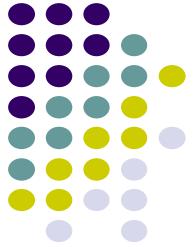
nytimespackets - Ethereal

File Edit Capture Display Tools Help

No.	Time	Source	Destination	Protocol	Info
1	0.000000	128.100.11.13	128.100.100.128	DNS	Standard query A www.nytimes.com
2	0.129976	128.100.100.128	128.100.11.13	DNS	Standard query response A 64.15.247.200 A 64
3	0.131524	128.100.11.13	64.15.247.200	TCP	1127 > http [SYN] Seq=3638689752 Ack=0 Win=1
4	0.168286	64.15.247.200	128.100.11.13	TCP	http > 1127 [SYN, ACK] Seq=1396200325 Ack=36
5	0.168320	128.100.11.13	64.15.247.200	TCP	1127 > http [ACK] Seq=3638689753 Ack=1396200
6	0.168688	128.100.11.13	64.15.247.200	HTTP	GET / HTTP/1.1
7	0.205439	64.15.247.200	128.100.11.13	TCP	http > 1127 [ACK] Seq=1396200326 Ack=3638690
8	0.236676	64.15.247.200	128.100.11.13	HTTP	HTTP/1.1 200 OK

Frame 8 (284 bytes on wire, 284 bytes captured)
Ethernet II, Src: 00:e0:52:ea:b5:00, Dst: 00:90:27:96:b8:07
Internet Protocol, Src Addr: 64.15.247.200 (64.15.247.200), Dst Addr: 128.100.11.13 (128.100.11.13)
Transmission Control Protocol, src Port: http (80), Dst Port: 1127 (1127), seq: 1396200326, Ack: 363869040
Hypertext Transfer Protocol
HTTP/1.1 200 OK\r\nServer: Netscape-Enterprise/4.1\r\nDate: Sat, 02 Nov 2002 02:53:48 GMT\r\nSet-cookie: spopunder=1; path=/; domain=.nytimes.com\r\nCache-control: no-cache\r\nPragma: no-cache\r\nContent-type: text/html\r\nConnection: close\r\n\r\n

0000 00 90 27 96 b8 07 00 e0 52 ea b5 00 08 00 45 00 R . . . E.
0010 01 0e b3 93 40 00 ed 06 16 0d 40 0f f7 c8 80 64 . . . @ . . . @ . . d
0020 0b 0d 00 50 04 67 53 38 53 86 d8 e2 02 62 50 18 . . . P . gS8 S . . . bP.
0030 7f ff 8a f6 00 00 48 54 54 50 2f 31 2e 31 20 32 I . . . HT TP/1.1 2
0040 30 30 20 4f 4b 0d 0a 53 65 72 76 65 72 3a 20 4e 00 OK..S erver: N
0050 65 74 73 63 61 70 65 2d 45 6e 74 65 72 70 72 69 etscape- Enterpri
..... 77 ee 7f 74 7a 71 dd .. 44 e1 74 ee 77 70 e1 .. 44 1 .. 77 ..
Filter: Reset Apply File: nytimespackets



Cookies and Web Sessions

- Cookies are data exchanged by clients & servers as header lines
- Since HTTP stateless, cookies can provide context for HTTP interaction
- Set *cookie* header line in reply message from server + unique ID number for client
- If client accepts cookie, cookie added to client's cookie file (must include expiration date)
- Henceforth client requests include ID
- Server site can track client interactions, store these in a separate database, and access database to prepare appropriate responses

Cookie Header Line; ID is 24 hexadecimal numeral



nytimespackets - Ethereal

File Edit Capture Display Tools Help

No.	Time	Source	Destination	Protocol	Info
1	0.000000	128.100.11.13	128.100.100.128	DNS	Standard query A www.nytimes.com
2	0.129976	128.100.100.128	128.100.11.13	DNS	Standard query response A 64.15.247.200 A 64
3	0.131524	128.100.11.13	64.15.247.200	TCP	1127 > http [SYN] Seq=3638689752 Ack=0 Win=1
4	0.168286	64.15.247.200	128.100.11.13	TCP	http > 1127 [SYN, ACK] seq=1396200325 Ack=36
5	0.168320	128.100.11.13	64.15.247.200	TCP	1127 > http [ACK] Seq=3638689753 Ack=1396200
6	0.168688	128.100.11.13	64.15.247.200	HTTP	GET / HTTP/1.1
7	0.205439	64.15.247.200	128.100.11.13	TCP	http > 1127 [ACK] Seq=1396200326 Ack=3638690
8	0.236676	64.15.247.200	128.100.11.13	HTTP	HTTP/1.1 200 OK

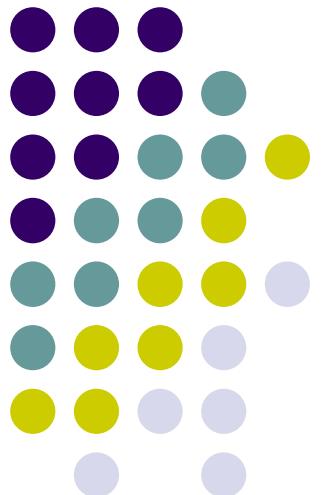
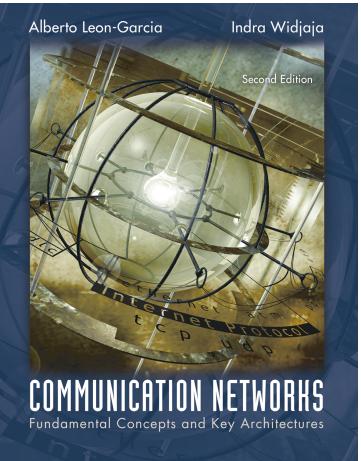
Frame 6 (703 bytes on wire, 703 bytes captured)
Ethernet II, Src: 00:90:27:96:b8:07, Dst: 00:e0:52:ea:b5:00
Internet Protocol, Src Addr: 128.100.11.13 (128.100.11.13), Dst Addr: 64.15.247.200 (64.15.247.200)
Transmission Control Protocol, Src Port: 1127 (1127), Dst Port: http (80), Seq: 3638689753, Ack: 139620032
Hypertext Transfer Protocol
GET / HTTP/1.1\r\nAccept: image/gif, image/x-bitmap, image/jpeg, image/pjpeg, application/vnd.ms-powerpoint, application/
Accept-Language: en-us\r\nAccept-Encoding: gzip, deflate\r\nUser-Agent: Mozilla/4.0 (compatible; MSIE 6.0; Windows NT 5.0)\r\nHost: www.nytimes.com\r
Connection: Keep-Alive\r\nCookie: RMID=80e7478f5a393db9fc19f2c4; NYT-S=1002xv091grjagxb2AZ90xq4lqdEe/irdKSU3XUnLr287eqe2QOMe5m08R6\r\n\r

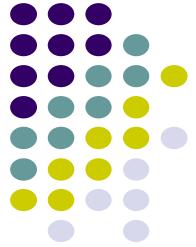
0000 00 e0 52 ea b5 00 00 90 27 96 b8 07 08 00 45 00 ..R..... E.
0010 02 b1 54 45 40 00 80 06 e0 b8 80 64 0b 0d 40 0f ..TE@... d..@.
0020 f7 c8 04 67 00 50 d8 e1 ff d9 53 38 53 86 50 18 ...g.P.. .S85.P.
0030 43 a4 87 81 00 00 47 45 54 20 2f 20 48 54 54 50 C....GE T / HTTP
0040 2f 31 2e 31 0d 0a 41 63 63 65 70 74 3a 20 69 6d /1.1..Ac cept: im
0050 61 67 65 2f 67 69 66 2c 20 69 6d 61 67 65 2f 78 age/gif, image/x
^~^ ^~^ ^~^ ^~^ ^~^ ^~^ ^~^ ^~^ ^~^ ^~^ ^~^ ^~^ ^~^ ^~^ ^~^ ^~^ ^~^ ^~^

Filter: Reset File: nytimespackets

Voice Over IP

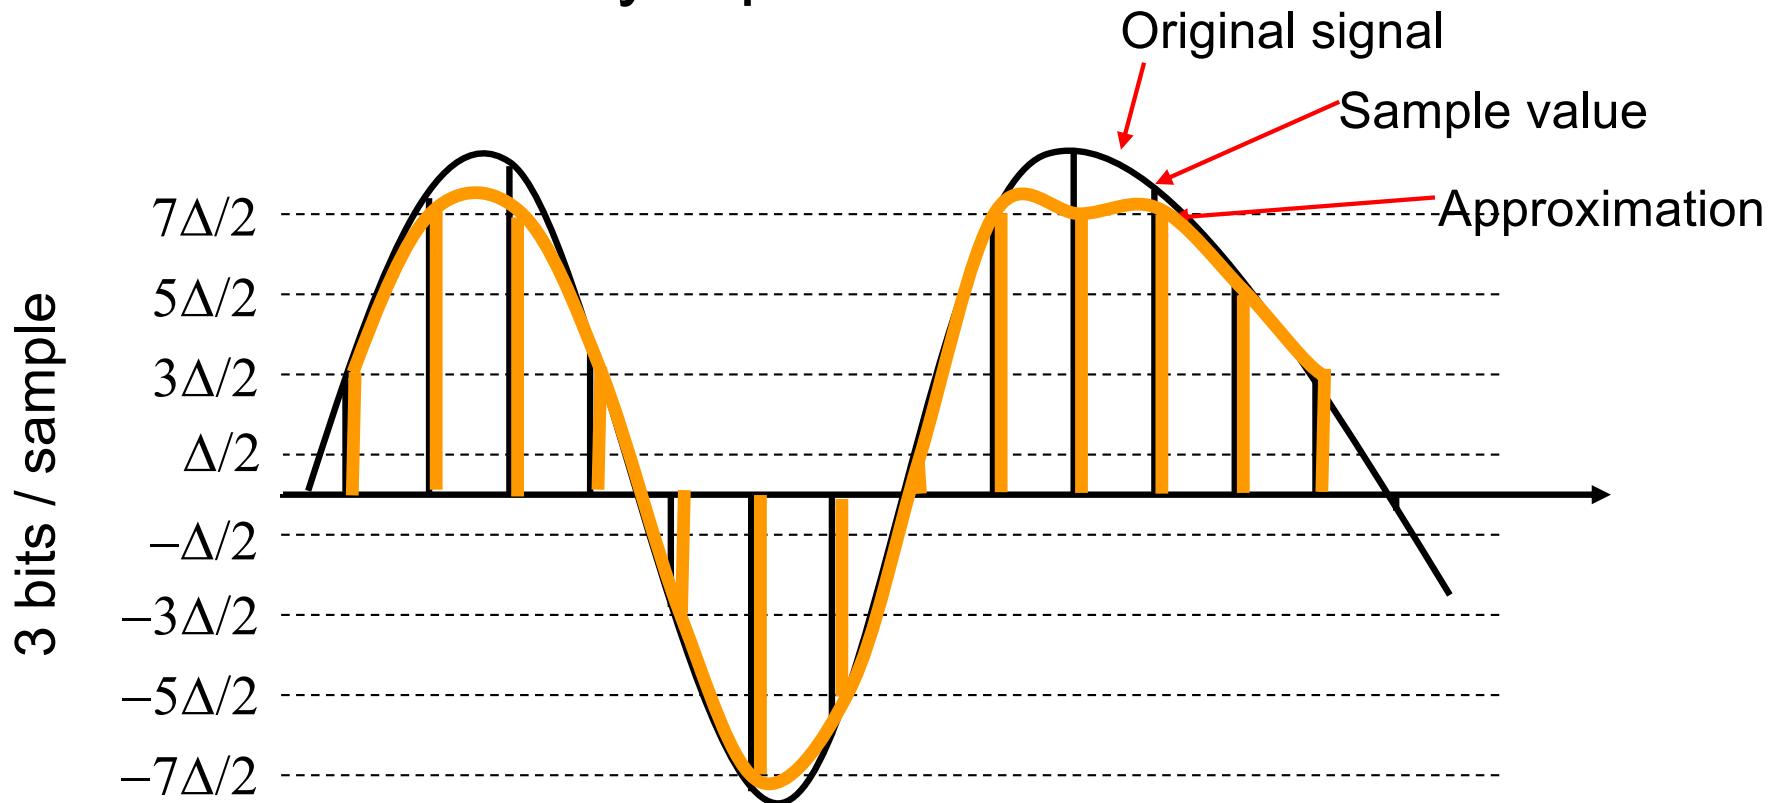
Digitization
Timing Recovery
RTP over UDP



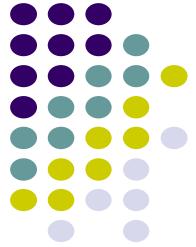


Digitization of Analog Signal

- Sample analog signal in time and amplitude
- Find closest binary representation



$$\begin{aligned} R_s &= \text{Bit rate} = \# \text{ bits/sample} \times \# \text{ samples/second} \\ &= n \text{ bits/sample} 2 W \text{ samples/sec} \end{aligned}$$



Example: Voice & Audio

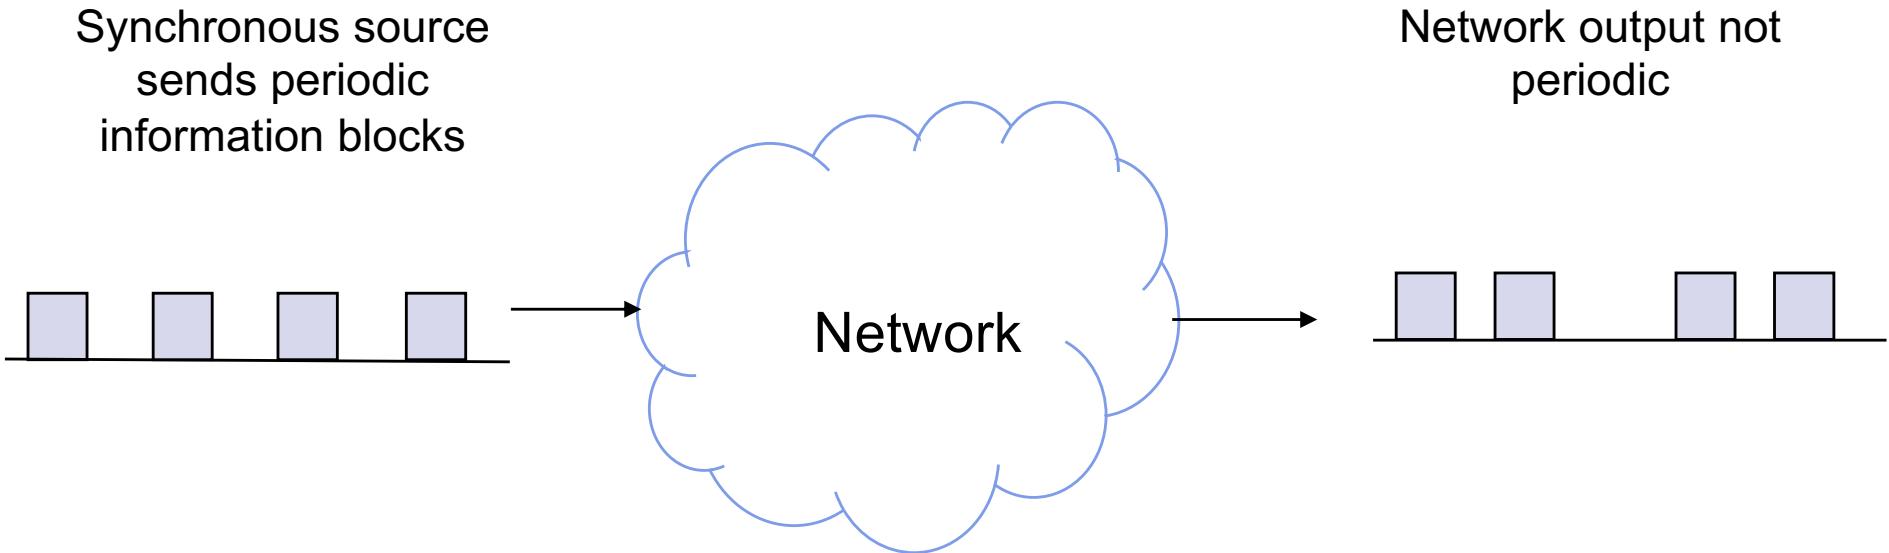
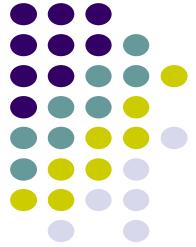
Telephone voice

- $W_s = 4 \text{ kHz} \rightarrow 8000 \text{ samples/sec}$
- 8 bits/sample
- $R_s = 8 \times 8000 = 64 \text{ kbps}$
- Cellular phones use more powerful compression algorithms: 8-12 kbps

CD Audio

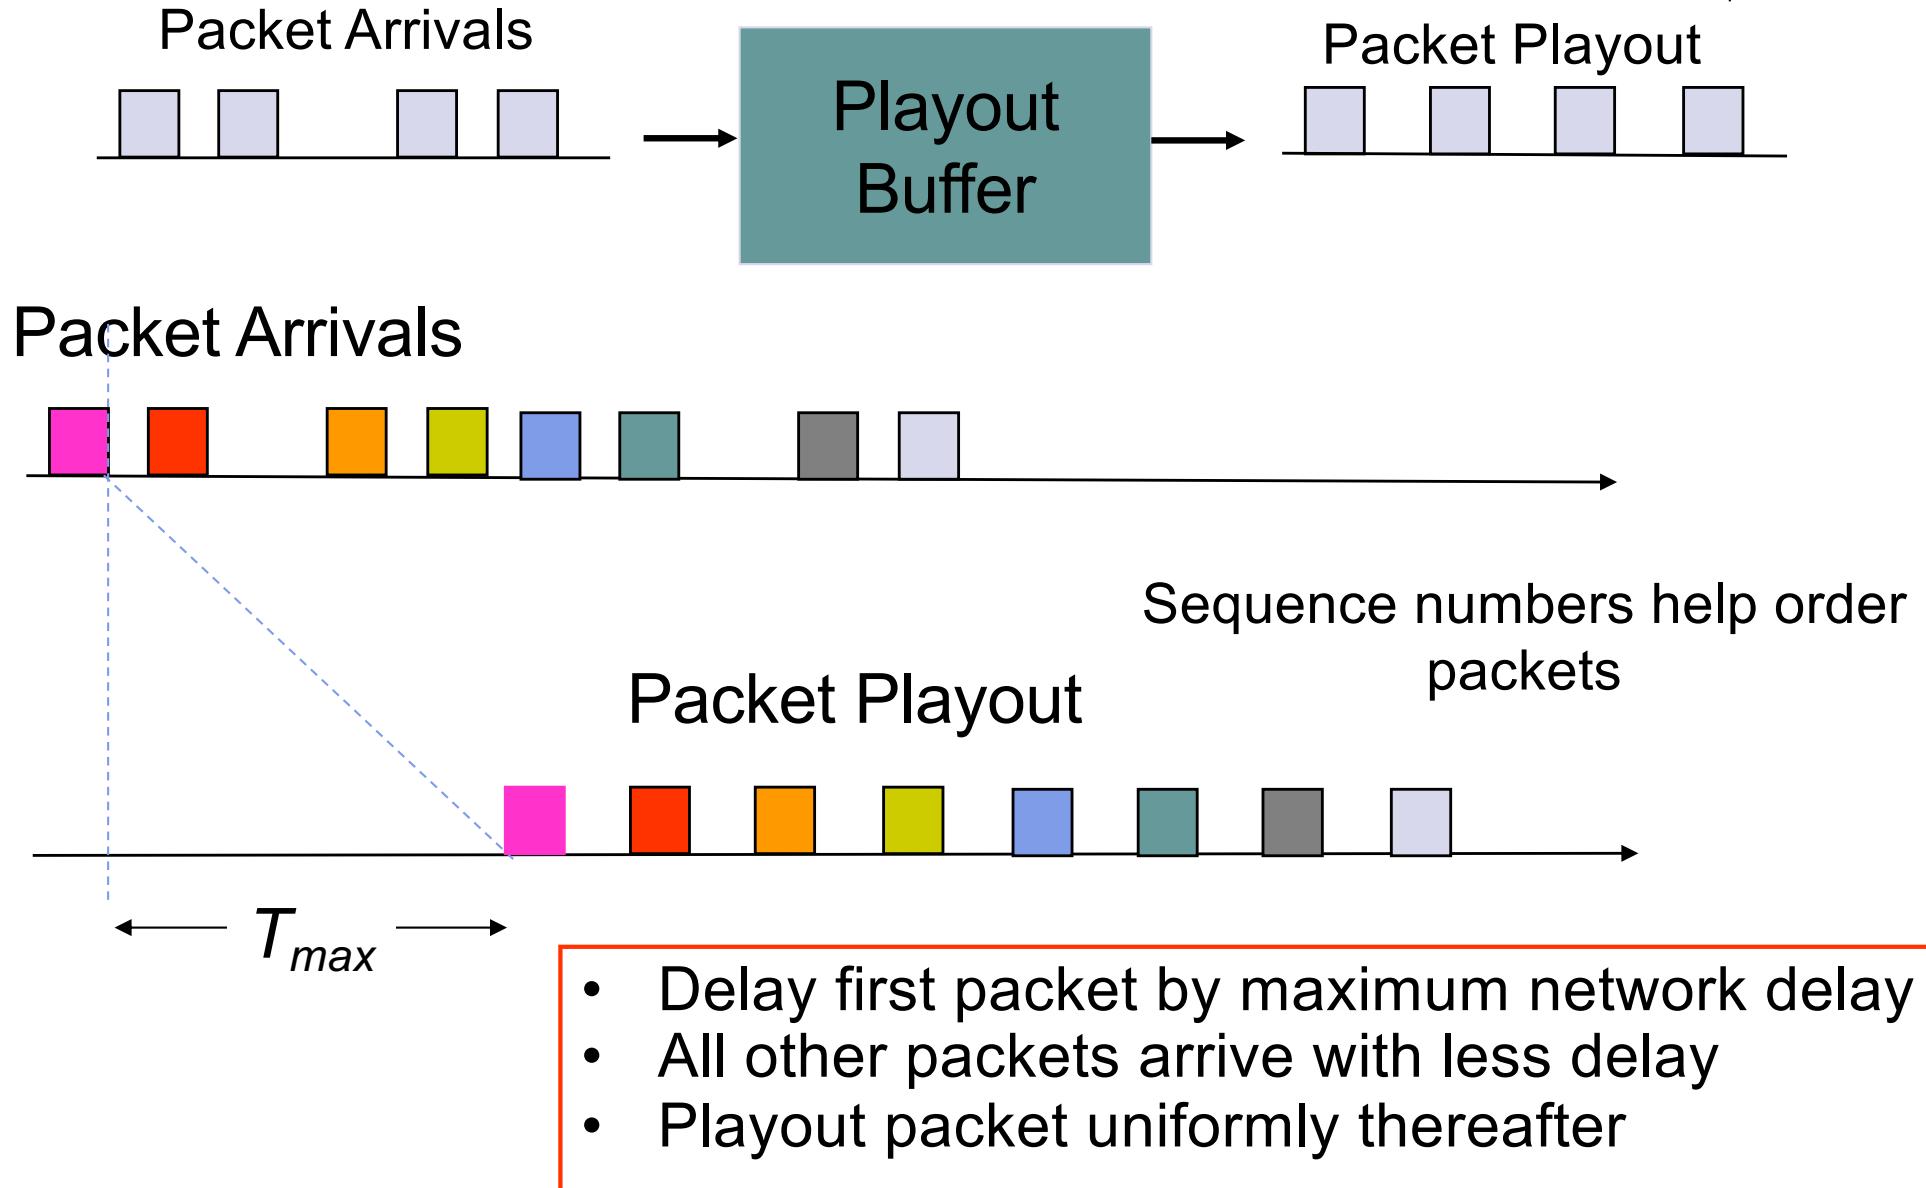
- $W_s = 22 \text{ kHz} \rightarrow 44000 \text{ samples/sec}$
- 16 bits/sample
- $R_s = 16 \times 44000 = 704 \text{ kbps}$ per audio channel
- MP3 uses more powerful compression algorithms: 50 kbps per audio channel

Timing Recovery for Synchronous Services

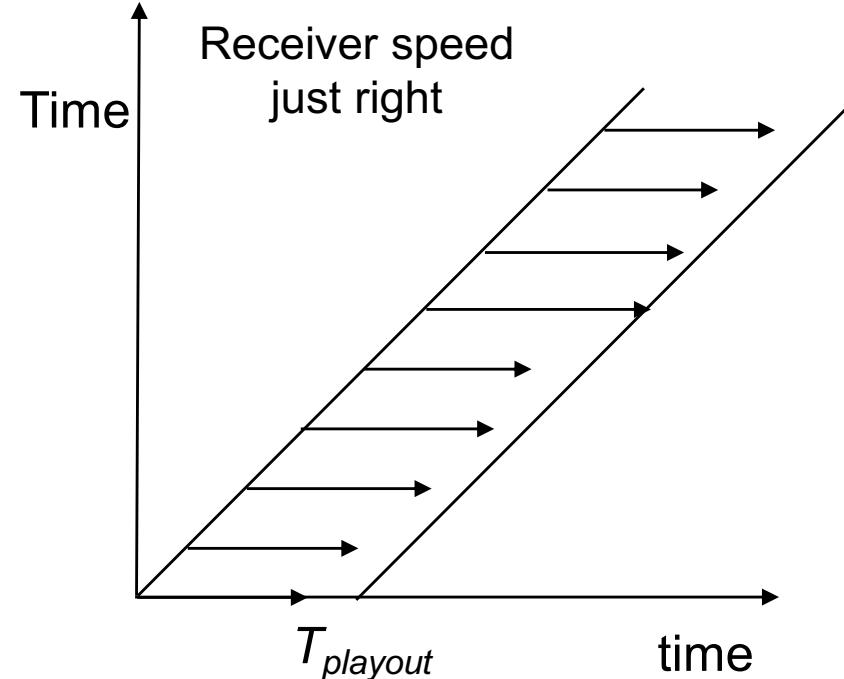
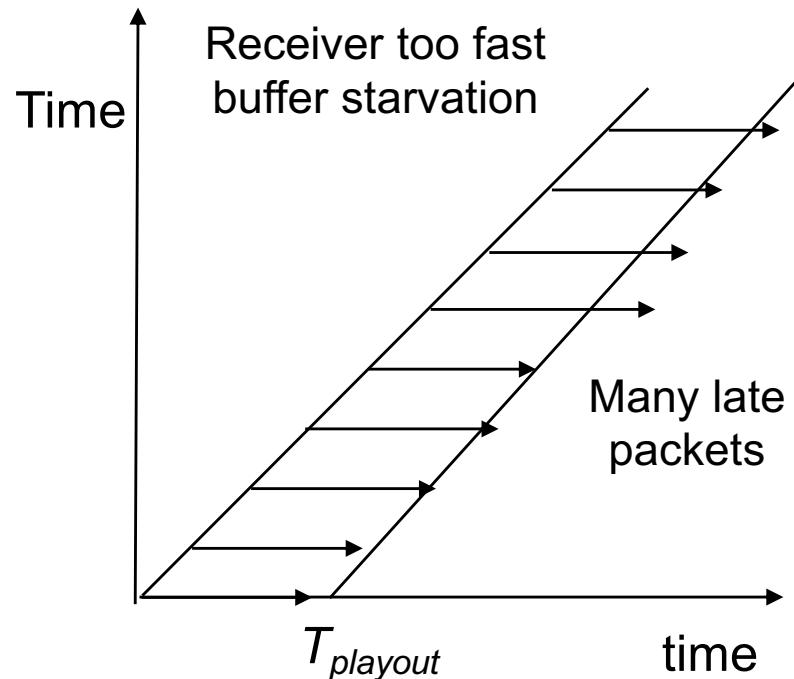
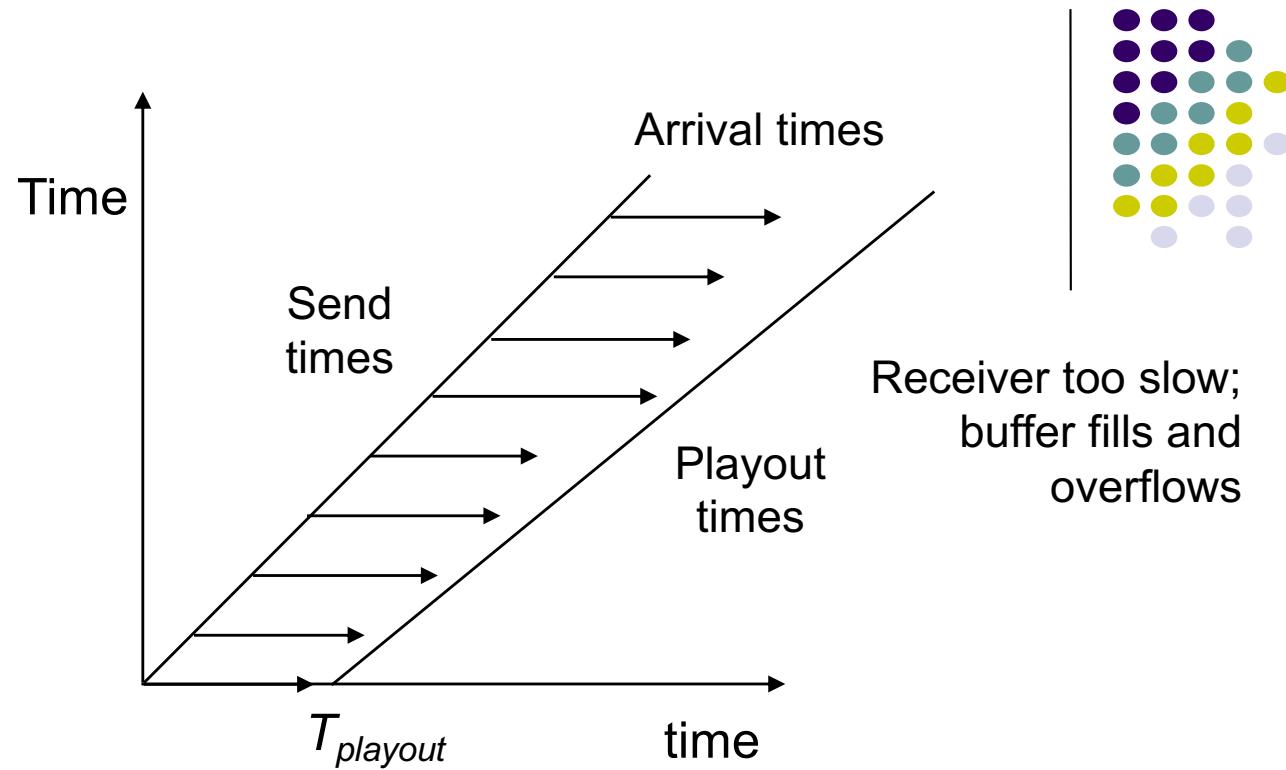


- Voice, audio, or video generate a synchronous information stream
- Information carried by equally-spaced fixed-length packets
- Packet switching introduces random delays
- Timing recovery is needed to re-establish the synchronous nature of the stream

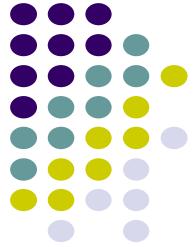
Introduce Playout Buffer



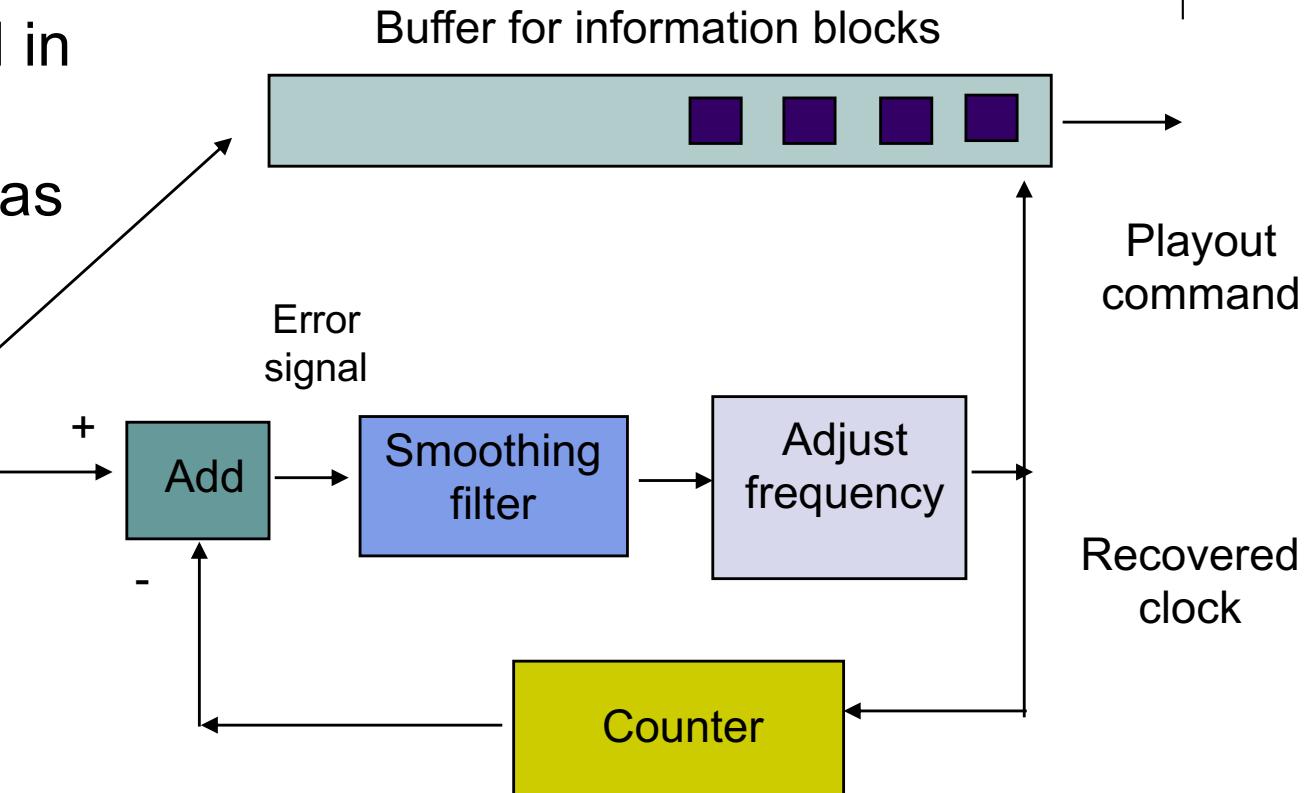
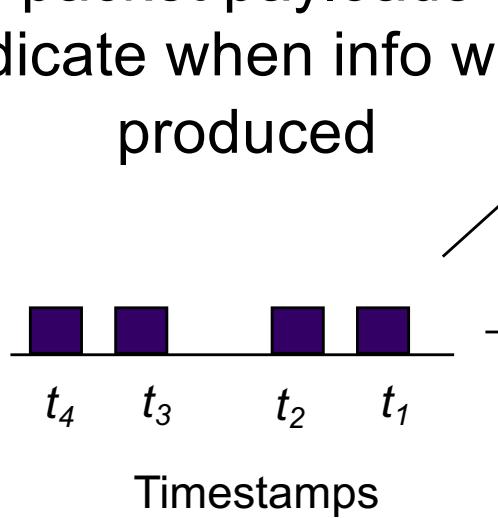
*Playout clock must
be synchronized to
transmitter clock*



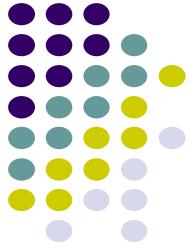
Clock Recovery



Timestamps inserted in packet payloads indicate when info was produced

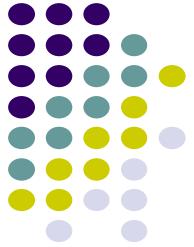


- Counter attempts to replicate transmitter clock
- Frequency of counter is adjusted according to arriving timestamps
- Jitter introduced by network causes fluctuations in buffer & in local clock

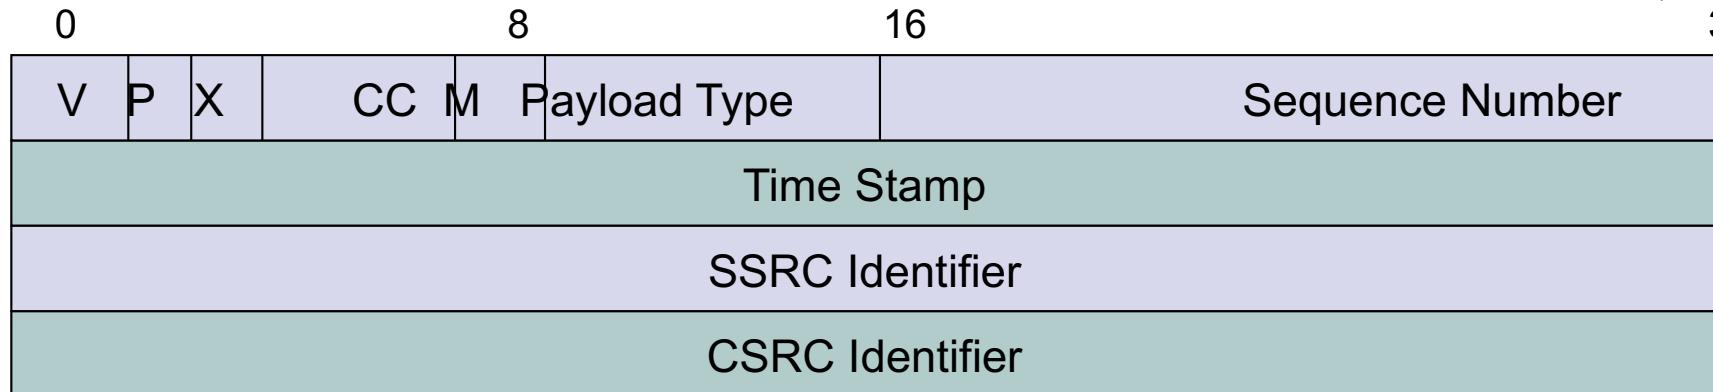


Real-Time Protocol

- RTP designed to support real-time applications such as voice, audio, video
- RTP provides means to carry:
 - Type of information source
 - Sequence numbers
 - Timestamps
- Timing recovery done by higher layer protocol
 - MPEG for video, MP3 for audio
- RTP runs over UDP



RTP Packet Format



- Payload Type: e.g. PCM, MPEG2, ...
- Sequence Number: detect packet loss
- Timestamp: sampling instant of first byte
- Synchronization Source: ID for synch source
- CSRC List: contributing sources to payload

RTP Packet



rtp capture - Ethereal

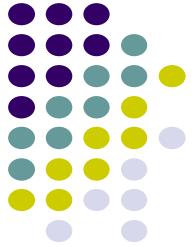
File Edit Capture Display Tools Help

No.	Time	Source	Destination	Protocol	Info
228	19.629299	192.168.2.2	192.168.2.17	RTP	Payload type=ITU-T G.723, SS
229	19.650622	192.168.2.2	192.168.2.17	RTP	Payload type=ITU-T G.723, SS
230	19.687930	192.168.2.2	192.168.2.17	RTP	Payload type=ITU-T G.723, SS
231	19.709274	192.168.2.2	192.168.2.17	RTP	Payload type=ITU-T G.723, SS
232	19.739368	192.168.2.17	192.168.2.2	RTCP	Receiver Report
233	19.746579	192.168.2.2	192.168.2.17	RTP	Payload type=ITU-T G.723, SS
234	19.767155	192.168.2.2	192.168.2.17	RTP	Payload type=ITU-T G.723, SS
225	19.800817	192.168.2.2	192.168.2.17	RTD	Payload type=ITU-T G.723, SS

Frame 230 (78 bytes on wire, 78 bytes captured)
Ethernet II, Src: 00:07:e9:c1:8e:0a, Dst: 00:e0:18:8c:49:b1
Internet Protocol, Src Addr: 192.168.2.2 (192.168.2.2), Dst Addr: 192.168.2.17 (192.168.2.17)
User Datagram Protocol, Src Port: 49608 (49608), Dst Port: 49608 (49608)
Real-Time Transport Protocol
Version: RFC 1889 version (2)
Padding: False
Extension: False
Contributing source identifiers count: 0
Marker: False
Payload type: ITU-T G.723 (4)
Sequence number: 7443
Timestamp: 25064
Synchronization Source identifier: 1704218765
Payload: A8D2F5100B94397121A948BA1448BC6E...

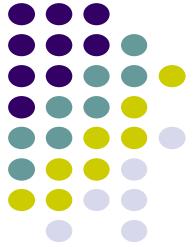
Hex	Dec	ASCII
0000	00 e0 18 8c 49 b1 00 07	e9 c1 8e 0a 08 00 45 a0
0010	00 40 51 19 00 00 80 11	00 00 c0 a8 02 02 c0 a8
0020	02 11 c1 c8 c1 c8 00 2c	f4 c8 80 04 1d 13 00 00
0030	61 e8 65 94 50 8d a8 d2	f5 10 0b 94 39 71 21 a9
0040	48 ba 14 48 bc 6e b2 e5	23 c5 48 fc 0f 0c

Filter: / Reset Apply File: rtp capture



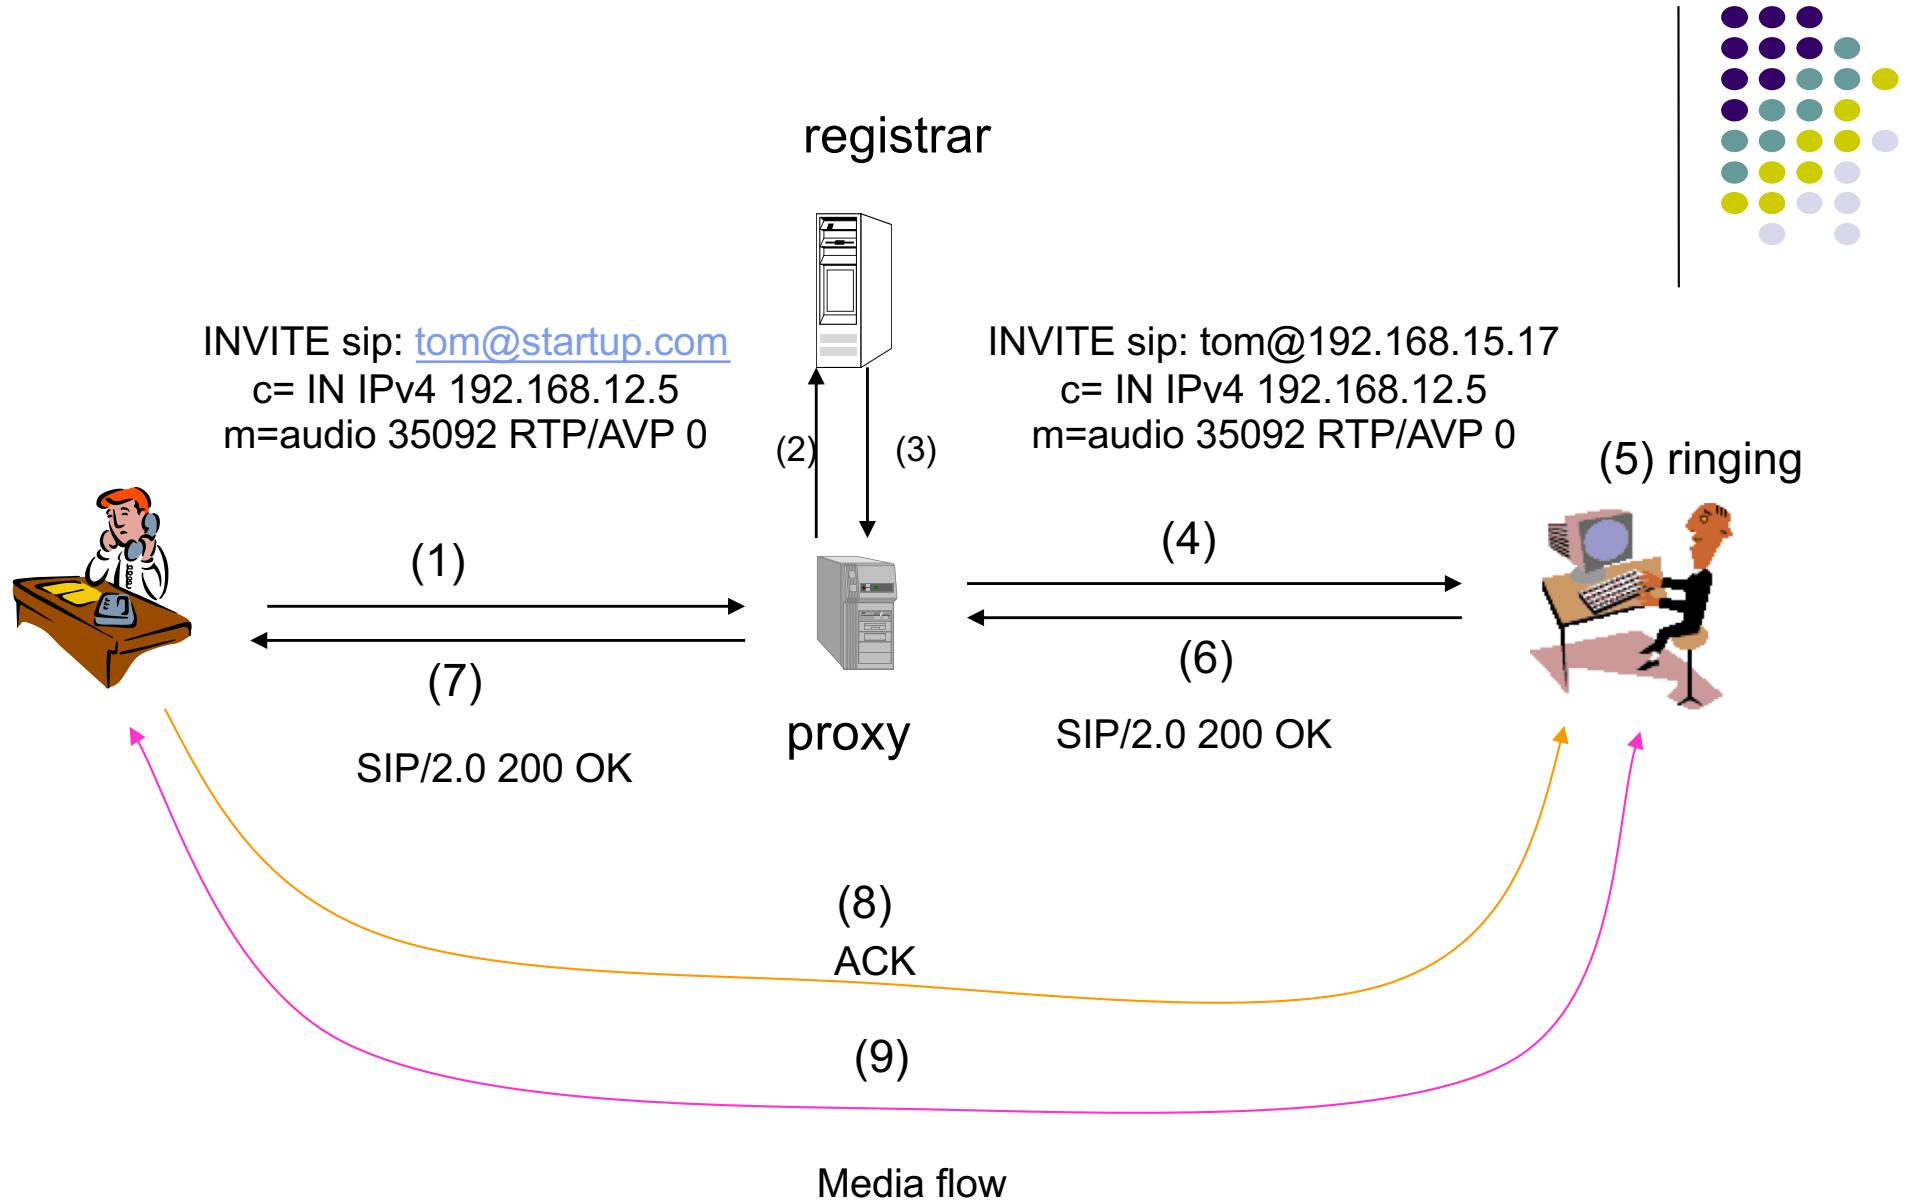
RTP Control Protocol (RTCP)

- RTP companion protocol
- Monitors quality of service at receivers
- Conveys monitored info to senders
- Canonical Name CNAME for each participant
- RTCP Packets
 - Sender Report Packet
 - Receiver Report Packet
 - Source Description (SDES)
 - BYE: end of participation by sender
 - APP: application specific functions



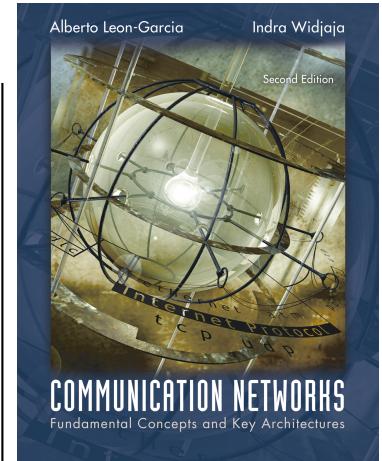
Session Initiation Protocol

- Session: association involving exchange of data between Internet end systems
 - Internet telephone call; multimedia videoconference; instant messaging; event notification
- Session Initiation Protocol
 - Setting up, maintaining, terminating session
 - People & media devices
 - Multicast or mesh of unicast connections
 - Support for user mobility
 - Over UDP or TCP

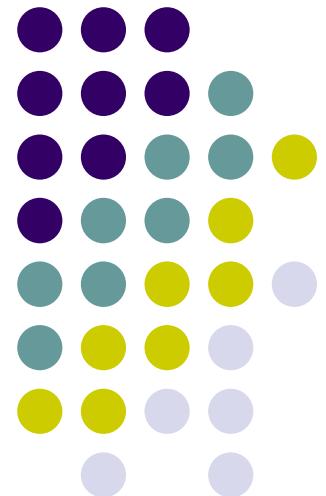


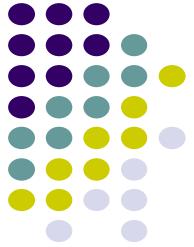
Chapter 2

Applications and Layered Architectures



Sockets

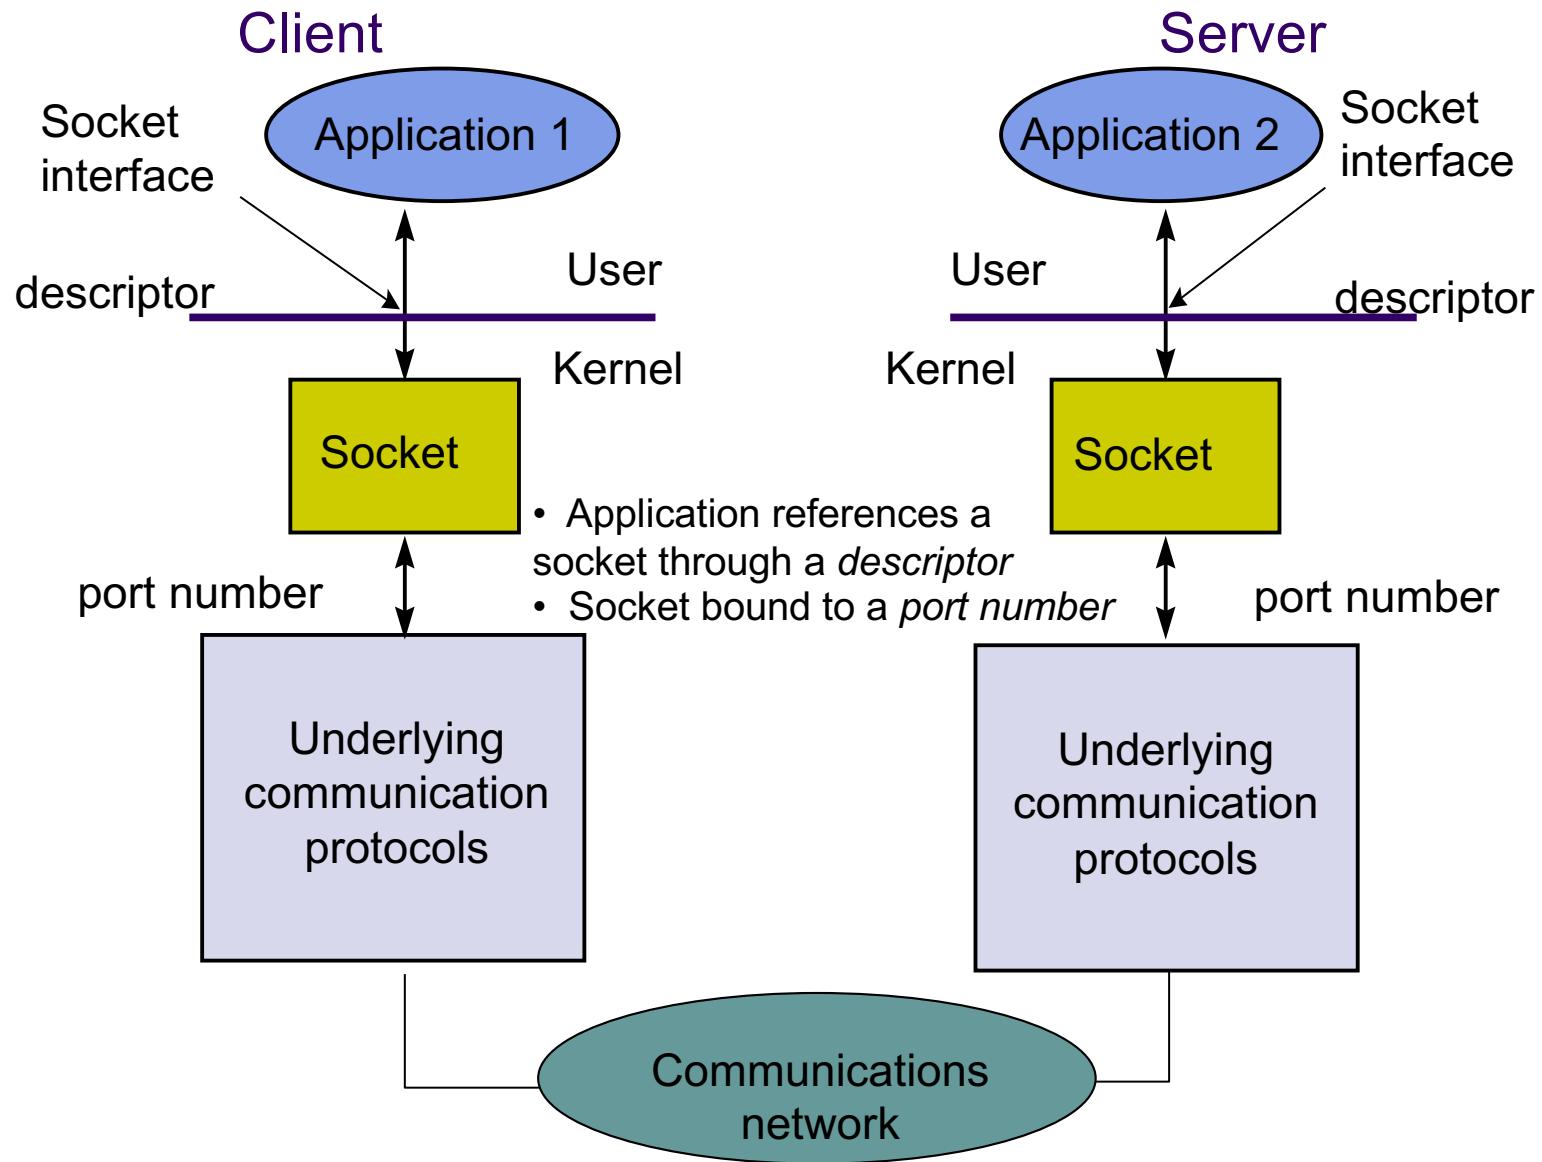
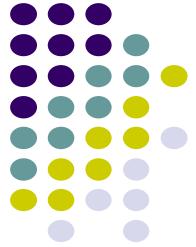


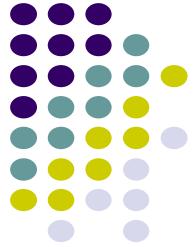


Socket API

- API (Application Programming Interface)
 - Provides a standard set of functions that can be called by applications
- Berkeley UNIX Sockets API
 - Abstraction for applications to send & receive data
 - Applications create sockets that “plug into” network
 - Applications write/read to/from sockets
 - Implemented in the kernel
 - Facilitates development of network applications
 - Hides details of underlying protocols & mechanisms
- Also in Windows, Linux, and other OS's

Communications through Socket Interface



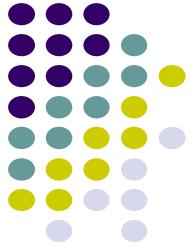


Stream mode of service

Connection-oriented

- First, setup connection between two peer application processes
- Then, reliable bidirectional in-sequence transfer of *byte stream* (boundaries not preserved in transfer)
- Multiple write/read between peer processes
- Finally, connection release
- Uses TCP

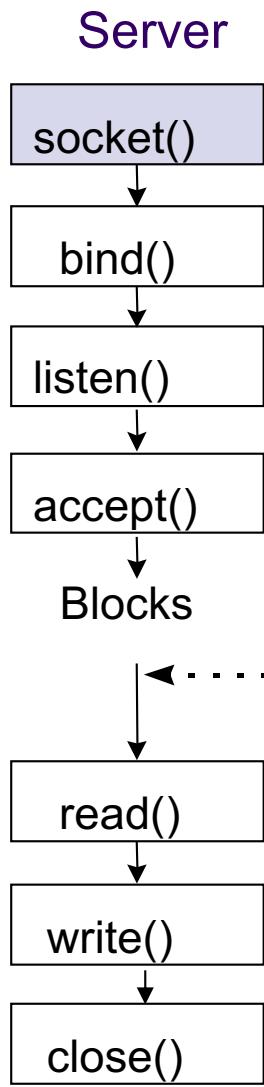
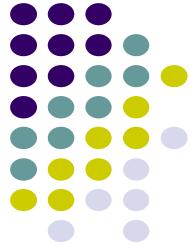
- Connectionless
- Immediate transfer of one block of information (boundaries preserved)
- No setup overhead & delay
- Destination address with each block
- Send/receive to/from multiple peer processes
- Best-effort service only
 - Possible out-of-order
 - Possible loss
- Uses UDP



Client & Server Differences

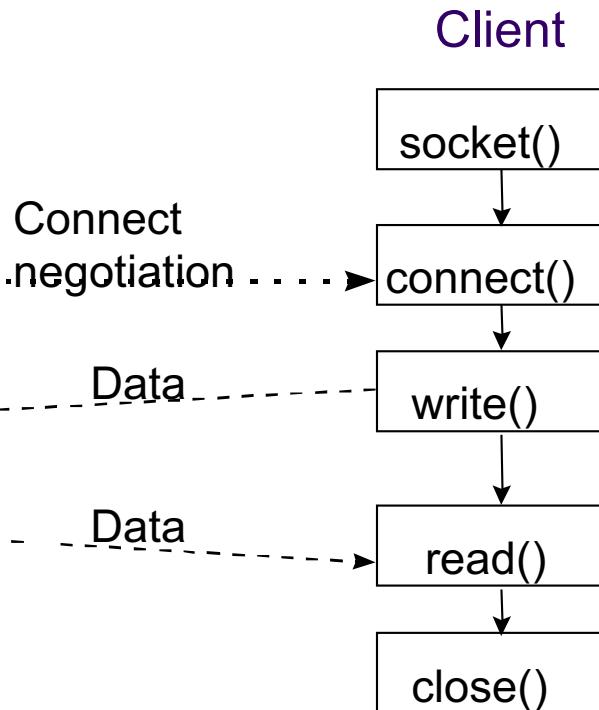
- **Server**
 - Specifies well-known port # when creating socket
 - May have multiple IP addresses (net interfaces)
 - Waits passively for client requests
- **Client**
 - Assigned ephemeral port #
 - Initiates communications with server
 - Needs to know server's IP address & port #
 - DNS for URL & server well-known port #
 - Server learns client's address & port #

Socket Calls for Connection-Oriented Mode

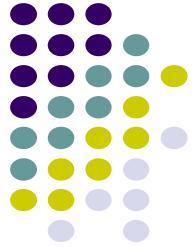


Server does Passive Open

- **socket** creates socket to *listen* for connection requests
- Server specifies type: TCP (stream)
- **socket** call returns: non-negative integer *descriptor*; or -1 if unsuccessful

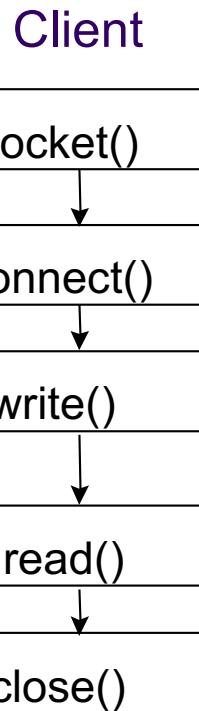
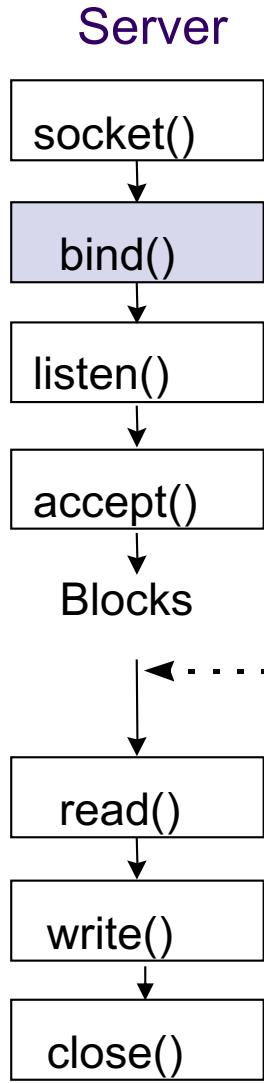


Socket Calls for Connection-Oriented Mode



Server does Passive Open

- **bind** assigns local address & port # to socket with specified descriptor
- Can wildcard IP address for multiple net interfaces
- **bind** call returns: 0 (success); or -1 (failure)
- Failure if port # already in use or if reuse option not set

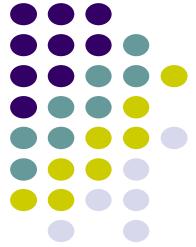


Connect
negotiation

Data

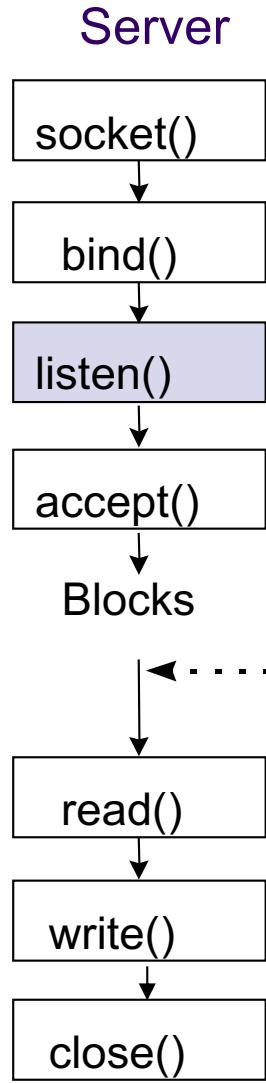
Data

Socket Calls for Connection-Oriented Mode

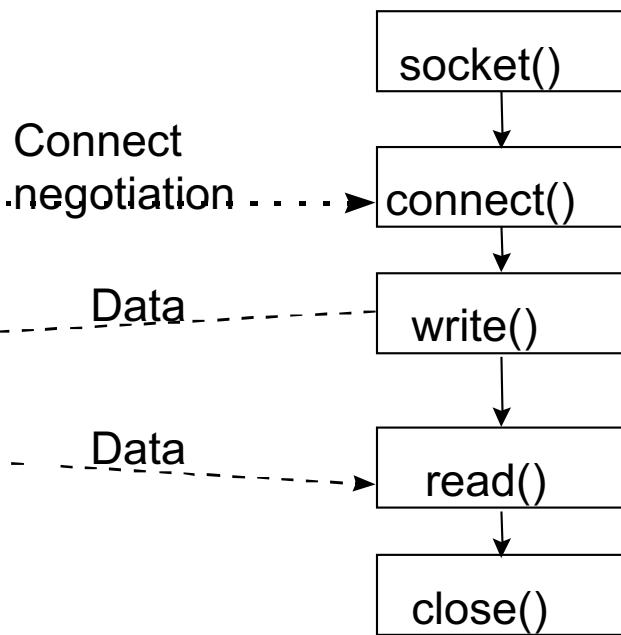


Server does Passive Open

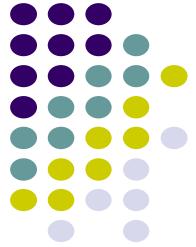
- **listen** indicates to TCP readiness to receive connection requests for socket with given descriptor
- Parameter specifies max number of requests that may be queued while waiting for server to accept them
- **listen** call returns: 0 (success); or -1 (failure)



Client

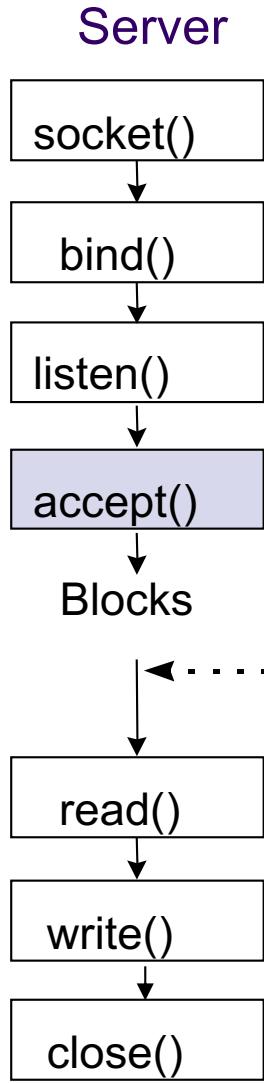


Socket Calls for Connection-Oriented Mode

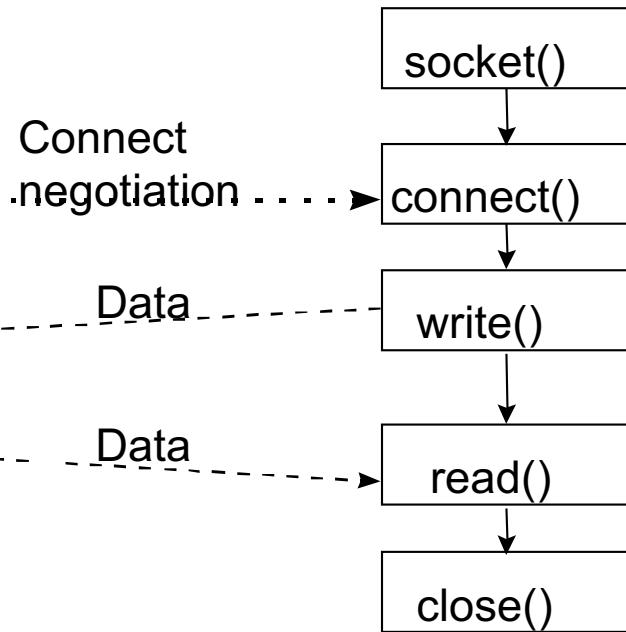


Server does Passive Open

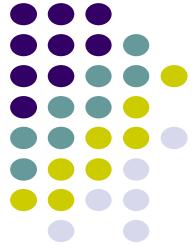
- Server calls `accept` to accept incoming requests
- `accept` blocks if queue is empty



Client

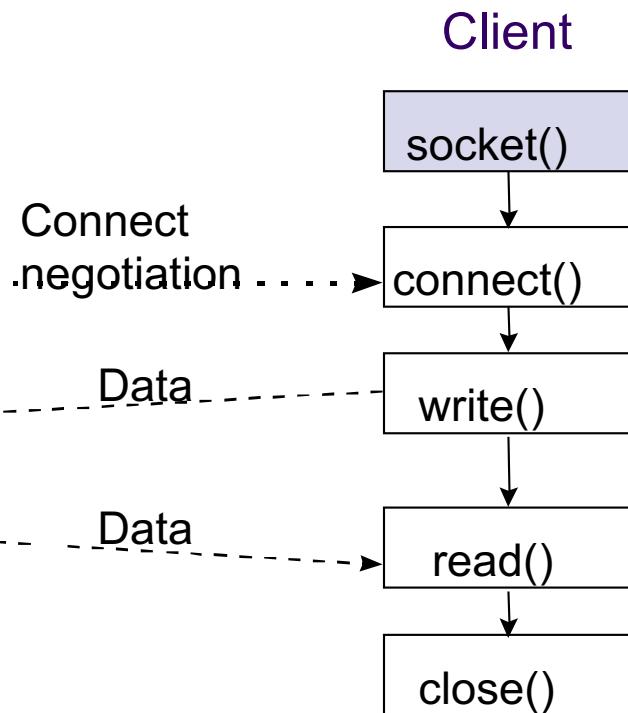
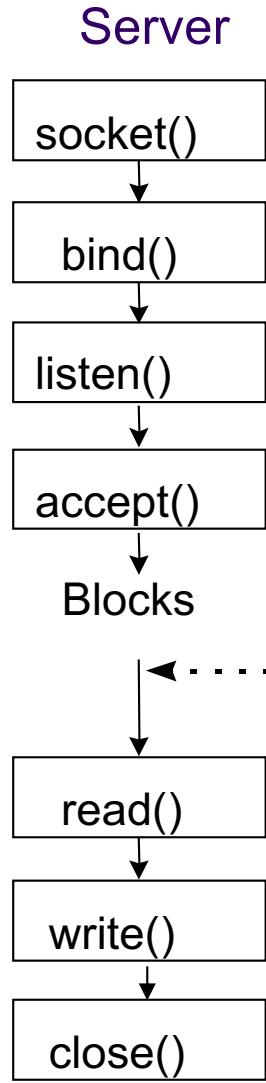


Socket Calls for Connection-Oriented Mode

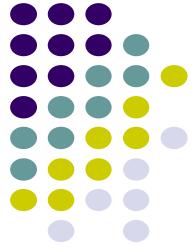


Client does Active Open

- **socket** creates socket to connect to server
- Client specifies type: TCP (stream)
- **socket** call returns: non-negative integer *descriptor*, or -1 if unsuccessful

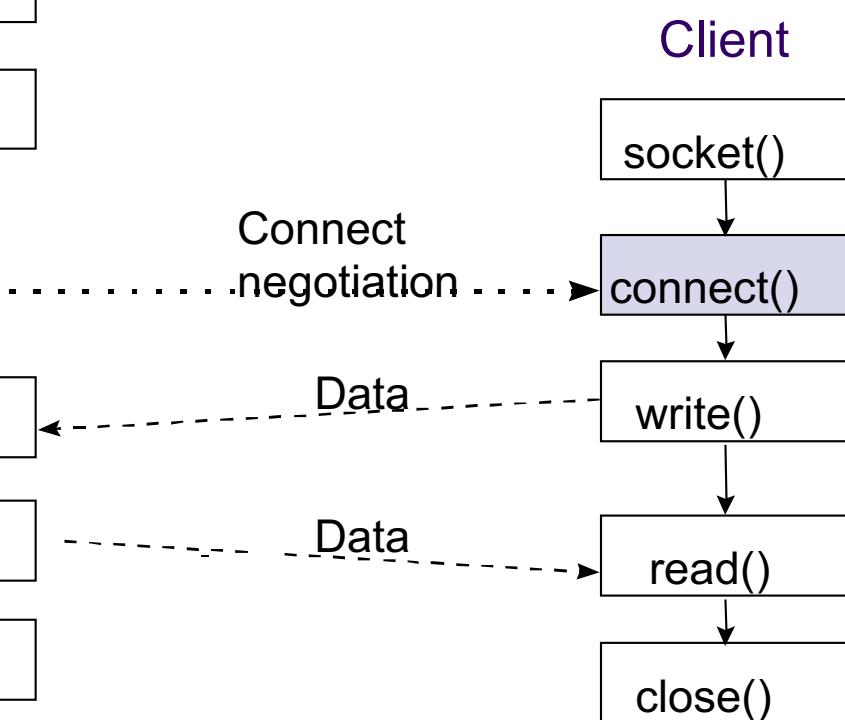
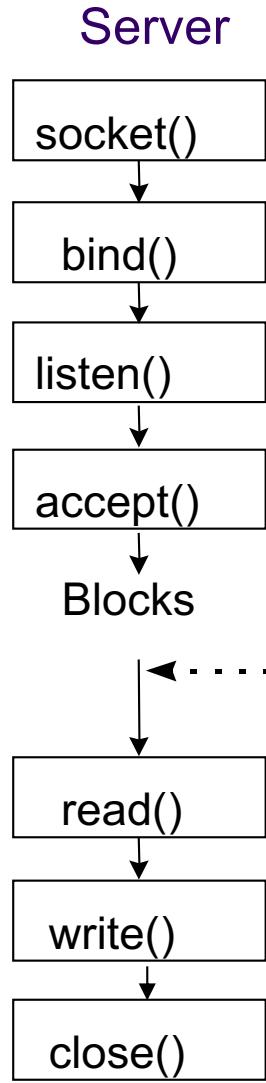


Socket Calls for Connection-Oriented Mode



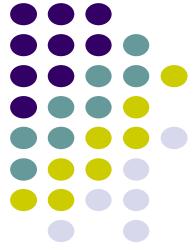
Client does Active Open

- **connect** establishes a connection on the local socket with the specified descriptor to the specified remote address and port #
- **connect** returns 0 if successful; -1 if unsuccessful

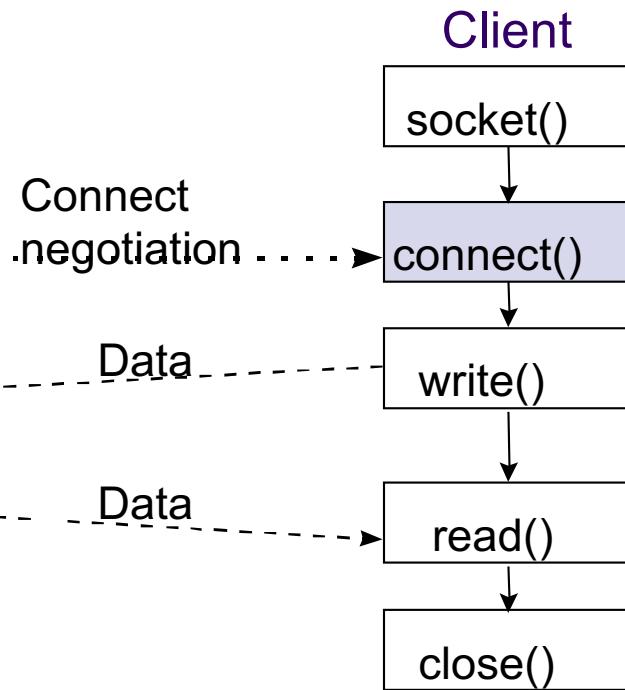
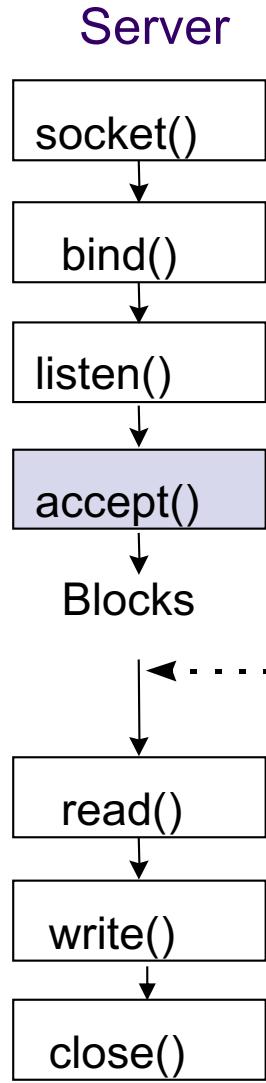


Note: **connect** initiates TCP three-way handshake

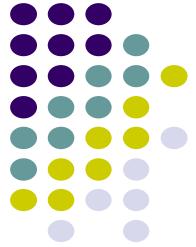
Socket Calls for Connection-Oriented Mode



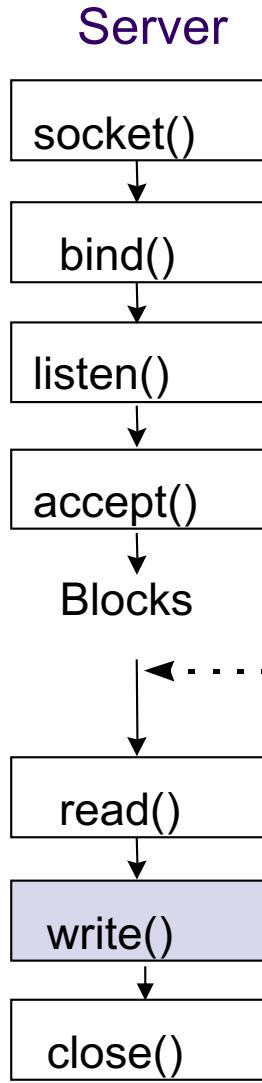
- **accept** wakes with incoming connection request
- **accept** fills client address & port # into address structure
- **accept** call returns: *descriptor of new connection socket* (success); or -1 (failure)
- Client & server use new socket for data transfer
- Original socket continues to listen for new requests



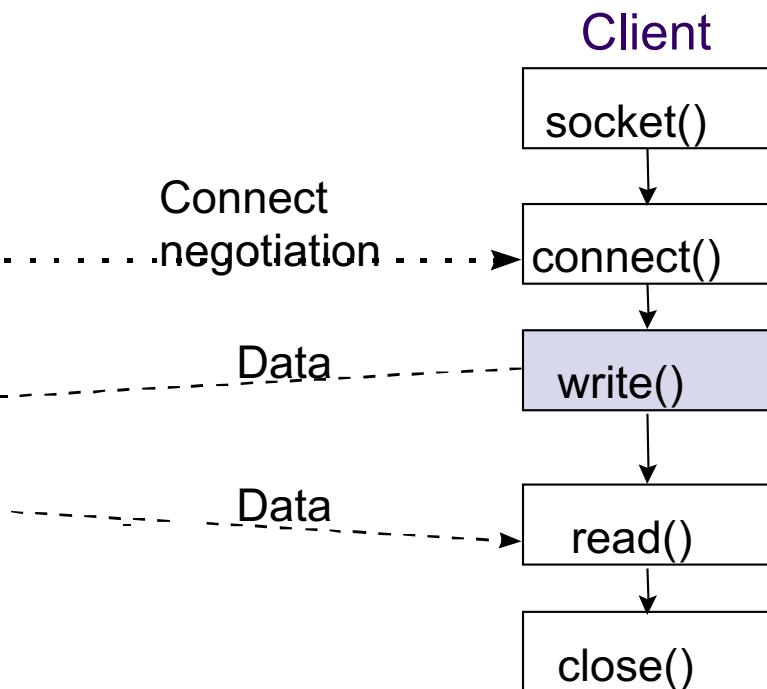
Socket Calls for Connection-Oriented Mode



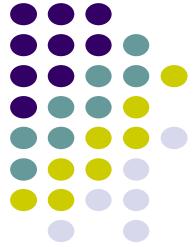
Data Transfer



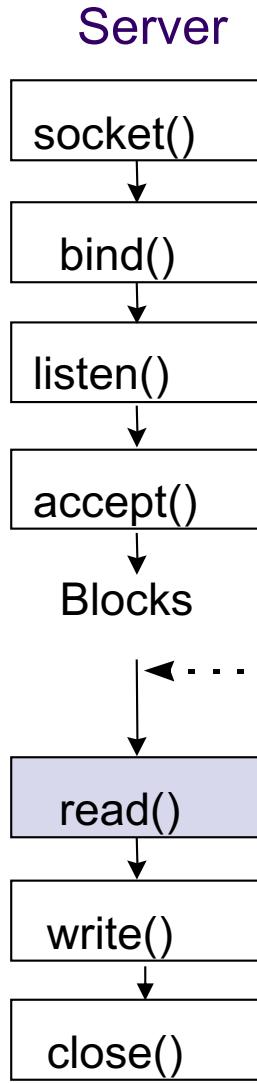
- Client or server call **write** to transmit data into a connected socket
- **write** specifies: socket descriptor; pointer to a buffer; amount of data; flags to control transmission behavior
- **write** call returns: # bytes transferred (success); or -1 (failure); blocks until all data transferred



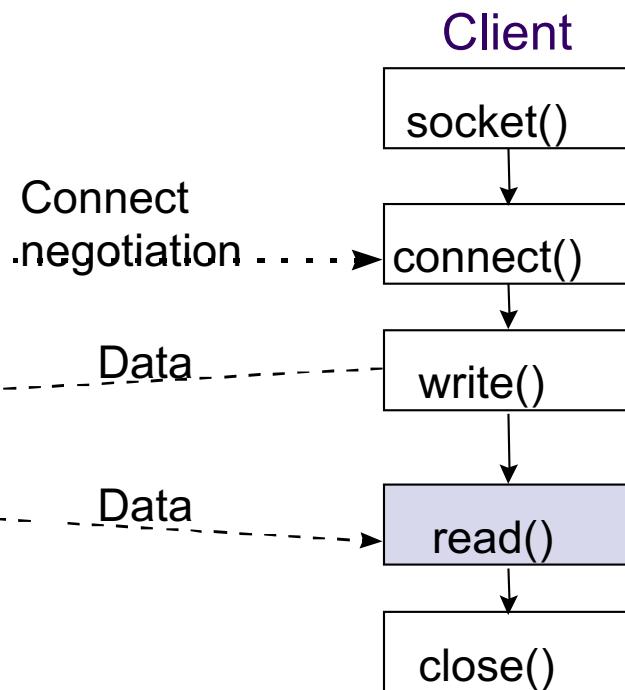
Socket Calls for Connection-Oriented Mode



Data Transfer

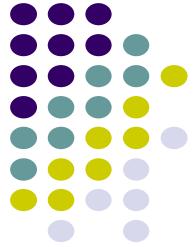


- Client or server call `read` to receive data from a connected socket
- `read` specifies: socket descriptor; pointer to a buffer; amount of data
- `read` call returns: # bytes read (success); or -1 (failure); blocks if no data arrives



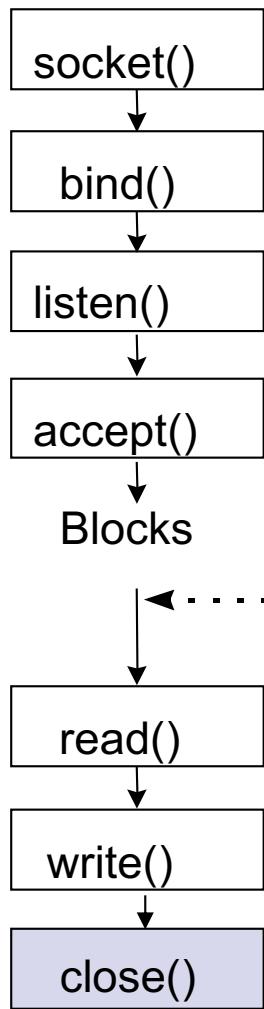
Note: `write` and `read` can be called multiple times to transfer byte streams in both directions

Socket Calls for Connection-Oriented Mode



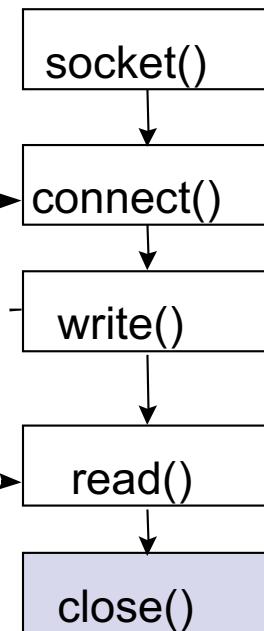
Connection Termination

Server



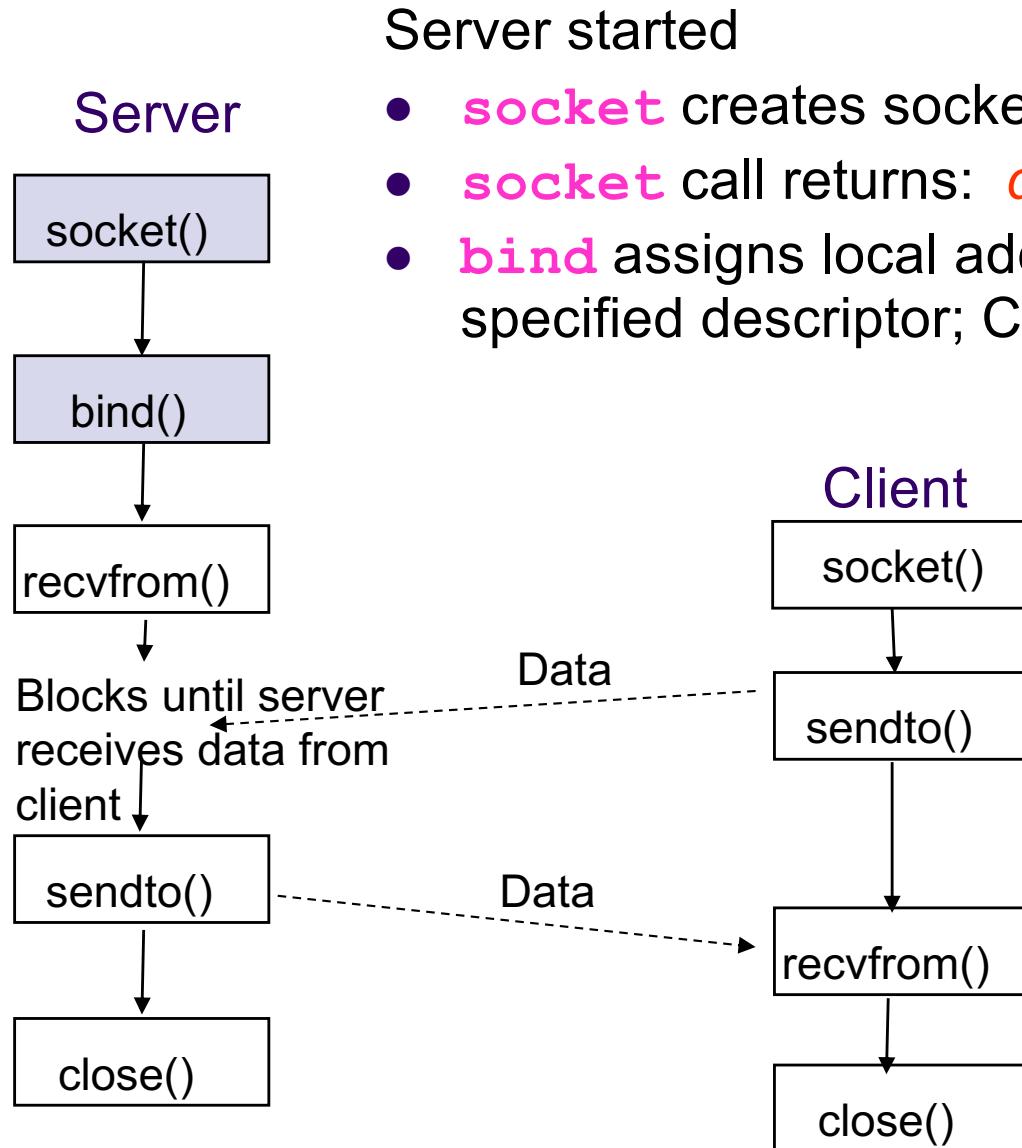
- Client or server call **close** when socket is no longer needed
- **close** specifies the socket descriptor
- **close** call returns: 0 (success); or -1 (failure)

Client



Note: **close** initiates TCP graceful close sequence

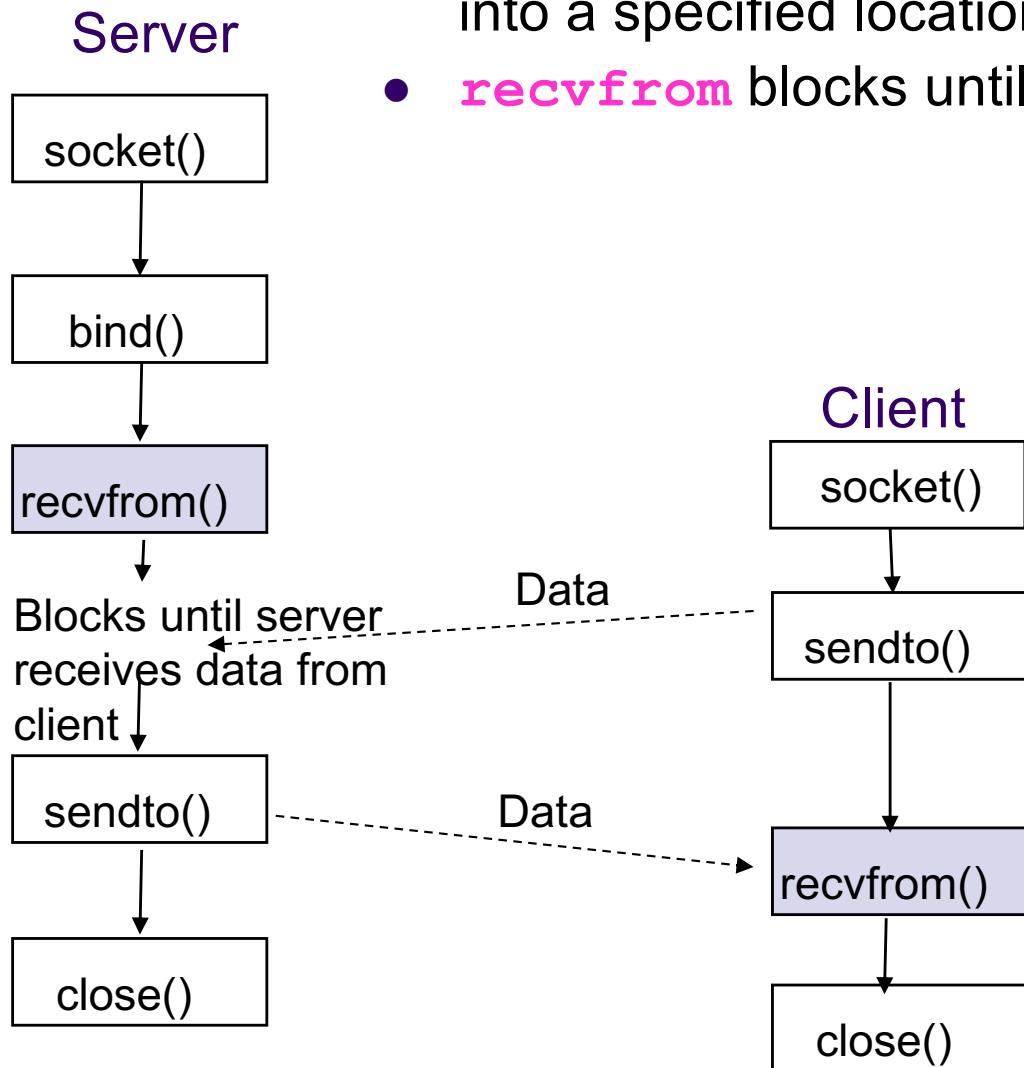
Socket Calls for Connection-Less Mode



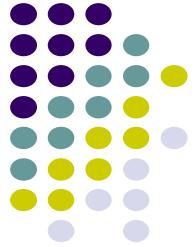
Socket Calls for Connection-Less Mode



- **recvfrom** copies bytes received in specified socket into a specified location
- **recvfrom** blocks until data arrives

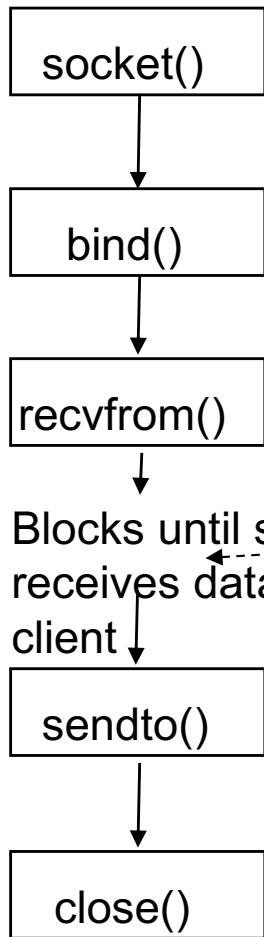


Socket Calls for Connection-Less Mode



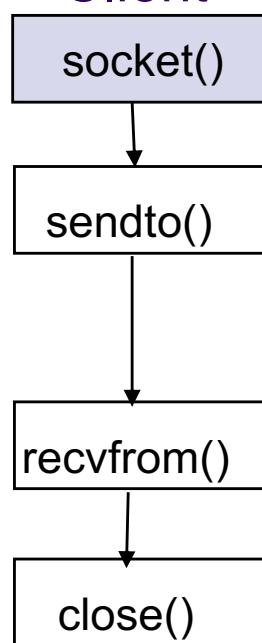
Client started

Server



- `socket` creates socket of type UDP (datagram)
- `socket` call returns: *descriptor*, or -1 if unsuccessful

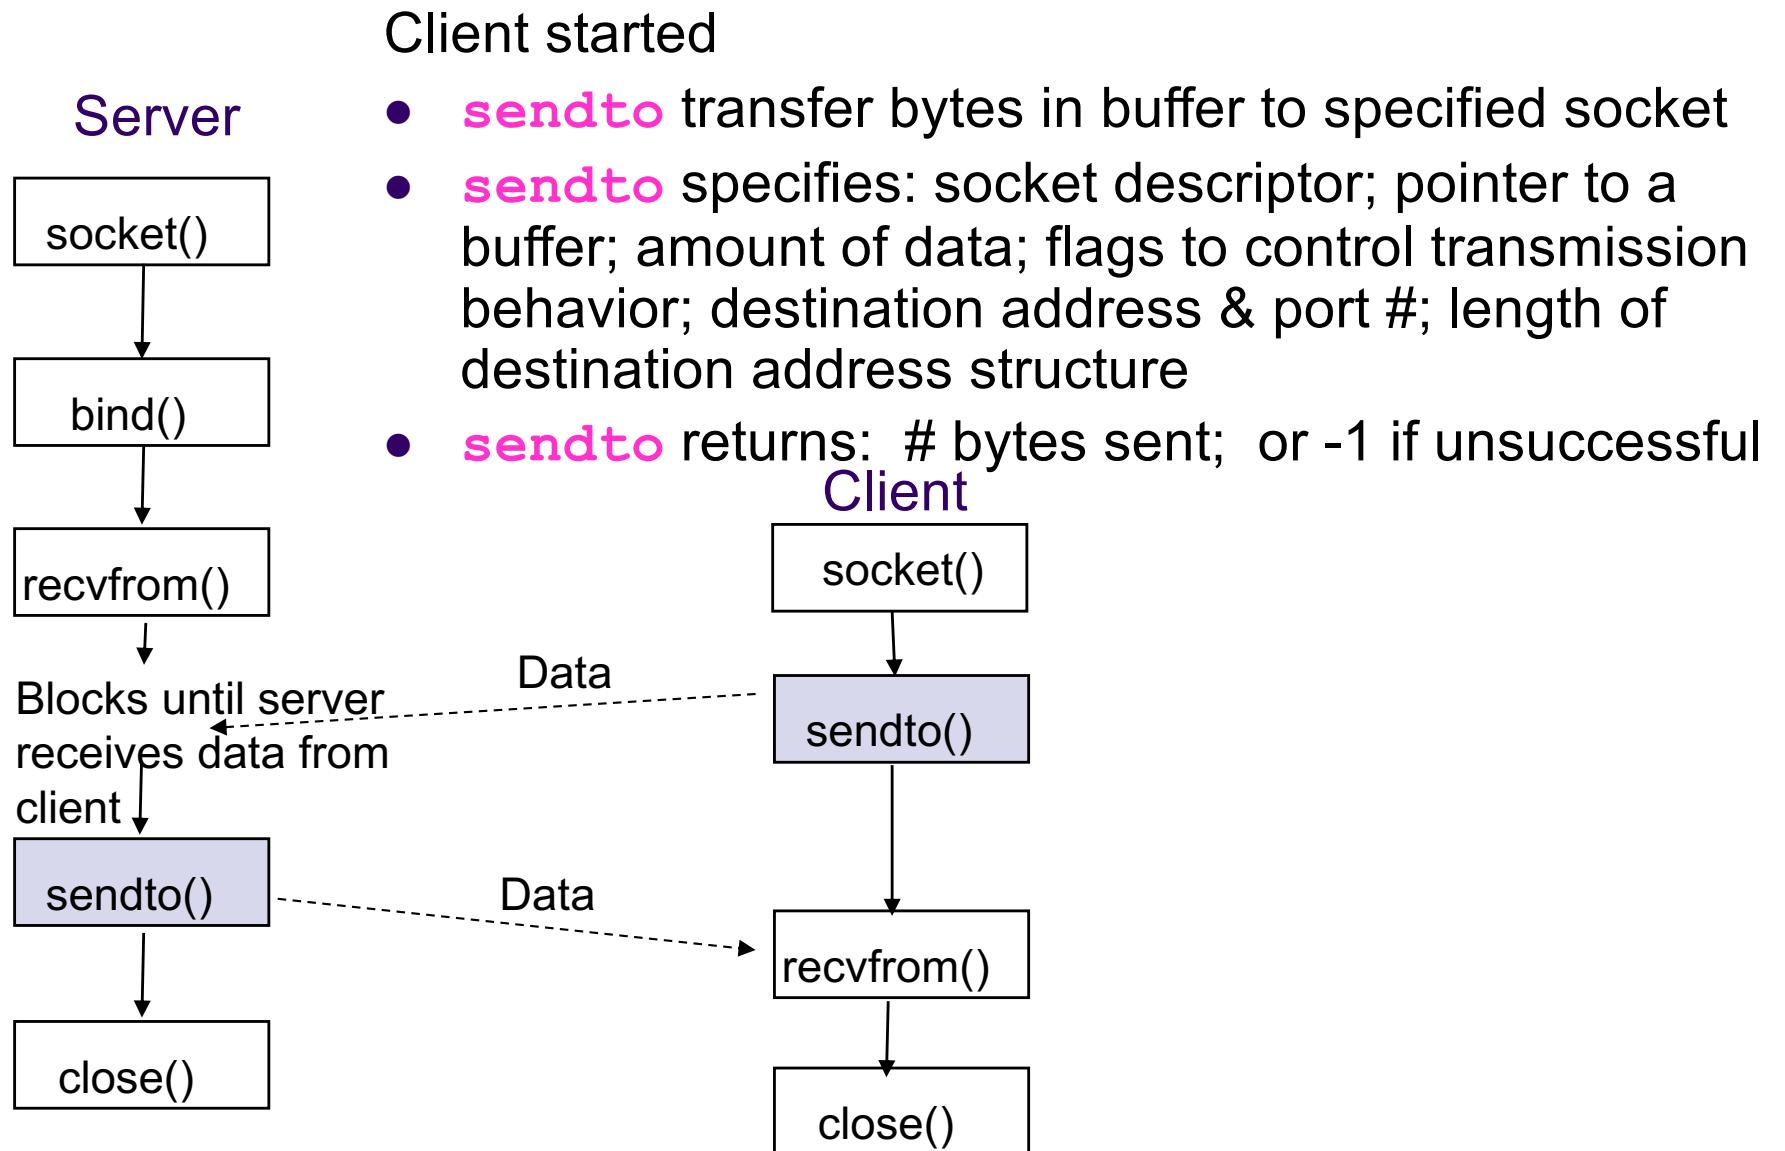
Client



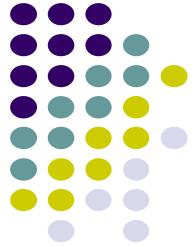
Data

Data

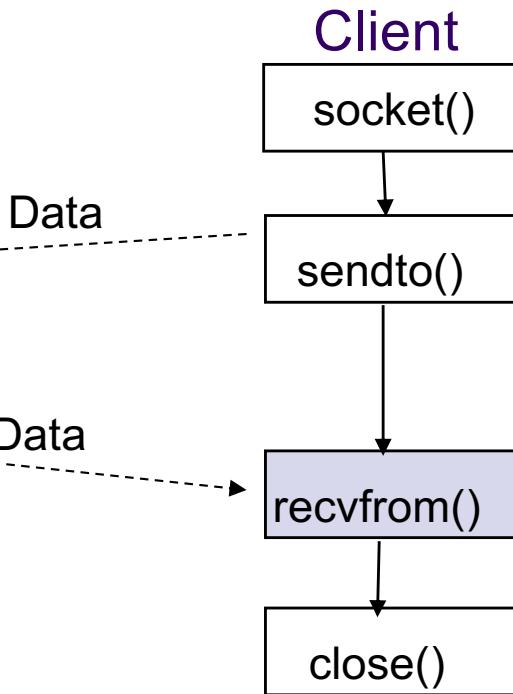
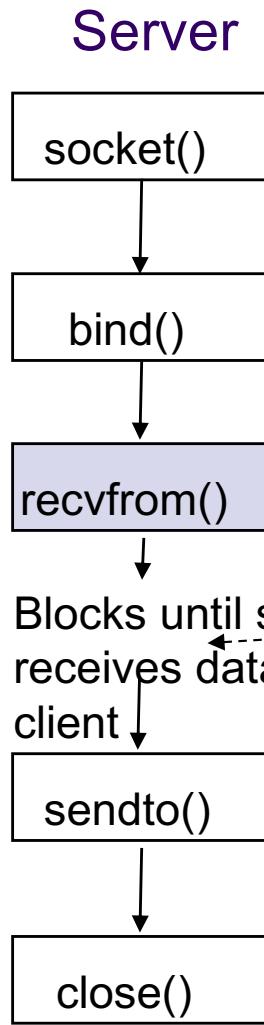
Socket Calls for Connection-Less Mode



Socket Calls for Connection-Less Mode

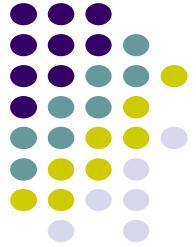


- **recvfrom** wakes when data arrives
- **recvfrom** specifies: socket descriptor; pointer to a buffer to put data; max # bytes to put in buffer; control flags; copies: sender address & port #; length of sender address structure
- **recvfrom** returns # bytes received or -1 (failure)



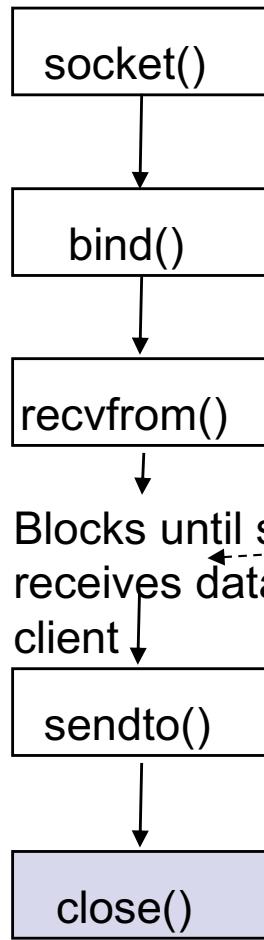
Note: **receivefrom** returns data from at most one **send**, i.e. from one datagram

Socket Calls for Connection-Less Mode



Socket Close

Server



- Client or server call `close` when socket is no longer needed
- `close` specifies the socket descriptor
- `close` call returns: 0 (success); or -1 (failure)

Client

